

The Setup: Paving the path to the OS

Prof. Dr. Antônio Augusto Fröhlich
UFSC / LISHA

<https://lisha.ufsc.br/~guto>

Almost there at the OS

- BIOS brought the system on
 - BIST, POST, hooks
 - First instruction fetched – 0x7c00
 - Lots of “jmp” so far, no calls, no stack
- **Bootstrap** loaded the OS
 - 16-bit assembly code
 - rudimentary process model
 - Configured memory and bus (A20)
 - Entered 32-bit Protected Mode (for GCC code)
- **Setup**
 - Bring the machine into a usable state

The Role of SETUP

- It “sets up” extremely architecture-dependent data structures
 - For the x86 PC
 - IDT – Interrupt Descriptor Table
 - GDB – Global Descriptor Table
 - Paging
 - PCI
 - FPU
 - Timers and interrupts
 - etc...
- Delivers
 - A basic memory model (flat)
 - A basic process model (single task)

System Information

Boot Map



```
struct Boot_Map
{
    Phy_Addr mem_base;           // Memory base
    unsigned int mem_size;       // Memory size (in bytes)
    int cpu_type;                // Processor type
    int cpu_clock;               // Processor clock frequency in Hz
    int n_threads;               // Max number of threads
    int n_tasks;                 // Max number of tasks
    unsigned short host_id;      // The local host id (-1 => RARP)
    unsigned short n_nodes;      // Number of nodes in SAN
    int img_size;                // Boot image size in bytes
    int setup_off;               // SETUP offset in the boot image
    int system_off;              // OS offset in the boot image
    int loader_off;              // LOADER offset in the boot image
    int app_off;                 // APPs offset in the boot image
};
```

System Information

Physical Memory Map



```
struct Physical_Memory_Map
{
    Phy_Addr app_lo;           // Application memory's lowest address
    Phy_Addr app_hi;           // Application memory's highest address
    Phy_Addr int_vec;          // Interrupt Vector
    Phy_Addr sys_pt;           // System Page Table
    Phy_Addr sys_pd;           // System Page Directory
    Phy_Addr sys_info;         // System Info
    Phy_Addr phy_mem_pts;      // Page tables to map the whole phy memory
    Phy_Addr io_mem_pts;       // Page tables to map the IO address space
    Phy_Addr sys_code;         // OS Code Segment
    Phy_Addr sys_data;         // OS Data Segment
    Phy_Addr sys_stack;        // OS Stack Segment
    Phy_Addr free;             // Free memory base
    unsigned int free_size;    // Free memory size (in frames)
    Phy_Addr mach1;            // Machine specific entries
    Phy_Addr mach2;
    Phy_Addr mach3;
    Phy_Addr free1_base;       // First free memory chunk base address
    Phy_Addr free1_top;        // First free memory chunk top address
    Phy_Addr free2_base;       // Second free memory chunk base address
    Phy_Addr free2_top;        // Second free memory chunk top address
};
```

System Information

Logical Memory Map



```
struct Logical_Memory_Map
{
    Log_Addr base;           // Lowest valid logical address
    Log_Addr top;            // Highest valid logical address
    Log_Addr app_lo;         // Application memory lowest address
    Log_Addr app_entry;      // First application's entry point
    Log_Addr app_code;       // First application's code base address
    Log_Addr app_data;       // First application's data base address
    Log_Addr app_hi;         // Application memory highest address
    Log_Addr phy_mem;        // Whole physical memory (contiguous)
    Log_Addr io_mem;         // IO address space
    Log_Addr int_vec;        // Interrupt Vector
    Log_Addr sys_pt;         // System Page Table
    Log_Addr sys_pd;         // System Page Directory
    Log_Addr sys_info;       // System Info
    Log_Addr sys_code;       // OS Code Segment
    Log_Addr sys_data;       // OS Data Segment
    Log_Addr sys_stack;      // OS Stack Segment
    Log_Addr mach1;          // Machine specific entries
    Log_Addr mach2;
    Log_Addr mach3;
};
```

System Information

I/O Memory Map



```
struct IO_Memory_Map
{
    int locator;
    Phy_Addr phy_addr;
    Log_Addr log_addr;
    unsigned int size;
};
```

System Information

```
class System_Info
{
public:
    typedef unsigned int Log_Addr;
    typedef unsigned int Phy_Addr;

    unsigned int mem_size;    // Memory size (in pages)
    unsigned int mem_free;    // Free memory (in pages)
    unsigned int iomm_size;
    Boot_Map bm;
    Physical_Memory_Map pmm;
    Logical_Memory_Map lmm;
    IO_Memory_Map iomm[];
};
```

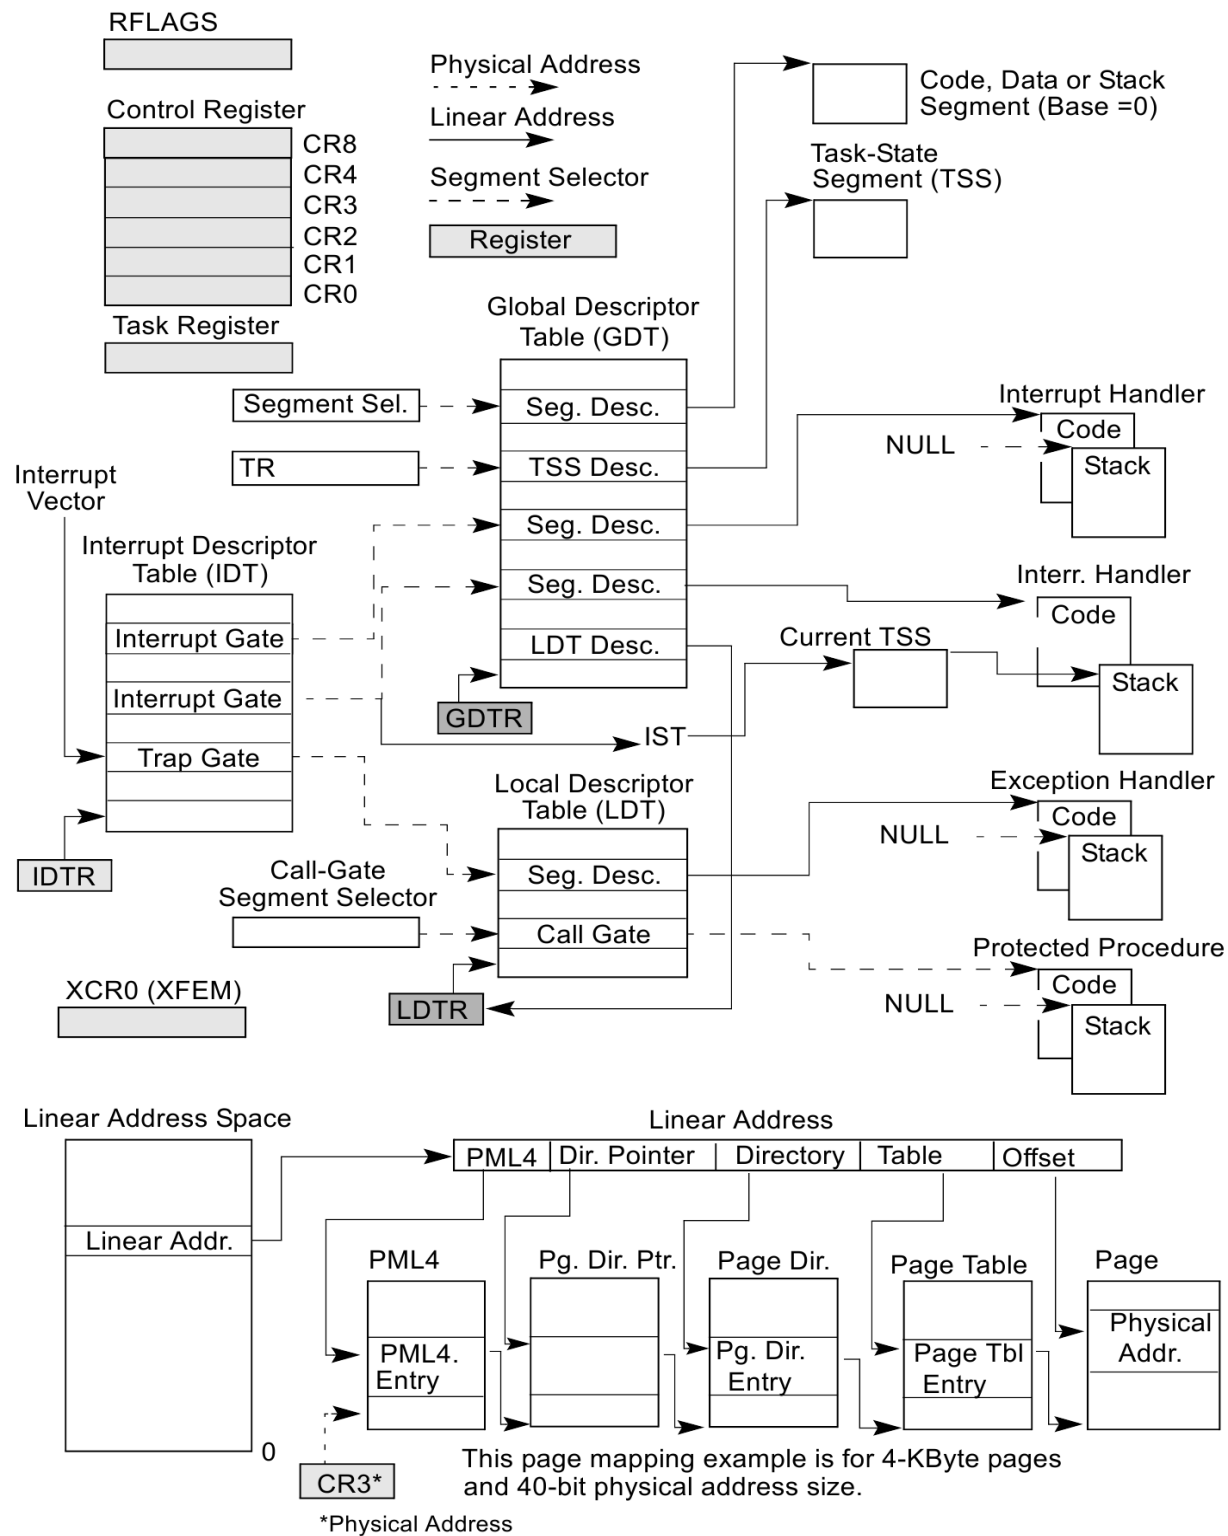

Starting SETUP

- SETUP begins running where the bootstrap loaded it
 - Outside the address space it was compiled for
 - Using only PC-relative addressing
 - But it has access to System_Info
 - Boot image at BOOT_IMAGE_ADDR
- Verify system image
 - Validate ELF image
 - Check if ELF entry point is accessible
- Relocate the SETUP (as it runs)
 - Move boot image to after SETUP
 - Setup a single-page stack for SETUP
- Call actual SETUP main

System Memory

- Allocate (reserve) memory for all entities we have to setup
 - Start at the highest address
 - Flat memory model
 - Reserve memory for
 - IDT – Interrupt Descriptor Table
 - GDT – Global Descriptor Table
 - SysPT
 - SysPD
 - System_Info
 - IO space
 - OS code, data, and stack
 - Mark rest of memory as free to applications

x86



x86 GDT

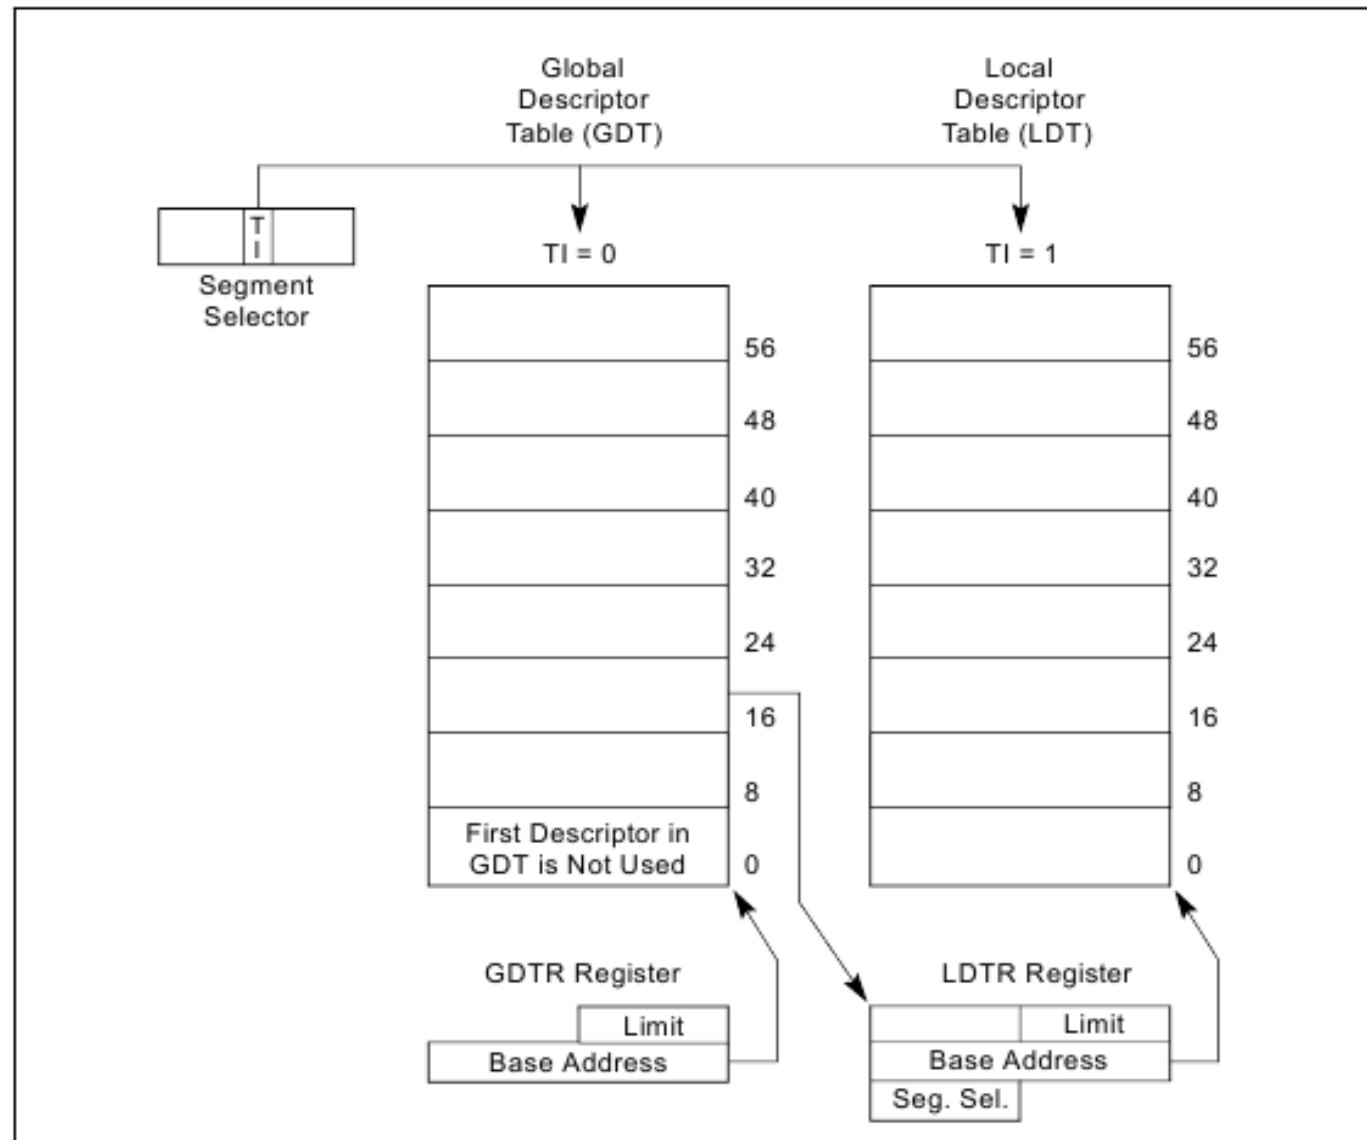
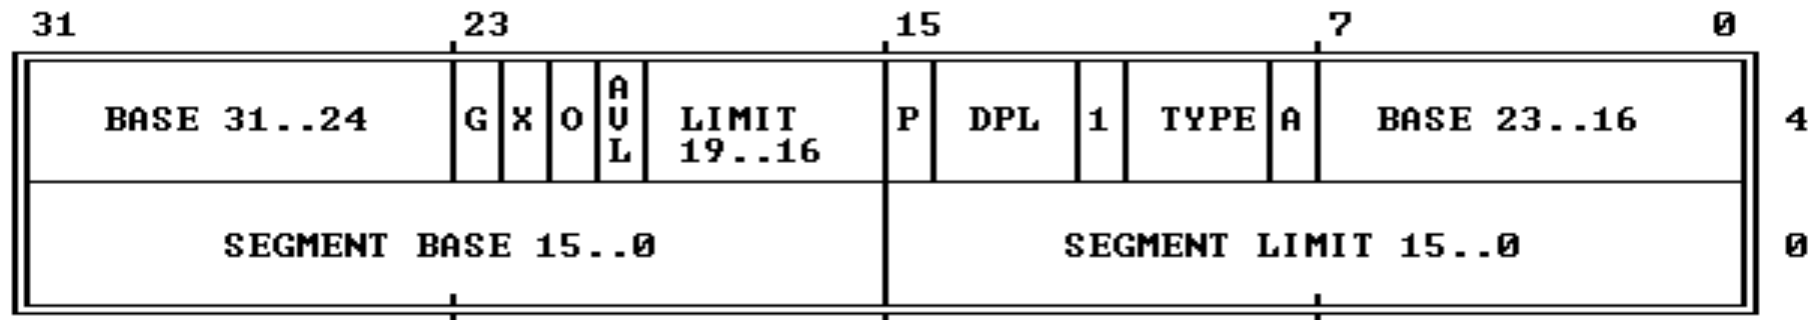


Figure 3-10. Global and Local Descriptor Tables

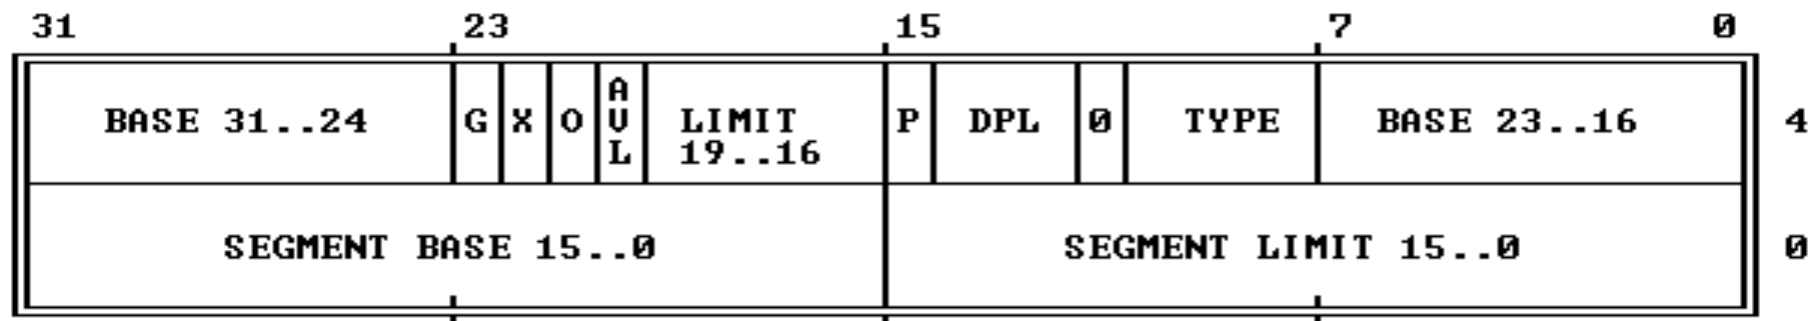
x86 GDT entries

Figure 5-3. General Segment-Descriptor Format

DESCRIPTORS USED FOR APPLICATIONS CODE AND DATA SEGMENTS

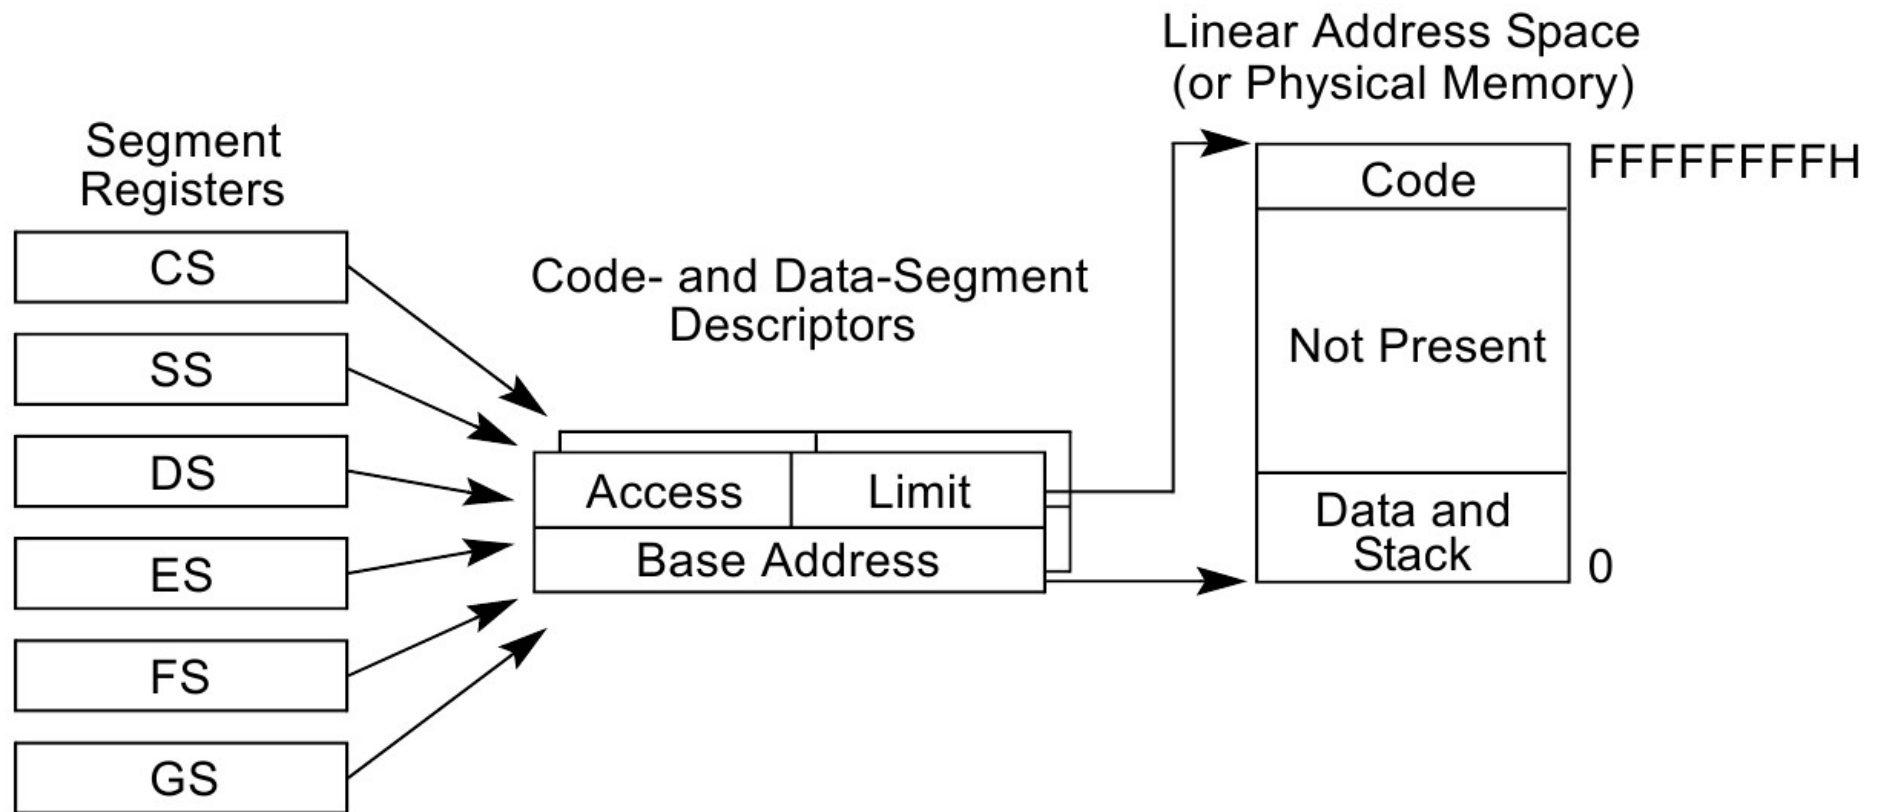


DESCRIPTORS USED FOR SPECIAL SYSTEM SEGMENTS



- A — ACCESSED
- AUL — AVAILABLE FOR USE BY SYSTEMS PROGRAMMERS
- DPL — DESCRIPTOR PRIVILEGE LEVEL
- G — GRANULARITY
- P — SEGMENT PRESENT

x86 Flat Memory Model



x86 Paging

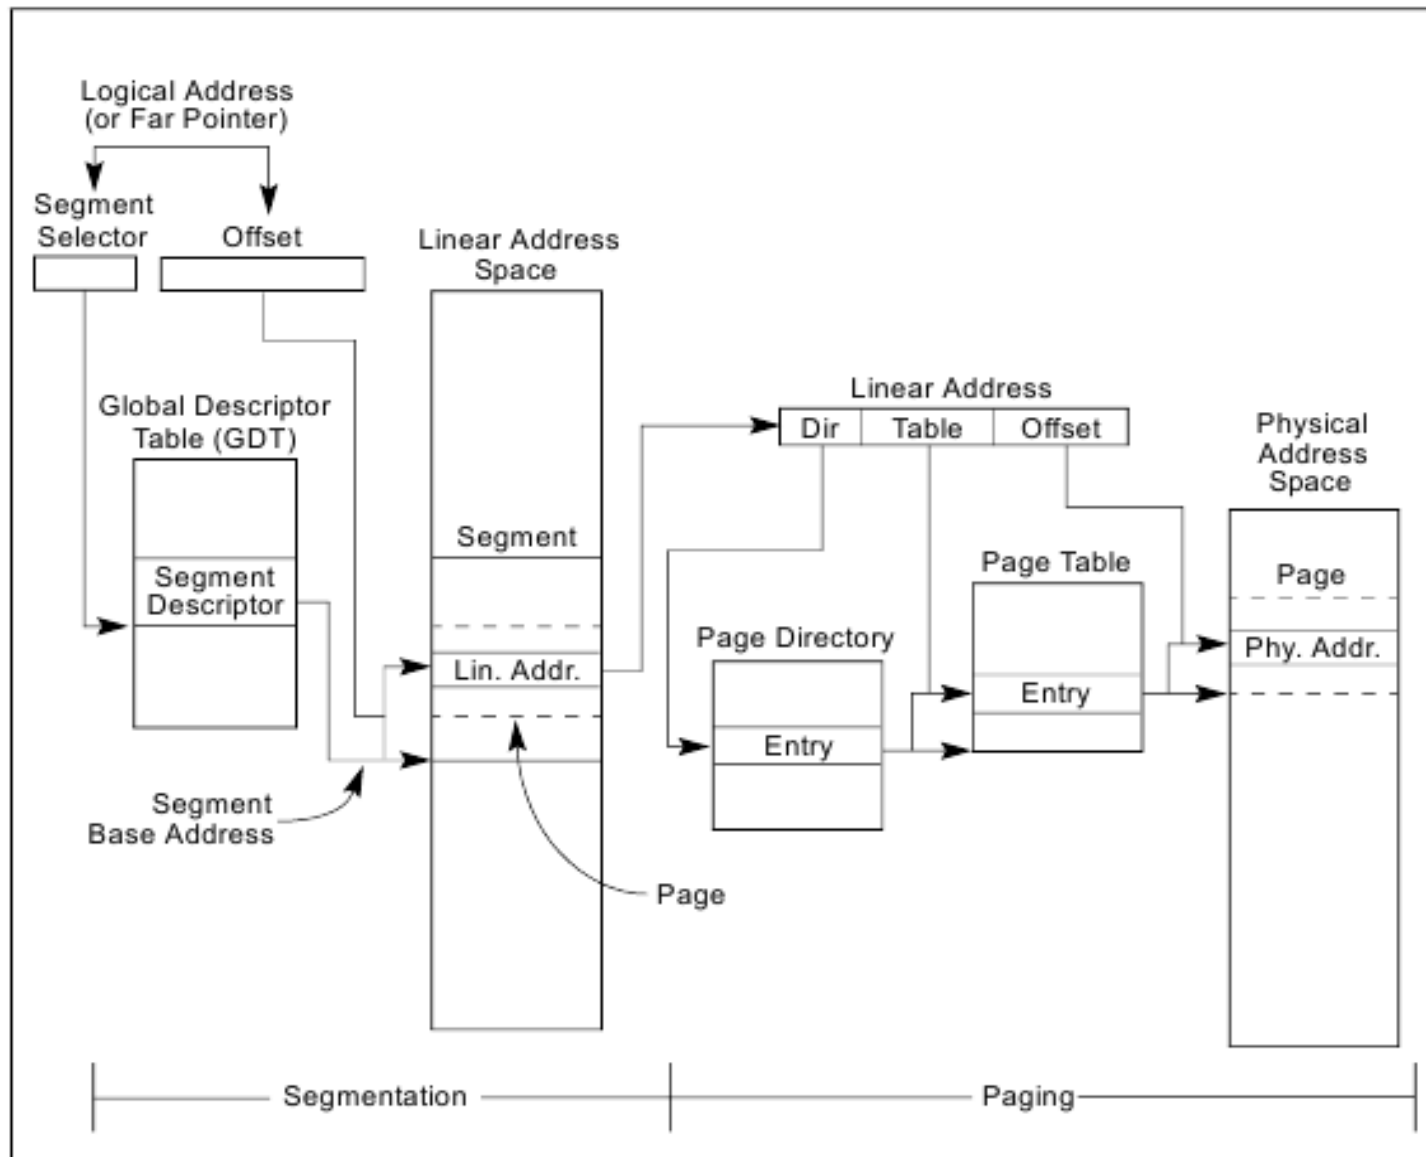


Figure 3-1. Segmentation and Paging

x86 Page Tables

Figure 5-8. Format of a Linear Address

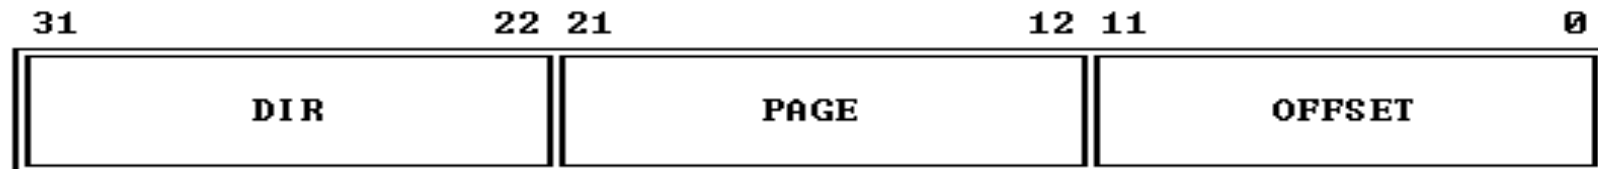
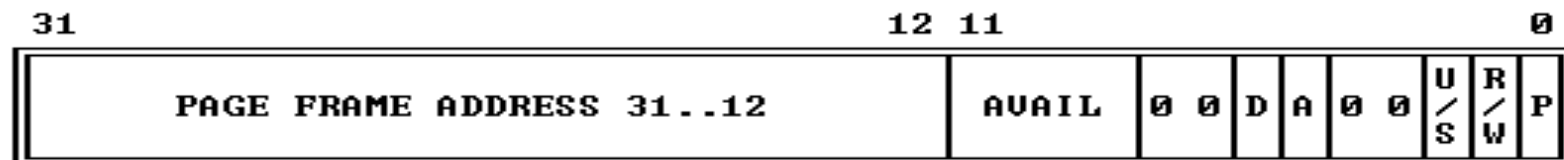


Figure 5-10. Format of a Page Table Entry



P - PRESENT
 R/W - READ/WRITE
 U/S - USER/SUPERVISOR
 D - DIRTY
 AVAIL - AVAILABLE FOR SYSTEMS PROGRAMMER USE
 NOTE: 0 INDICATES INTEL RESERVED. DO NOT DEFINE.

x86 Pate Tables by SETUP

- System Page Directory (SysPD)
 - First level page table
 - Point to 2nd level page tables
- Page Tables
 - Map system code, data, and stack
 - Map all the Physical Memory
- Mapping performed by MMU

PCI Bus

- Scan the PCI bus
 - Look for devices with memory mapped regions
- Fill the IO_Memory_Map in System_Info

```
for(unsigned int i = 0; i < si->iomm_size; i++)  
    si->iomm[i].log_addr = MM::IO_MEM + (si->iomm[i].phy_addr - base);
```

APIC

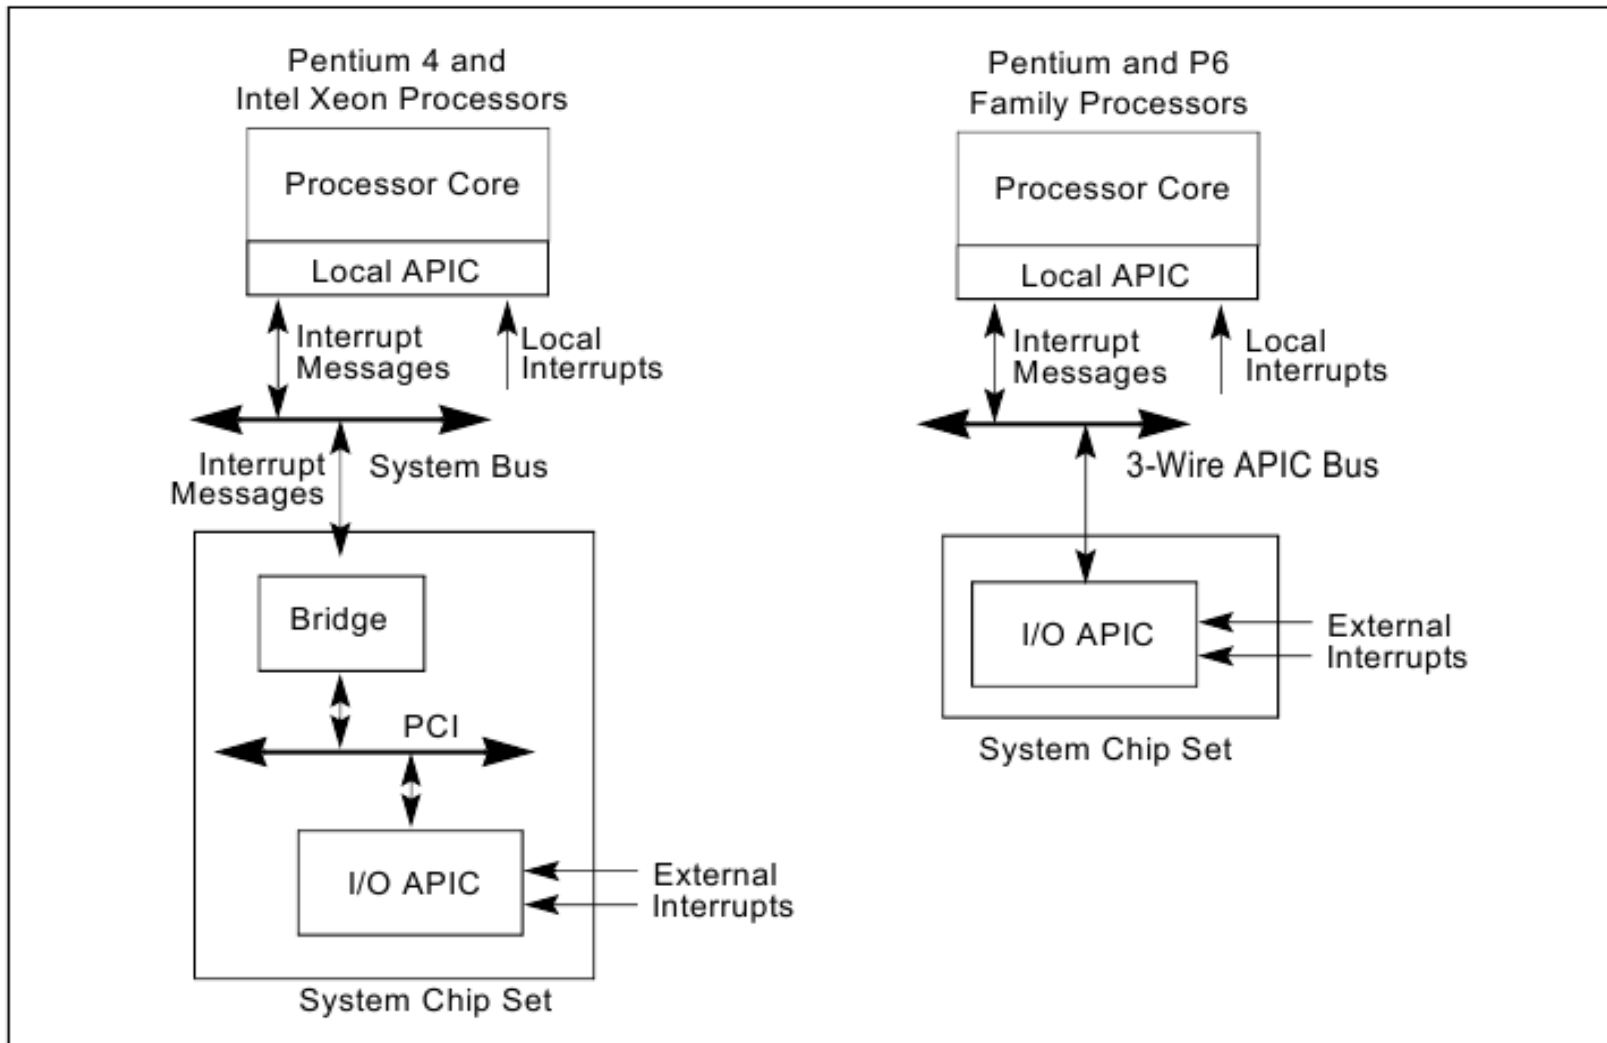


Figure 10-1. Relationship of Local APIC and I/O APIC In Single-Processor Systems

Intel 3A

Local APIC

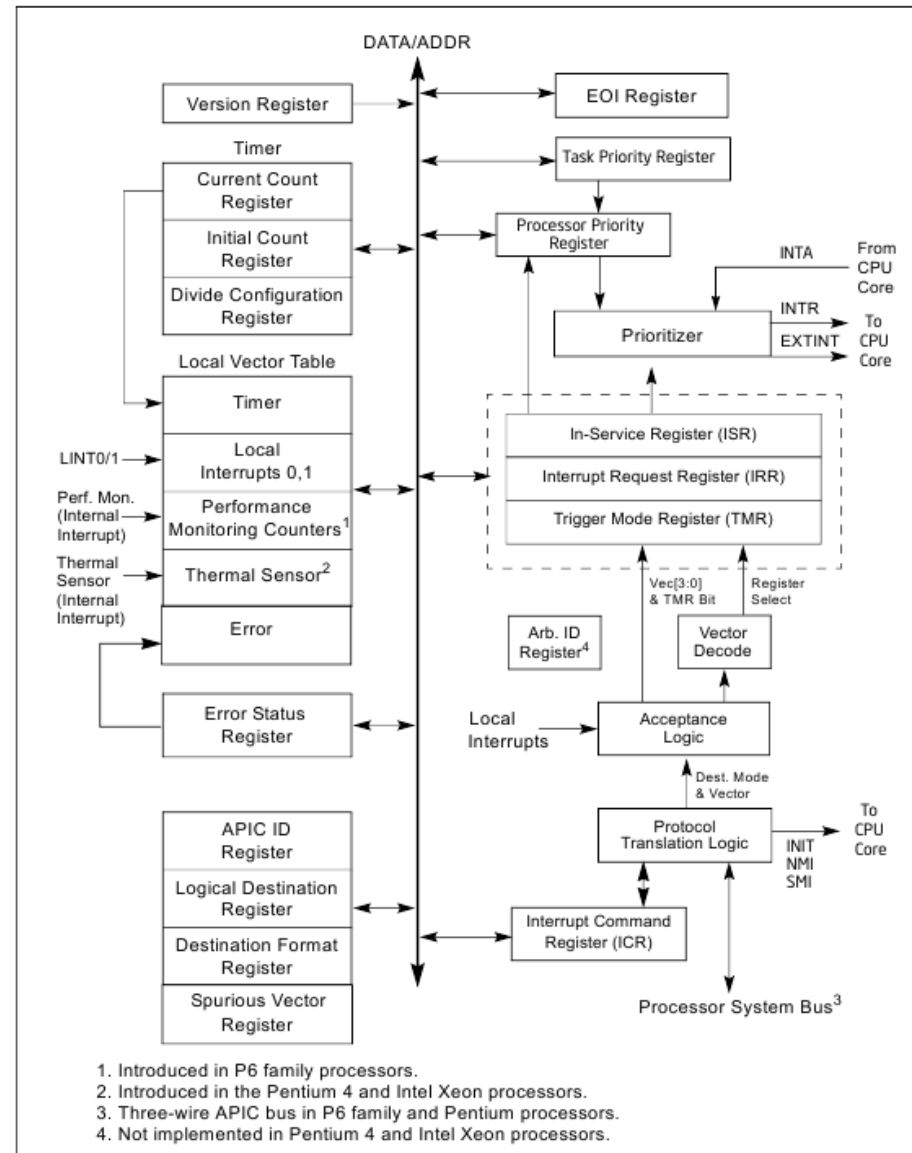


Figure 10-4. Local APIC Structure

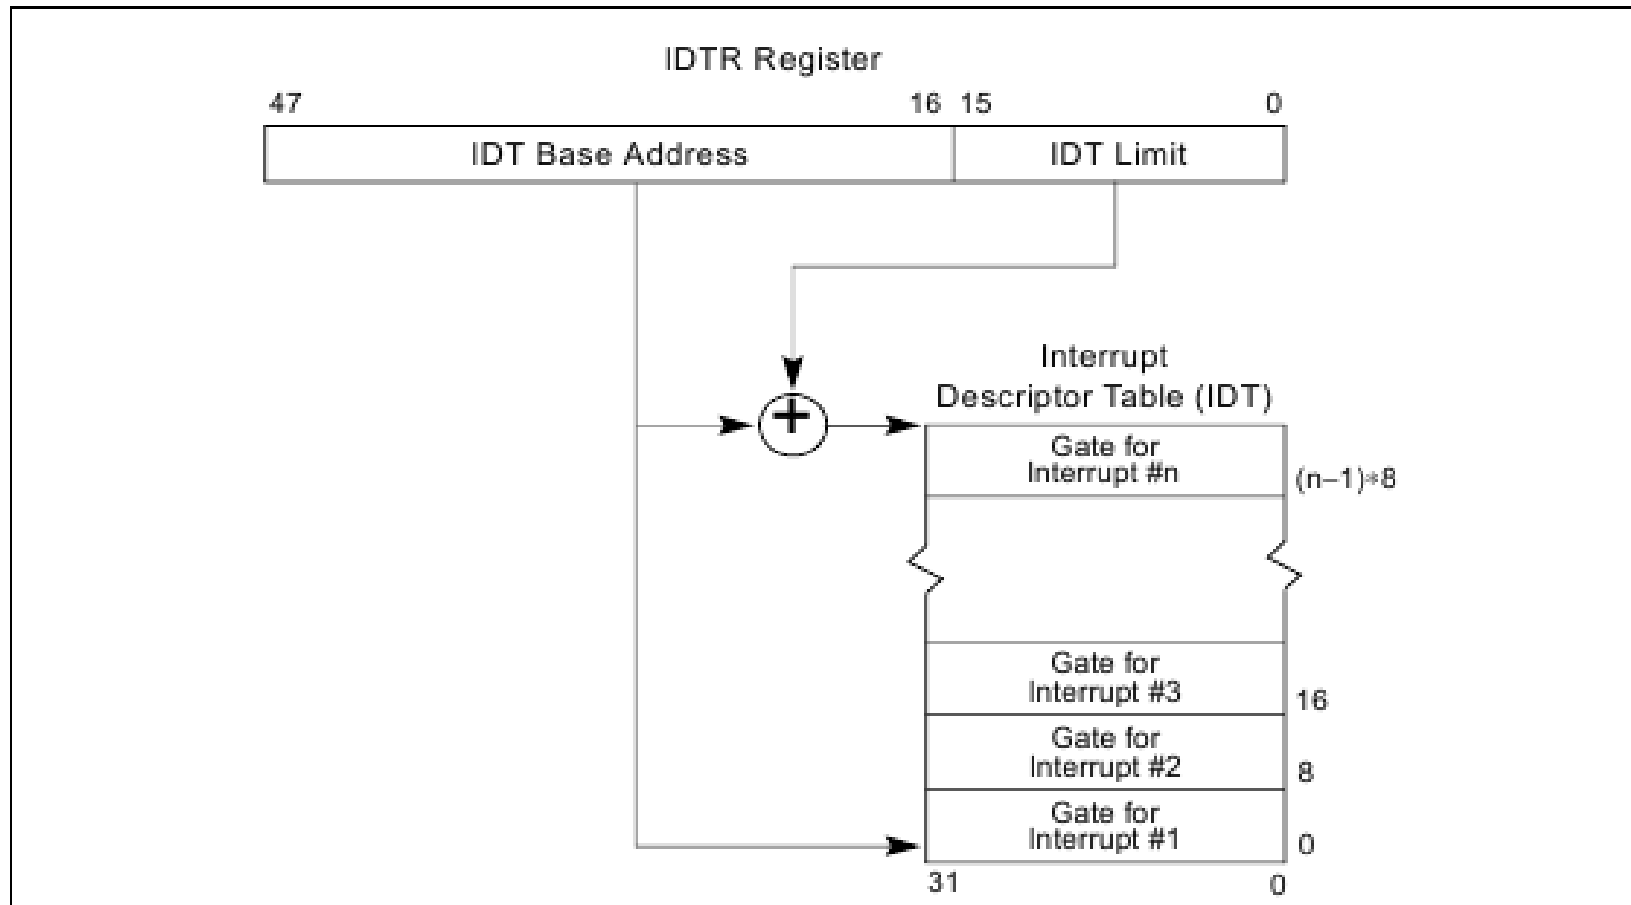


Figure 6-1. Relationship of the IDTR and IDT

Finishing SETUP

- Make previous configuration active
 - Set IDTR
 - Set GDTR
 - Set CR3 (page directory)
 - Enable paging
 - Break pre-fetch queue to activate configuration
 - Reload segment registers (flat model)
 - Remap pointers to their logical addresses
 - Stack
 - System_Info
 - Flush TLB (MMU – make sure new config is in use)
- Load OS and application(s)
- Enable interrupts
- Call system/application