

# Embedded Parallel Operating System

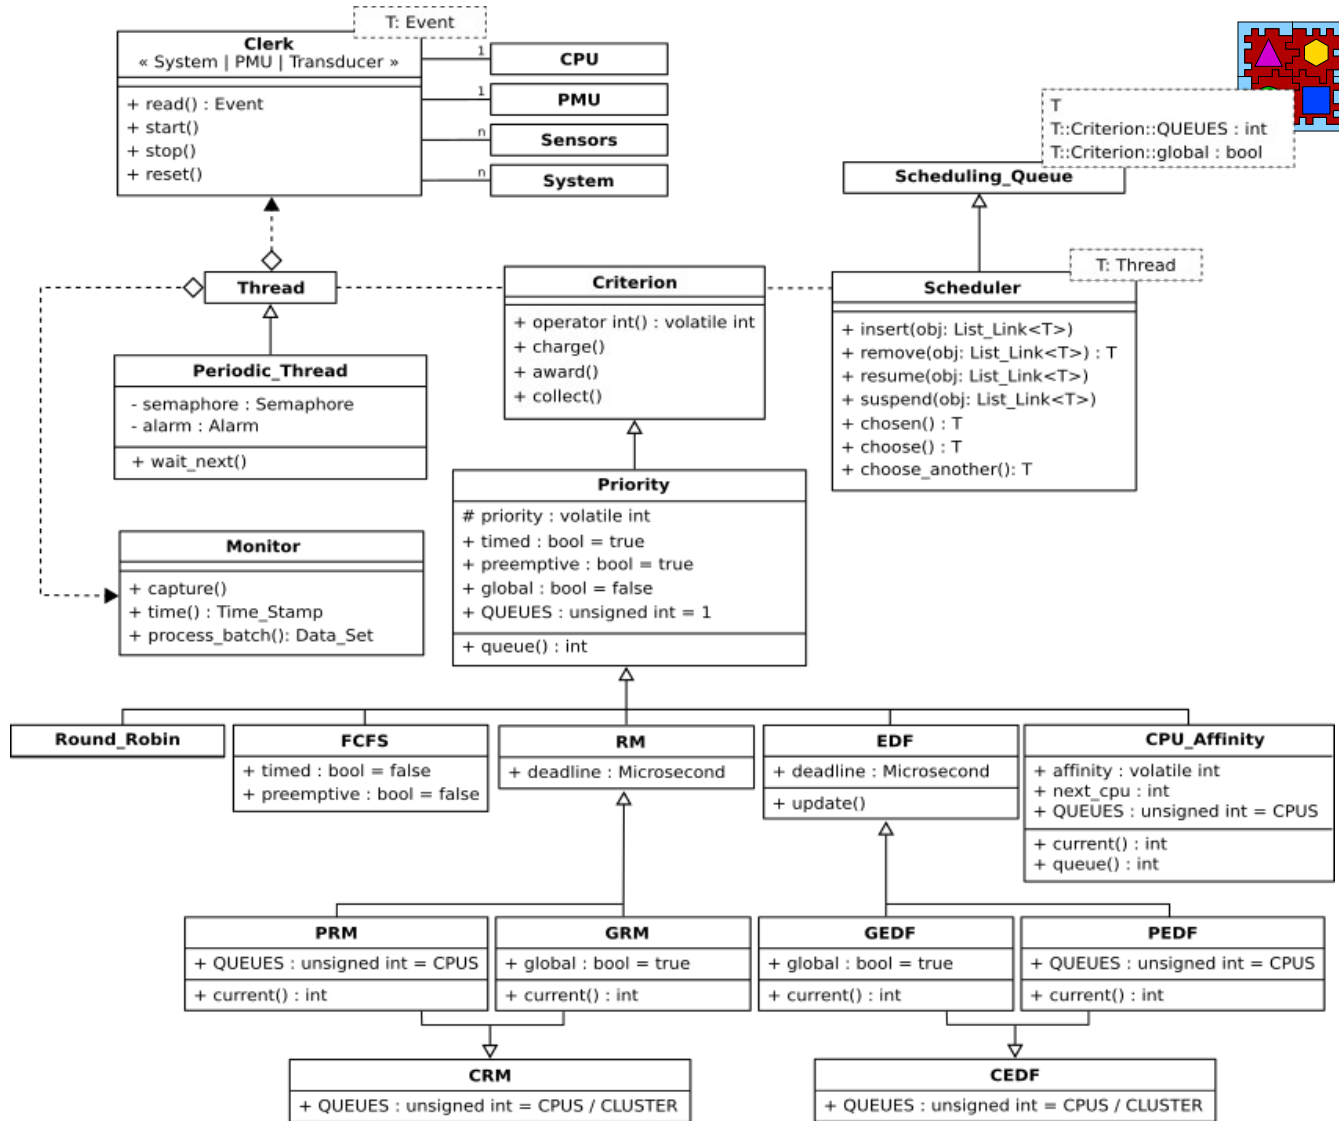
**Coding Journey through OS Design  
–from co-routines to a multicore kernel–**

**Prof. Antônio Augusto Fröhlich, Ph.D.**

**UFSC / LISHA  
July 12, 2021**

## Flexible architecture

- Framework
- Powerful List
- Policies mapped into ordering



```

template<typename T, typename R = Rank>
class Ranked
{
public:
    typedef T Object_Type;
    typedef R Rank_Type;
    typedef Ranked Element;

public:
    Ranked(const T * o, const R & r = 0):
        _object(o), _rank(r) {}

    T * object() const {
        return const_cast<T *>(_object); }

    const R & rank() const { return _rank; }
    const R & key() const { return _rank; }
    void rank(const R & r) { _rank = r; }
    int promote(const R & n = 1) {
        _rank -= n; return _rank; }
    int demote(const R & n = 1) {
        _rank += n; return _rank; }

private:
    const T * _object;
    R _rank;
};

```

```

template<typename T, typename R = Rank>
class Doubly_Linked_Ordered
{
public:
    typedef T Object_Type;
    typedef Rank Rank_Type;
    typedef Doubly_Linked_Ordered Element;

public:
    Doubly_Linked_Ordered(const T * o,
        const R & r = 0):
        _object(o), _rank(r), _prev(0), _next(0) {}


    T * object() const {
        return const_cast<T *>(_object); }

    Element * prev() const { return _prev; }
    Element * next() const { return _next; }
    void prev(Element * e) { _prev = e; }
    void next(Element * e) { _next = e; }

    const R & rank() const { return _rank; }
    void rank(const R & r) { _rank = r; }

private:
    const T * _object;
    R _rank;
    Element * _prev;
    Element * _next;
};

```



```

template<typename T, typename R = Rank>
class Ranked
{
public:
    typedef T Object_Type;
    typedef R Rank_Type;
    typedef Ranked Element;

public:
    Ranked(const T * o, const R & r = 0):
        _object(o), _rank(r) {}

    T * object() const {
        return const_cast<T *>(_object); }

    const R & rank() const { return _rank; }
    const R & key() const { return _rank; }
    void rank(const R & r) { _rank = r; }
    int promote(const R & n = 1) {
        _rank -= n; return _rank; }
    int demote(const R & n = 1) {
        _rank += n; return _rank; }

private:
    const T * _object;
    R _rank;
};

```

```

template<typename T, typename R = Rank>
class
{
public:
    typedef T Object_Type;
    typedef Rank Rank_Type;
    typedef Doubly_Linked_Ordered Element;

public:
    Doubly_Linked_Ordered(const T * o,
        const R & r = 0):
        _object(o), _rank(r), _prev(0), _next(0) {}

    T * object() const {
        return const_cast<T *>(_object); }


    Element * prev() const { return _prev; }
    Element * next() const { return _next; }
    void prev(Element * e) { _prev = e; }
    void next(Element * e) { _next = e; }

    const R & rank() const { return _rank; }
    void rank(const R & r) { _rank = r; }

private:
    const T * _object;
    R _rank;
    Element * _prev;
    Element * _next;
};

```

list order



## list ordering

Rank>

```
{
public:
    typedef T Object_Type;
    typedef R Rank_Type;
    typedef Ranked Element;

public:
    Ranked(const T * o, const R & r = 0):
        _object(o), _rank(r) {}

    T * object() const {
        return const_cast<T *>(_object); }

    const R & rank() const { return _rank; }
    const R & key() const { return _rank; }
    void rank(const R & r) { _rank = r; }
    int promote(const R & n = 1) {
        _rank -= n; return _rank; }
    int demote(const R & n = 1) {
        _rank += n; return _rank; }

private:
    const T * _object;
    R _rank;
};
```

template<typename T, typename R = Rank>

```
class Doubly_Linked_Ordered
{
public:
    typedef T Object_Type;
    typedef Rank Rank_Type;
    typedef Doubly_Linked_Ordered Element;


public:
    Doubly_Linked_Ordered(const T * o,
        const R & r = 0):
        _object(o), _rank(r), _prev(0), _next(0) {}

    T * object() const {
        return const_cast<T *>(_object); }

    Element * prev() const { return _prev; }
    Element * next() const { return _next; }
    void prev(Element * e) { _prev = e; }
    void next(Element * e) { _next = e; }

    const R & rank() const { return _rank; }
    void rank(const R & r) { _rank = r; }

private:
    const T * _object;
    R _rank;
    Element * _prev;
    Element * _next;
};
```



```

class Priority
{
public:
    enum : int {
        MAIN    = 0,
        HIGH    = 1,
        NORMAL   = (unsigned(1)<<(sizeof(int)*8-1))-3,
        LOW     = (unsigned(1)<<(sizeof(int)*8-1))-2,
        IDLE    = (unsigned(1)<<(sizeof(int)*8-1))-1
    };

    enum {
        PERIODIC      = HIGH,
        APERIODIC     = NORMAL,
        SPORADIC      = NORMAL
    };

    static const bool timed = false;
    static const bool dynamic = false;
    static const bool preemptive = true;

public:
    template <typename ... Tn>
    Priority(int p = NORMAL, Tn & ... an):
        _priority(p) {}
    operator const volatile int() const volatile {
        return _priority; }

protected:
    volatile int _priority;
};


```

```

class RR: public Priority
{
public:
    static const bool timed = true;
    static const bool dynamic = false;
    static const bool preemptive = true;

public:
    template <typename ... Tn>
    RR(int p = NORMAL, Tn & ... an): Priority(p) {}
};

```



```

class FCFS: public Priority
{
public:
    static const bool timed = false;
    static const bool dynamic = false;
    static const bool preemptive = false;

public:
    template <typename ... Tn>
    FCFS(int p = NORMAL, Tn & ... an): Priority(p){}
};

```

```

class Priority
{
public:
    enum : int {
        MAIN    = 0,
        HIGH     = 1,
        NORMAL   = (unsigned(1)<<(sizeof(int)*8-1))-3,
        LOW      = (unsigned(1)<<(sizeof(int)*8-1))-2,
        IDLE     = (unsigned(1)<<(sizeof(int)*8-1))-1
    };

    enum {
        PERIODIC      = HIGH,
        APERIODIC      = NORMAL,
        SPORADIC       = NORMAL
    };

    static const bool timed = false;
    static const bool dynamic = false;
    static const bool preemptive = true;

public:
    template <typename ... Tn>
    Priority(int p = NORMAL, Tn & ... an):
        _priority(p) {}
    operator const volatile int() const volatile {
        return _priority; }

protected:
    volatile int _priority;
};

```

```

class
{
public:
    static const bool timed = true;
    static const bool dynamic = false;
    static const bool preemptive = true;

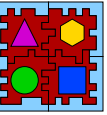
public:
    template <typename ... Tn>
    RR(int p = NORMAL, Tn & ... an): Priority(p) {}
};

class FCFS: public Priority
{
public:
    static const bool timed = false;
    static const bool dynamic = false;
    static const bool preemptive = false;

public:
    template <typename ... Tn>
    FCFS(int p = NORMAL, Tn & ... an): Priority(p){}
};

```

Priority-based policies



```

class Priority
{
public:
    enum : int {
        MAIN    = 0,
        HIGH    = 1,
        NORMAL   = (unsigned(1)<<(sizeof(int)*8-1))-3,
        LOW     = (unsigned(1)<<(sizeof(int)*8-1))-2,
        IDLE    = (unsigned(1)<<(sizeof(int)*8-1))-1
    };


    enum {
        PERIODIC    = HIGH,
        APERIODIC   = NORMAL,
        SPORADIC    = NORMAL
    };

    static const bool timed = false;
    static const bool dynamic = false;
    static const bool preemptive = true;

public:
    template <typename ... Tn>
    Priority(int p = NORMAL, Tn & ... an):
        _priority(p) {}
    operator const volatile int() const volatile {
        return _priority; }

protected:
    volatile int _priority;
};

```



```

class RR: public Priority
{
public:
    static const
    static const bool dynamic = false;
    static const bool preemptive = true;

public:
    template <typename ... Tn>
    RR(int p = NORMAL, Tn & ... an): Priority(p) {}
};

class FCFS: public Priority
{
public:
    static const bool timed = false;
    static const bool dynamic = false;
    static const bool preemptive = false;

public:
    template <typename ... Tn>
    FCFS(int p = NORMAL, Tn & ... an): Priority(p){}
};

```



```

class Priority
{
public:
    enum : int {
        MAIN    = 0,
        HIGH     = 1,
        NORMAL   = (unsigned(1)<<(sizeof(int)*8-1))-3,
        LOW      = (unsigned(1)<<(sizeof(int)*8-1))-2,
        IDLE     = (unsigned(1)<<(sizeof(int)*8-1))-1
    };

    enum {
        PERIODIC      = HIGH,
        APERIODIC     = NORMAL,
        SPORADIC      = NORMAL
    };

    static const bool timed = false;
    static const bool dynamic = false;
    static const bool preemptive = true;

public:
    template <typename ... Tn>
    Priority(int p = NORMAL, Tn & ... an):
        _priority(p) {}
    operator const volatile int() const volatile {
        return _priority; }

protected:
    volatile int _priority;
};


```

```

class RR: public Priority
{
public:
    static const bool timed = true;
    static const bool dynamic = false;
    static const bool preemptive = true;

public:
    template <typename ... Tn>
    RR(int p = NORMAL, Tn & ... an): Priority(p) {}
};

```



policy types

```

class FCFS: public Priority
{
public:
    static const bool timed = false;
    static const bool dynamic = false;
    static const bool preemptive = false;

public:
    template <typename ... Tn>
    FCFS(int p = NORMAL, Tn & ... an): Priority(p){}
};

```

```

class Priority
{
public:
    enum : int {
        MAIN    = 0,
        HIGH     = 1,
        NORMAL   = (unsigned(1)<<(sizeof(int)*8-1))-3,
        LOW      = (unsigned(1)<<(sizeof(int)*8-1))-2,
        IDLE     = (unsigned(1)<<(sizeof(int)*8-1))-1
    };

    enum {
        PERIODIC      = HIGH,
        APERIODIC      = NORMAL,
        SPORADIC       = NORMAL
    };

    static const bool timed = false;
    static const bool dynamic = false;
    static const bool preemptive = true;

public:
    template <typename ... Tn>
    Priority(int p = NORMAL, Tn & ... an):
        _priority(p) {}
    operator const volatile int() const volatile {
        return _priority; }

protected:
    volatile int _priority;
};


```

```

class RR: public Priority
{
public:
    static const bool timed = true;
    static const bool dynamic = false;
    static const bool preemptive = true;

public:
    template <typename ... Tn>
    RR(int p = NORMAL, Tn & ... an): Priority(p) {}
};

```



```

class FCFS:

```

policy traits

```

{
public:
    static const bool timed = false;
    static const bool dynamic = false;
    static const bool preemptive = false;

public:
    template <typename ... Tn>
    FCFS(int p = NORMAL, Tn & ... an): Priority(p){}
};

```

```

class Priority
{
public:
    enum : int {
        MAIN    = 0,
        HIGH     = 1,
        NORMAL   = (unsigned(1)<<(sizeof(int)*8-1))-3,
        LOW      = (unsigned(1)<<(sizeof(int)*8-1))-2,
        IDLE     = (unsigned(1)<<(sizeof(int)*8-1))-1
    };

    enum {
        PERIODIC      = HIGH,
        APERIODIC      = NORMAL,
        SPORADIC       = NORMAL
    };

    static const bool timed = false;
    static const bool dynamic = false;
    static const bool preemptive = true;

public:
    template <typename ... Tn>
    Priority(int p = NORMAL, Tn & ... an):
        _priority(p) {}
    operator const volatile int() const volatile {
        return _priority; }

protected:
    volatile int _priority;
};


```

```

class RR: public Priority
{
public:
    static const bool timed = true;
    static const bool dynamic = false;
    static const bool preemptive = true;

public:
    template <typename ... Tn>
    RR(int p = NORMAL, Tn & ... an): Priority(p) {}
};

```



```

class FCFS: public Priority
{
public:
    static const bool timed = false;
    static const bool dynamic = false;
    static const bool preemptive = false;

public:
    template <typename ... Tn>
    FCFS(int p = NORMAL, Tn & ... an): Priority(p) {}
};

```

scheduling criteria must define operator int() with the semantics of returning the desired order of a given object within the scheduling list (or, alternatively, declare all the ordering operators)

## Round-Robin


```
enum : int {  
    MAIN    = 0,  
    HIGH    = 1,  
    NORMAL  = (unsigned(1)<<(sizeof(int)*8-1))-3,  
    LOW     = (unsigned(1)<<(sizeof(int)*8-1))-2,  
    IDLE    = (unsigned(1)<<(sizeof(int)*8-1))-1  
};  
  
enum {  
    PERIODIC    = HIGH,  
    APERIODIC   = NORMAL,  
    SPORADIC    = NORMAL  
};  
  
static const bool timed = false;  
static const bool dynamic = false;  
static const bool preemptive = true;  
  
public:  
    template <typename ... Tn>  
    Priority(int p = NORMAL, Tn & ... an):  
        _priority(p) {}  
    operator const volatile int() const volatile {  
        return _priority; }  
  
protected:  
    volatile int _priority;  
};
```

```
class RR: public Priority
```

```
{  
public:  
    static const bool timed = true;  
    static const bool dynamic = false;  
    static const bool preemptive = true;  
  
public:  
    template <typename ... Tn>  
    RR(int p = NORMAL, Tn & ... an): Priority(p) {}  
};
```

```
class FCFS: public Priority
```

```
{  
public:  
    static const bool timed = false;  
    static const bool dynamic = false;  
    static const bool preemptive = false;  
  
public:  
    template <typename ... Tn>  
    FCFS(int p = NORMAL, Tn & ... an): Priority(p){}  
};
```



```

class Priority
{
public:
    enum : int {
        MAIN    = 0,
        HIGH    = 1,
        NORMAL   = (unsigned(1)<<(sizeof(int)*8-1))-3,
        LOW     = (unsigned(1)<<(sizeof(int)*8-1))-2,
        IDLE    = (unsigned(1)<<(sizeof(int)*8-1))-1
    };

    enum {
        PERIODIC      = HIGH,
        APERIODIC     = NORMAL,
        SPORADIC      = NORMAL
    };

```

FCFS

```

public:
    template <typename ... Tn>
    Priority(int p = NORMAL, Tn & ... an):
        _priority(p) {}
    operator const volatile int() const volatile {
        return _priority; }

protected:
    volatile int _priority;
};


```

```

class RR: public Priority
{
public:
    static const bool timed = true;
    static const bool dynamic = false;
    static const bool preemptive = true;

public:
    template <typename ... Tn>
    RR(int p = NORMAL, Tn & ... an): Priority(p) {}
};

```



```

class FCFS: public Priority
{
public:
    static const bool timed = false;
    static const bool dynamic = false;
    static const bool preemptive = false;

public:
    template <typename ... Tn>
    FCFS(int p = NORMAL, Tn & ... an): Priority(p){}
};

```

```
template<typename T, typename R = typename T::Criterion>
class Scheduling_Queue: public Scheduling_List<T> {};
```

Scheduling\_Queue is the heart  
of EPOS scheduler

```
template<typename T>
class Scheduler: public Scheduling_Queue<T>
{
private:
    typedef Scheduling_Queue<T> Base;
```

```
public:
    typedef typename T::Criterion Criterion;
    typedef Scheduling_List<T, Criterion> Queue;
    typedef typename Queue::Element Element;
```

```
public:
    Scheduler() {}
```

```
    unsigned int schedulables() { return Base::size(); }
```


```
    T * volatile chosen() { return const_cast<T * volatile>(Base::chosen()->object()); }
```

```
    void insert(T * obj) { Base::insert(obj->link()); }
    T * remove(T * obj) { return Base::remove(obj->link()) ? obj : 0; }
```

```
    void suspend(T * obj) { Base::remove(obj->link()); }
    void resume(T * obj) { Base::insert(obj->link()); }
```

```
    T * choose() { return Base::choose()->object(); }
    T * choose_another() { return Base::choose_another()->object(); }
    T * choose(T * obj) { return Base::choose(obj->link()) ? obj : 0; }
```

```
};
```



```
template<typename T, typename R = typename T::Criterion>
class Scheduling_Queue: public Scheduling_List<T> {};
```

```
template<typename T>
class Scheduler: public Scheduling_Queue<T>
{
private:
    typedef Scheduling_Queue<T> Base;
```

```
public:
    typedef typename T::Criterion Criterion;
    typedef Scheduling_List<T, Criterion> Queue;
    typedef typename Queue::Element Element;
```

Criterion defines ordering

```
public:
    Scheduler() {}
```

```
    unsigned int schedulables() { return Base::size(); }
```


```
    T * volatile chosen() { return const_cast<T * volatile>(Base::chosen()->object()); }
```

```
    void insert(T * obj) { Base::insert(obj->link()); }
    T * remove(T * obj) { return Base::remove(obj->link()) ? obj : 0; }
```

```
    void suspend(T * obj) { Base::remove(obj->link()); }
    void resume(T * obj) { Base::insert(obj->link()); }
```

```
    T * choose() { return Base::choose()->object(); }
    T * choose_another() { return Base::choose_another()->object(); }
    T * choose(T * obj) { return Base::choose(obj->link()) ? obj : 0; }
```

```
};
```



```
template<typename T, typename R = typename T::Criterion>
class Scheduling_Queue: public Scheduling_List<T> {};
```

```
template<typename T>
class Scheduler: public Scheduling_Queue<T>
{
private:
    typedef Scheduling_Queue<T> Base;

public:
    typedef typename T::Criterion Criterion;
    typedef Scheduling_List<T, Criterion> Queue;
    typedef typename Queue::Element Element;
```

```
public:
    Scheduler() {}
```

```
    unsigned int schedulables() { return Base::size(); }
```

```
    T * volatile chosen() { return const_cast<T * volatile>(Base::chosen()->object()); }
```

```
    void insert(T * obj) { Base::insert(obj->link()); }
    T * remove(T * obj) { return Base::remove(obj->link()) ? obj : 0; }
```

```
    void suspend(T * obj) { Base::remove(obj->link()); }
    void resume(T * obj) { Base::insert(obj->link()); }
```


```
    T * choose() { return Base::choose()->object(); }
    T * choose_another() { return Base::choose_another()->object(); }
    T * choose(T * obj) { return Base::choose(obj->link()) ? obj : 0; }
```

```
};
```

define the queue (instead of  
importing it)



```
template<typename T, typename R = typename T::Criterion>
class Scheduling_Queue: public Scheduling_List<T> {};
```



```
template<typename T>
class Scheduler: public Scheduling_Queue<T>
{
private:
    typedef Scheduling_Queue<T> Base;
```

T must provide link()

```
public:
    typedef typename T::Criterion Criterion;
    typedef Scheduling_List<T, Criterion> Queue;
    typedef typename Queue::Element Element;

public:
    Scheduler() {}

    unsigned int schedulables() { return Base::size(); }


    T * volatile chosen() { return const_cast<T * volatile>(Base::chosen()->object()); }

    void insert(T * obj) { Base::insert(obj->link()); }
    T * remove(T * obj) { return Base::remove(obj->link()) ? obj : 0; }

    void suspend(T * obj) { Base::remove(obj->link()); }
    void resume(T * obj) { Base::insert(obj->link()); }

    T * choose() { return Base::choose()->object(); }
    T * choose_another() { return Base::choose_another()->object(); }
    T * choose(T * obj) { return Base::choose(obj->link()) ? obj : 0; }

};
```



```
template<typename T, typename R = typename T::Criterion>
class Scheduling_Queue: public Scheduling_List<T> {};
```

```
template<typename T>
class Scheduler: public Scheduling_Queue<T>
{
private:
    typedef Scheduling_Queue<T> Base;

public:
    typedef typename T::Criterion Criterion;
    typedef Scheduling_List<T, Criterion> Queue;
    typedef typename Queue::Element Element;
```

```
public:
    Scheduler() {}

    unsigned int schedulables() { return Base::size(); }

    T * volatile chosen() { return const_cast<T * volatile>(Base::chosen()->object()); }


    void insert(T * obj) { Base::insert(obj->link()); }
    T * remove(T * obj) { return Base::remove(obj->link()) ? obj : 0; }

    void suspend(T * obj) { Base::remove(obj->link()); }
    void resume(T * obj) { Base::insert(obj->link()); }

    T * choose() { return Base::choose()->object(); }
    T * choose_another() { return Base::choose_another()->object(); }
    T * choose(T * obj) { return Base::choose(obj->link()) ? obj : 0; }

};
```

suspend() and resume() are key for  
hardware implementations



```
template<typename T, typename R = typename T::Criterion>
class Scheduling_Queue: public Scheduling_List<T> {};
```

```
template<typename T>
class Scheduler: public Scheduling_Queue<T>
{
private:
    typedef Scheduling_Queue<T> Base;
```

```
public:
    typedef typename T::Criterion Criterion;
    typedef Scheduling_List<T, Criterion> Queue;
    typedef typename Queue::Element Element;
```

```
public:
    Scheduler() {}
```

```
    unsigned int schedulables() { return Base::size(); }
```

```
    T * volatile chosen() { return const_cast<T * volatile>(Base::chosen()->object()); }
```


```
    void insert(T * obj) { Base::insert(obj->link()); }
    T * remove(T * obj) { return Base::remove(obj->link()) ? obj : 0; }
```

```
    void suspend(T * obj) { Base::remove(obj->link()); }
    void resume(T * obj) { Base::insert(obj->link()); }
```

```
    T * choose() { return Base::choose()->object(); }
    T * choose_another() { return Base::choose_another()->object(); }
    T * choose(T * obj) { return Base::choose(obj->link()) ? obj : 0; }
```

```
};
```

policy based on queue ordering



```
class Thread
{
    friend class Scheduler<Thread>;

protected:
    static const bool preemptive =
        Traits<Thread>::Criterion::preemptive;

public:
    typedef Scheduling_Criteria::Priority Priority;

    typedef Traits<Thread>::Criterion Criterion;
    enum {
        HIGH      = Criterion::HIGH,
        NORMAL    = Criterion::NORMAL,
        LOW       = Criterion::LOW,
        MAIN      = Criterion::MAIN,
        IDLE      = Criterion::IDLE
    };

    typedef Ordered_Queue<Thread, Criterion,
        Scheduler<Thread>::Element> Queue;

protected:
    Criterion & criterion() {
        return const_cast<Criterion &>(_link.rank()); }
    Queue::Element * link() { return &_link; }

    static Thread * volatile running() {
        return _scheduler.chosen(); }
    static void dispatch(Thread * prev, Thread * next,
        bool charge = true);
```

```
protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Thread * volatile _joining;
    Queue::Element _link;

    static volatile unsigned int _thread_count;
    static Scheduler_Timer * _timer;
    static Scheduler<Thread> _scheduler;
};

void Thread::constructor_prologue(...)
{
    lock();
    _thread_count++;
    _scheduler.insert(this);
    _stack = new (SYSTEM) char[stack_size];
}

void Thread::constructor_epilogue(...)
{
    if((_state != READY) && (_state != RUNNING))
        _scheduler.suspend(this);
    if(preemptive && (_state == READY) &&
        (_link.rank() != IDLE))
        reschedule();
    else
        unlock();
}
```

```

class Thread
{
    friend class Scheduler<Thread>;

protected:
    static const bool preemptive =
        Traits<Thread>::Criterion::preemptive;

public:
    typedef Scheduling_Criteria::Priority Priority;

    typedef Traits<Thread>::Criterion Criterion;
    enum {
        HIGH      = Criterion::HIGH,
        NORMAL    = Criterion::NORMAL,
        LOW       = Criterion::LOW,
        MAIN      = Criterion::MAIN,
        IDLE      = Criterion::IDLE
    };

    typedef Ordered_Queue<Thread, Criterion,
        Scheduler<Thread>::Element> Queue;

protected:
    Criterion & criterion() {
        return const_cast<Criterion &>(_link.rank()); }
    Queue::Element * link() { return &_link; }

    static Thread * volatile running() {
        return _scheduler.chosen(); }
    static void dispatch(Thread * prev, Thread * next,
        bool charge = true);

```

protected

```

char * volatile _context;
volatile State _state;
Queue * _waiting;
Thread * volatile _joining;
Queue::Element _link;


static volatile unsigned int _thread_count;
static Scheduler_Timer * _timer;
static Scheduler<Thread> _scheduler;
};


void Thread::constructor_prologue(...)
{
    lock();
    _thread_count++;
    _scheduler.insert(this);
    _stack = new (SYSTEM) char[stack_size];
}

void Thread::constructor_epilogue(...)
{
    if((_state != READY) && (_state != RUNNING))
        _scheduler.suspend(this);
    if(preemptive && (_state == READY) &&
        (_link.rank() != IDLE))
        reschedule();
    else
        unlock();
}

```

isolated scheduler





```
class Thread
{
    friend class Scheduler<Thread>;

protected:
    static const bool preemptive =
        Traits<Thread>::Criterion::preemptive;

public:
    typedef Scheduling_Criteria::Priority Priority;

    typedef Traits<Thread>::Criterion Criterion;
    enum {
        HIGH      = Criterion::HIGH,
        NORMAL    = Criterion::NORMAL,
        LOW       = Criterion::LOW,
        MAIN      = Criterion::MAIN,
        IDLE      = Criterion::IDLE
    };

    typedef Ordered_Queue<Thread, Criterion,
        Scheduler<Thread>::Element> Queue;

protected:
    Criterion & criterion() {
        return const_cast<Criterion &>(_link.rank()); }
    Queue::Element * link() { return &_link; }

    static Thread * volatile running() {
        return _scheduler.chosen(); }
    static void dispatch(Thread * prev, Thread * next,
        bool charge = true);
```


```
protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * waiting;
    Thread * _next;
    Queue * _prev;

    static volatile unsigned int _thread_count;
    static Scheduler_Timer * _timer;
    static Scheduler<Thread> _scheduler;
};

void Thread::constructor_prologue(...)
{
    lock();
    _thread_count++;
    _scheduler.insert(this);
    _stack = new (SYSTEM) char[stack_size];
}

void Thread::constructor_epilogue(...)
{
    if((_state != READY) && (_state != RUNNING))
        _scheduler.suspend(this);
    if(preemptive && (_state == READY) &&
        (_link.rank() != IDLE))
        reschedule();
    else
        unlock();
}
```

Criterion defines priorities and policy  
Traits



```
class Thread
{
    friend class Scheduler<Thread>;

protected:
    static const bool preemptive =
        Traits<Thread>::Criterion::preemptive;

public:
    typedef Scheduling_Criteria::Priority Priority;

    typedef Traits<Thread>::Criterion Criterion;
    enum {
        HIGH      = Criterion::HIGH,
        NORMAL    = Criterion::NORMAL,
        LOW        = Criterion::LOW,
        MAIN       = Criterion::MAIN,
        IDLE       = Criterion::IDLE
    };

    typedef Ordered_Queue<Thread, Criterion,
        Scheduler<Thread>::Element> Queue;

protected:
    Criterion & criterion() {
        return const_cast<Criterion &>(_link.rank()); }
    Queue::Element * link() { return &_link; }

    static Thread * volatile running() {
        return _scheduler.chosen(); }
    static void dispatch(Thread * prev, Thread * next,
        bool charge = true);
```


```
protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Thread * volatile _joining;
    Queue::Element _link;

    static volatile unsigned int _thread_count;
    static Scheduler_Timer * _timer;
    static Scheduler<Thread> _scheduler;
};

void Thread::constructor_prologue(...)
{
    lock()
    _thread
    _sched
    _stack = new (SYSTEM) char[stack_size];
}

void Thread::constructor_epilogue(...)
{
    if((_state != READY) && (_state != RUNNING))
        _scheduler.suspend(this);
    if(preemptive && (_state == READY) &&
        (_link.rank() != IDLE))
        reschedule();
    else
        unlock();
}
```

Queue for uses other than scheduling, with Thread and Criterion, but Element imported for interoperability



```
class Thread
{
    friend class Scheduler<Thread>;

protected:
    static const bool preemptive =
        Traits<Thread>::Criterion::preemptive;

public:
    typedef Scheduling_Criteria::Priority Priority;

    typedef Traits<Thread>::Criterion Criterion;
    enum {
        HIGH      = Criterion::HIGH,
        NORMAL    = Criterion::NORMAL,
        LOW       = Criterion::LOW,
        MAIN      = Criterion::MAIN,
        IDLE      = Criterion::IDLE
    };

    typedef Ordered_Queue<Thread, Criterion,
        Scheduler<Thread>::Element> Queue;

protected:
    Criterion & criterion() {
        return const_cast<Criterion &>(_link.rank()); }
    Queue::Element * link() { return &_link; }

    static Thread * volatile running() {
        return _scheduler.chosen(); }
    static void dispatch(Thread * prev, Thread * next,
        bool charge = true);
```

```
protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Thread * volatile _joining;
    Queue::Element _link;


    static volatile unsigned int _thread_count;
    static Scheduler_Timer * _timer;
    static Scheduler<Thread> _scheduler;
};

void Thread::constructor_prologue(...)
{
    lock();
    _thread_count++;
    _scheduler.insert(this);
    _stack = new (SYSTEM) char[stack_size];
}

void Thread::constructor_epilogue(...)
{
    if(!_scheduler.is_running())
        return;
    if(preemptive &
        (_link.rank() != IDLE))
        reschedule();
    else
        unlock();
}
```

getters for the scheduler





```
class Thread
{
    friend class Scheduler<Thread>;

protected:
    static const bool preemptive =
        Traits<Thread>::Criterion::preemptive;

public:
    typedef Scheduling_Criteria::Priority Priority;

    typedef Traits<Thread>::Criterion Criterion;
    enum {
        HIGH      = Criterion::HIGH,
        NORMAL    = Criterion::NORMAL,
        LOW       = Criterion::LOW,
        MAIN      = Criterion::MAIN,
        IDLE      = Criterion::IDLE
    };

    typedef Ordered_Queue<Thread, Criterion,
        Scheduler<Thread>::Element> Queue;

protected:
    Criterion & criterion() {
        return const_cast<Criterion &>(_link.rank()); }
    Queue::Element * link() { return &_link; }


    static Thread * volatile running() {
        return _scheduler.chosen(); }
    static void dispatch(Thread * prev, Thread * next,
        bool charge = true);
```

```
protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Thread * volatile _joining;
    Queue::Element _link;

    static volatile unsigned int _thread_count;
    static Scheduler_Timer * _timer;
    static Scheduler<Thread> _scheduler;
};

void Thread::constructor_prologue(...)
{
    lock();
    _thread_count++;
    _scheduler.insert(this);
    _stack = new (SYSTEM) char[stack_size];
}

void Thread::running(...)
{
    _running is now the head of the scheduling list
    if((_state != READY) && (_state != RUNNING))
        _scheduler.suspend(this);
    if(preemptive && (_state == READY) &&
        (_link.rank() != IDLE))
        reschedule();
    else
        unlock();
}
```



```
class Thread
{
    friend class Scheduler<Thread>;

protected:
    static const bool preemptive =
        Traits<Thread>::Criterion::preemptive;

public:
    typedef Scheduling_Criteria::Priority Priority;

    typedef Traits<Thread>::Criterion Criterion;
    enum {
        HIGH      = Criterion::HIGH,
        NORMAL     = Criterion::NORMAL,
        LOW        = Criterion::LOW,
        MAIN       = Criterion::MAIN,
        IDLE       = Criterion::IDLE
    };

    typedef Ordered_Queue<Thread, Criterion,
        Scheduler<Thread>::Element> Queue;

protected:
    Criterion & criterion() {
        return const_cast<Criterion &>(_link.rank()); }
    Queue::Element * link() { return &_link; }


    static Thread * volatile running() {
        return _scheduler.chosen(); }
    static void dispatch(Thread * prev, Thread * next,
        bool charge = true);
```

```
protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Thread * volatile _joining;
    Queue::Element _link;

    static volatile unsigned int _thread_count;
    static Scheduler_Timer * _timer;
    static Scheduler<Thread> _scheduler;
};

void Thread::constructor_prologue(...)
{
    lock();
    _thread_count++;
    _scheduler.insert(this);
    _stack = new (SYSTEM) char[stack_size];
}

void Thread::constructor_epilogue(...)
{
    if((_state != READY) && (_state != RUNNING))
        _scheduler.suspend(this);
    if(preemptive && (_state == READY) &&
        (_link.resched() && dispatch() now features charging
        control for pass()
        else
            unlock();
    }
```



```
class Thread
{
    friend class Scheduler<Thread>;

protected:
    static const bool preemptive =
        Traits<Thread>::Criterion::preemptive;

public:
    typedef ... ty Priority;
    Scheduler::schedulables()
    typedef Traits<Thread>::Criterion Criterion;
    enum {
        HIGH      = Criterion::HIGH,
        NORMAL    = Criterion::NORMAL,
        LOW       = Criterion::LOW,
        MAIN      = Criterion::MAIN,
        IDLE      = Criterion::IDLE
    };

    typedef Ordered_Queue<Thread, Criterion,
        Scheduler<Thread>::Element> Queue;

protected:
    Criterion & criterion() {
        return const_cast<Criterion &>(_link.rank()); }
    Queue::Element * link() { return &_link; }


    static Thread * volatile running() {
        return _scheduler.chosen(); }
    static void dispatch(Thread * prev, Thread * next,
        bool charge = true);
```

```
protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Thread * volatile _joining;
    Queue::Element _link;

    static volatile unsigned int _thread_count;
    static Scheduler_Timer * _timer;
    static Scheduler<Thread> _scheduler;
};

void Thread::constructor_prologue(...)
{
    lock();
    _thread_count++;
    _scheduler.insert(this);
    _stack = new (SYSTEM) char[stack_size];
}

void Thread::constructor_epilogue(...)
{
    if((_state != READY) && (_state != RUNNING))
        _scheduler.suspend(this);
    if(preemptive && (_state == READY) &&
        (_link.rank() != IDLE))
        reschedule();
    else
        unlock();
}
```



```
class Thread
{
    friend class Scheduler<Thread>;

protected:
    static const bool preemptive =
        Traits<Thread>::Criterion::preemptive;

public:
    typedef Scheduling_Criteria::Priority Priority;

    typedef Traits<Thread>::Criterion Criterion;
    enum {
        always register the Thread with the
            Scheduler
        MAIN    = Criterion::MAIN,
        IDLE    = Criterion::IDLE
    };

    typedef Ordered_Queue<Thread, Criterion,
        Scheduler<Thread>::Element> Queue;

protected:
    Criterion & criterion() {
        return const_cast<Criterion &>(_link.rank()); }
    Queue::Element * link() { return &_link; }


    static Thread * volatile running() {
        return _scheduler.chosen(); }
    static void dispatch(Thread * prev, Thread * next,
        bool charge = true);
```

```
protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Thread * volatile _joining;
    Queue::Element _link;

    static volatile unsigned int _thread_count;
    static Scheduler_Timer * _timer;
    static Scheduler<Thread> _scheduler;
};

void Thread::constructor_prologue(...)
{
    lock();
    _thread_count++;
    _scheduler.insert(this);
    _stack = new (SYSTEM) char[stack_size];
}

void Thread::constructor_epilogue(...)
{
    if((_state != READY) && (_state != RUNNING))
        _scheduler.suspend(this);
    if(preemptive && (_state == READY) &&
        (_link.rank() != IDLE))
        reschedule();
    else
        unlock();
}
```



```
class Thread
{
    friend class Scheduler<Thread>;

protected:
    static const bool preemptive =
        Traits<Thread>::Criterion::preemptive;

public:
    typedef Scheduling_Criteria::Priority Priority;

    typedef Traits<Thread>::Criterion Criterion;
    enum {
        HIGH      = Criterion::HIGH,
        NORMAL    = Criterion::NORMAL,
        LOW       = Criterion::LOW,
        MAIN      = Criterion::MAIN,
        IDLE      = Criterion::IDLE
    };

    typedef Ordered_Queue<Thread, Criterion,
        suspend the Thread if not READY or RUNNING,
        Element> Queue;

    Criterion & criterion() {
        return const_cast<Criterion &>(_link.rank()); }
    Queue::Element * link() { return &_link; }

    static Thread * volatile running() {
        return _scheduler.chosen(); }
    static void dispatch(Thread * prev, Thread * next,
        bool charge = true);
```

```
protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Thread * volatile _joining;
    Queue::Element _link;


    static volatile unsigned int _thread_count;
    static Scheduler_Timer * _timer;
    static Scheduler<Thread> _scheduler;
};


void Thread::constructor_prologue(...)
{
    lock();
    _thread_count++;
    _scheduler.insert(this);
    _stack = new (SYSTEM) char[stack_size];
}

void Thread::constructor_epilogue(...)
{
    if((_state != READY) && (_state != RUNNING))
        _scheduler.suspend(this);
    if(preemptive && (_state == READY) &&
        (_link.rank() != IDLE))
        reschedule();
    else
        unlock();
}
```

Thread::~Thread()


```
{
    lock();
    assert(_state != RUNNING);
    switch(_state) {
    case RUNNING:
        exit(-1);
        break;
    case READY:
        _scheduler.remove(this);
        _thread_count--;
        break;
    case SUSPENDED:
        _scheduler.resume(this);
        _scheduler.remove(this);
        _thread_count--;
        break;
    case WAITING:
        _waiting->remove(this);
        _scheduler.resume(this);
        _scheduler.remove(this);
        _thread_count--;
        break;
    case FINISHING:
        break;
    }
    if(_joining)
        _joining->resume();
    unlock();
    delete _stack;
}
```





```
Thread::~~Thread()
{
    lock();
    assert(_state != RUNNING);
    switch(_state) {
    case RUNNING:
        exit(-1);
        break;
    case READY:
        _scheduler.remove(this);
        _thread_count--;
        break;
    case SUSPENDED:
        _scheduler.resume(this);
        _scheduler.remove(this);
        _thread_count--;
        break;
    case WAITING:
        _waiting->remove(this);
        _scheduler.resume(this);
        _scheduler.remove(this);
        _thread_count--;
        break;
    case FINISHING:
        break;
    }
    if(_joining)
        _joining->resume();
    unlock();
    delete _stack;
}
```

resume() before remove()  
(won't start running)



```
Thread::~~Thread()
{
    lock();
    assert(_state != RUNNING);
    switch(_state) {
    case RUNNING:
        exit(-1);
        break;
    case READY:
        _scheduler.remove(this);
        _thread_count--;
        break;
    case SUSPENDED:
        _scheduler.resume(this);
        _scheduler.remove(this);
        _thread_count--;
        break;
    case WAITING:
        _waiting->remove(this);
        _scheduler.resume(this);
        _scheduler.remove(this);
        _thread_count--;
        break;
    case FINISHING:
        break;
    }
    if(_joining)
        _joining->resume();
    unlock();
    delete _stack;
}
```


be careful with multiple queues!



```
Thread::~~Thread()
```

```
{  
    lock();  
    assert(_state != RUNNING);  
    switch(_state) {  
    case RUNNING:  
        exit(-1);  
        break;  
    case READY:  
        _scheduler.remove(this);  
        _thread_count--;  
        break;  
    case SUSPENDED:  
        _scheduler.resume(this);  
        _scheduler.remove(this);  
        _thread_count--;  
        break;  
    case WAITING:  
        _waiting->remove(this);  
        _scheduler.resume(this);  
        _scheduler.remove(this);  
        _thread_count--;  
        break;  
    case FINISHING:  
        break;  
    }  
    if(_joining)  
        _joining->resume();  
    unlock();  
    delete _stack;  
}
```

case *RUNNING*: // For switch completion only: the running thread would have deleted itself! Stack wouldn't have been released!



```

void Thread::priority(const Priority & c)
{
    lock();

    _link.rank(Criterion(c));

    if(_state != RUNNING) {
        _scheduler.remove(this);
        _scheduler.insert(this);
    }
    if(preemptive)
        reschedule();
    else
        unlock();
}

void Thread::pass()
{
    lock();

    Thread * prev = running();
    Thread * next = _scheduler.choose(this);

    if(next)
        dispatch(prev, next, false);
    else {
        db<Thread>(WRN) << "Thread::pass => thread
            (" << this << ") not ready!" << endl;
        unlock();
    }
}

```

```

void Thread::suspend(bool locked)
{
    if(!locked)
        lock();
    Thread * prev = running();
    _state = SUSPENDED;
    _scheduler.suspend(this);
    Thread * next = running();
    dispatch(prev, next);
}

void Thread::resume()
{
    lock();
    if(_state == SUSPENDED) {
        _state = READY;
        _scheduler.resume(this);
        if(preemptive)
            reschedule();
        else
            unlock();
    } else {
        db<Thread>(WRN) << "Resume called for
            unsuspended object!" << endl;
        unlock();
    }
}

```

```

void Thread::priority(const Priority & c)
{
    lock();

    _link.rank(Criterion(c));

    if(_state != RUNNING) {
        _scheduler.remove(this);
        _scheduler.insert(this);
    }
    if(preemptive)
        reschedule();
    else
        unlock();
}

void Thread::pass()
{
    lock();

    Thread * prev = running();
    Thread * next = _scheduler.choose(this);

    if(next)
        dispatch(prev, next, false);
    else {
        db<Thread>(WRN) << "Thread::pass => thread
            (" << this << ") not ready!" << endl;
        unlock();
    }
}

```

```

void Thread::suspend(bool locked)
{
    if(!locked)
        mind the conversions!
    Thread * prev = running();
    _state = SUSPENDED;
    _scheduler.suspend(this);
    Thread * next = running();
    dispatch(prev, next);
}

void Thread::resume()
{
    lock();
    if(_state == SUSPENDED) {
        _state = READY;
        _scheduler.resume(this);
        if(preemptive)
            reschedule();
        else
            unlock();
    } else {
        db<Thread>(WRN) << "Resume called for
            unsuspended object!" << endl;
        unlock();
    }
}

```

```

void Thread::priority(const Priority & c)
{
    lock();

    _link.rank(Criterion(c));

    if(_state != RUNNING) {
        _scheduler.remove(this);
        _scheduler.insert(this);
    }
    if(preemptive)
        reschedule();
    else
        unlock();
}

void Thread::pass()
{
    lock();

    Thread * prev = running();
    Thread * next = _scheduler.choose(this);

    if(next)
        dispatch(prev, next, false);
    else {
        db<Thread>(WRN) << "Thread::pass => thread
            (" << this << ") not ready!" << endl;
        unlock();
    }
}

```

```

void Thread::suspend(bool locked)
{
    if(!locked)
        lock();
    Thread * prev = running();
    _state = SUSPENDED;
    _scheduler.remove(this);
    Thread * next = _scheduler.choose(this);
    dispatch(prev, next, false);
}

```

reorder the scheduling queue

```

void Thread::resume()
{
    lock();
    if(_state == SUSPENDED) {
        _state = READY;
        _scheduler.resume(this);
        if(preemptive)
            reschedule();
        else
            unlock();
    } else {
        db<Thread>(WRN) << "Resume called for
            unsuspended object!" << endl;
        unlock();
    }
}

```

```

void Thread::priority(const Priority & c)
{
    lock();

    _link.rank(Criterion(c));

    if(_state != RUNNING) {
        _scheduler.remove(this);
        _scheduler.insert(this);
    }
    if(preemptive)
        reschedule();
    else
        unlock();
}

void Thread::pass()
{
    lock();

    Thread * prev = running();
    Thread * next = _scheduler.choose(this);

    if(next)
        dispatch(prev, next, false);
    else {
        db<Thread>(WRN) << "Thread::pass => thread
            (" << this << ") not ready!" << endl;
        unlock();
    }
}

```

```

void Thread::suspend(bool locked)
{
    if(!locked)
        lock();
    Thread * prev = running();
    _state = SUSPENDED;
    _scheduler.suspend(this);
    Thread * next = running();
    dispatch(prev, next);
}

void Thread::resume()
{
    lock();
    if(_state == SUSPENDED) {
        _state = READY;
        _scheduler.resume(this);
        if(preemptive)
            reschedule();
        else
            null if cannot be chosen
            unlock();
    } else {
        db<Thread>(WRN) << "Resume called for
            unsuspended object!" << endl;
        unlock();
    }
}

```

```

void Thread::priority(const Priority & c)
{
    lock();

    _link.rank(Criterion(c));

    if(_state != RUNNING) {
        _scheduler.remove(this);
        _scheduler.insert(this);
    }
    if(preemptive)
        reschedule();
    else
        unlock();
}

void Thread::pass()
{
    lock();

    Thread * prev = running();
    Thread * next = _scheduler.choose(this);

    if(next)
        dispatch(prev, next, false);
    else {
        db<Thread>(WRN) << "Thread::pass => thread
            (" << this << ") not ready!" << endl;
        unlock();
    }
}

```

```

void Thread::suspend(bool locked)
{
    if(!locked)
        lock();
    Thread * prev = running();
    _state = SUSPENDED;
    _scheduler.suspend(this);
    Thread * next = running();
    dispatch(prev, next);
}

void Thread::resume()
{
    lock();
    if(_state == SUSPENDED) {
        _state = READY;
        _scheduler.resume(this);
        if(preemptive)
            reschedule();
        else
            unlock();
    } else {
        db<Thread>(WRN) << "Resume called for
            unsuspended object!" << endl;
    }
}

```

don't charge the passing Thread  
(charge target)

```
void Thread::priority(const Priority & c)
{
    lock();

```

who holds a pointer to the previously running thread in a software implementation?  
(suspend() == remove())

```

        _scheduler.remove(this);
        _scheduler.insert(this);
    }
    if(preemptive)
        reschedule();
    else
        unlock();
}

void Thread::pass()
{
    lock();

    Thread * prev = running();
    Thread * next = _scheduler.choose(this);

    if(next)
        dispatch(prev, next, false);
    else {
        db<Thread>(WRN) << "Thread::pass => thread
            (" << this << ") not ready!" << endl;
        unlock();
    }
}

```

```

void Thread::suspend(bool locked)
{
    if(!locked)
        lock();
    Thread * prev = running();
    _state = SUSPENDED;
    _scheduler.suspend(this);
    Thread * next = running();
    dispatch(prev, next);
}

void Thread::resume()
{
    lock();
    if(_state == SUSPENDED) {
        _state = READY;
        _scheduler.resume(this);
        if(preemptive)
            reschedule();
        else
            unlock();
    } else {
        db<Thread>(WRN) << "Resume called for
            unsuspended object!" << endl;
        unlock();
    }
}

```

```

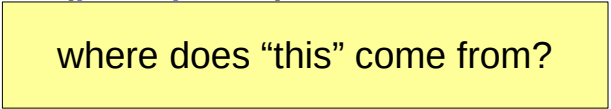
void Thread::priority(const Priority & c)
{
    lock();

    _link.rank(Criterion(c));

    if(_state != RUNNING) {
        _scheduler.remove(this);
        _scheduler.insert(this);
    }
    if(preemptive)
}

```

where does "this" come from?



```

void Thread::pass()
{
    lock();

    Thread * prev = running();
    Thread * next = _scheduler.choose(this);

    if(next)
        dispatch(prev, next, false);
    else {
        db<Thread>(WRN) << "Thread::pass => thread
            (" << this << ") not ready!" << endl;
        unlock();
    }
}

```

```

void Thread::suspend(bool locked)
{
    if(!locked)
        lock();
    Thread * prev = running();
    _state = SUSPENDED;
    _scheduler.suspend(this);
    Thread * next = running();
    dispatch(prev, next);
}


```

```

void Thread::resume()
{
    lock();
    if(_state == SUSPENDED) {
        _state = READY;
        _scheduler.resume(this);
        if(preemptive)
            reschedule();
        else
            unlock();
    } else {
        db<Thread>(WRN) << "Resume called for
            unsuspended object!" << endl;
        unlock();
    }
}

```





```
void Thread::yield()
{
    lock();
    Thread * prev = running();
    Thread * next = __scheduler.choose_another();
    dispatch(prev, next);
}

void Thread::reschedule()
{
    assert(locked());
    Thread * prev = running();
    Thread * next = __scheduler.choose();
    dispatch(prev, next);
}

void Thread::dispatch(Thread * prev, Thread * next,
                      bool charge)
{
    if(charge) {
        if(Criterion::timed)
            __timer->reset();
    }
    if(prev != next) {
        if(prev->_state == RUNNING)
            prev->_state = READY;
        next->_state = RUNNING;
        CPU::switch_context(const_cast<Context **>(&prev->_context), next->_context);
    }
    unlock();
}
```

```
class Init_First
{
public:
    Init_First() {
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::self()->_context->load();
    }
};
```

```

void Thread::yield()
{
    lock();
    Thread * prev = running();
    Thread * next = __scheduler.choose_another();
    dispatch(prev, next);
}

```

```

void Thread::reschedule()
{
    assert(locked());
    Thread * prev = running();
    Thread * next = __scheduler.choose();
    dispatch(prev, next);
}

```

```

void Thread::dispatch(Thread * prev, Thread * next,
                      bool charge)
{
    if(charge) {
        if(Criterion::timed)
            __timer->reset();
    }
    if(prev != next) {
        if(prev->_state == RUNNING)
            prev->_state = READY;
        next->_state = RUNNING;
        CPU::switch_context(const_cast<Context **>(&prev->_context), next->_context);
    }
    unlock();
}

```

```

class Init_First
{
public:

```


```


    Init_First() {
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::self()->_context->load();
    }
};

```

prev might be equal to next!





```
void Thread::yield()
{
    lock();
    Thread * prev = running();
    Thread * next = __scheduler.choose_another();
    dispatch(prev, next);
}
```


```
void Thread::reschedule()
{
    assert(locked());
    Thread * prev = running();
    Thread * next = __scheduler.choose();
    dispatch(prev, next);
}
```

```
void Thread::dispatch(Thread * prev, Thread * next,
                      bool charge)
{
    if(charge) {
        if(Criterion::timed)
            __timer->reset();
    }
    if(prev != next) {
        if(prev->_state == RUNNING)
            prev->_state = READY;
        next->_state = RUNNING;
        CPU::switch_context(const_cast<Context **>(&prev->_context), next->_context);
    }
    unlock();
}
```

```
class Init_First
{
public:
    Init_First() {
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::self()->_context->load();
    }
};
```

more elaborated charges can come here



```
void Thread::yield()
{
    lock();
    Thread * prev = running();
    Thread * next = __scheduler.choose_another();
    dispatch(prev, next);
}


void Thread::reschedule()
{
    assert(locked());
    Thread * prev = running();
    Thread * next = __scheduler.choose();
    dispatch(prev, next);
}

void Thread::dispatch(Thread * prev, Thread * next,
                      bool charge)
{
    if(charge) {
        if(Criterion::timed)
            __timer->reset();
    }
    if(prev != next) {
        if(prev->_state == RUNNING)
            prev->_state = READY;
        next->_state = RUNNING;
        CPU::switch_context(const_cast<Context **>(&prev->_context), next->_context);
    }
    unlock();
}
```

```
class Init_First
{
public:
    Init_First() {
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::self()->_context->load();
    }
};
```

will actually dispatch?



```
void Thread::yield()
{
    lock();
    Thread * prev = running();
    Thread * next = __scheduler.choose_another();
    dispatch(prev, next);
}

void Thread::reschedule()
{
    assert(locked());
    Thread * prev = running();
    Thread * next = __scheduler.choose();
    dispatch(prev, next);
}

void Thread::dispatch(Thread * prev, Thread * next,
                      bool charge)
{
    if(charge) {
        if(Criterion::timed)
            __timer->reset();
    }
    if(prev != next) {
        if(prev->_state == RUNNING)
            prev->_state = READY;
        next->_state = RUNNING;
        CPU::switch_context(const_cast<Context **>(&prev->_context), next->_context);
    }
    unlock();
}
```

```
class Init_First
{
public:
    Init_First() {
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::self()->_context->load();
    }
};
```

then, adjust states

`void Thread::yield()`

CRTs are run before main()

```
Thread * next = __scheduler.choose_another();
dispatch(prev, next);
}
```

`void Thread::reschedule()`

```
{
    assert(locked());
    Thread * prev = running();
    Thread * next = __scheduler.choose();
    dispatch(prev, next);
}
```


`void Thread::dispatch(Thread * prev, Thread * next, bool charge)`


```
{
    if(charge) {
        if(Criterion::timed)
            __timer->reset();
    }
    if(prev != next) {
        if(prev->_state == RUNNING)
            prev->_state = READY;
        next->_state = RUNNING;
        CPU::switch_context(const_cast<Context **>(&prev->_context), next->_context);
    }
    unlock();
}
```

`class Init_First`

```
{
public:
    Init_First() {
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::self()->_context->load();
    }
};
```






```
void Thread::yield()
{
    lock();
    Thread * prev = running();
    Thread * next = __scheduler.choose_another();
    dispatch(prev, next);
}
```

```
void
{
    Thread * prev = running();
    Thread * next = __scheduler.choose();
    dispatch(prev, next);
}
```

is it safe to call self() here?  
(running is set by the Scheduler at each  
insert())

```
void Thread::dispatch(Thread * prev, Thread * next,
                      bool charge)
{
    if(charge) {
        if(Criterion::timed)
            __timer->reset();
    }
    if(prev != next) {
        if(prev->_state == RUNNING)
            prev->_state = READY;
        next->_state = RUNNING;
        CPU::switch_context(const_cast<Context **>(&prev->_context), next->_context);
    }
    unlock();
}
```

```
class Init_First
{
public:
    Init_First() {
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }
        Thread::self()->_context->load();
    };
};
```



```
void Thread::yield()
{
    lock();
    Thread * prev = running();
    Thread * next = _scheduler.choose_another();
    dispatch(prev, next);
}
```

```
void Thread::reschedule()
```

```
{
    assert(locked());
    Thread * prev = running();
    choose();
}
```

no other context to switch from  
(what was then running so far?)

```
void Thread::dispatch(Thread * prev, Thread * next,
                      bool charge)
```

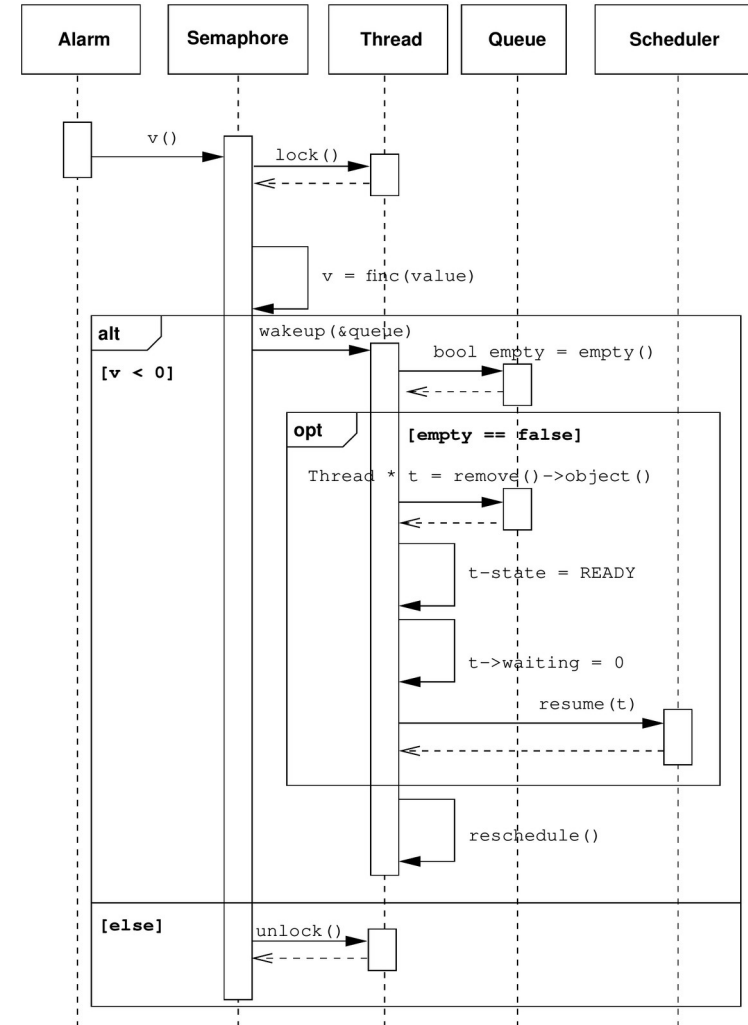
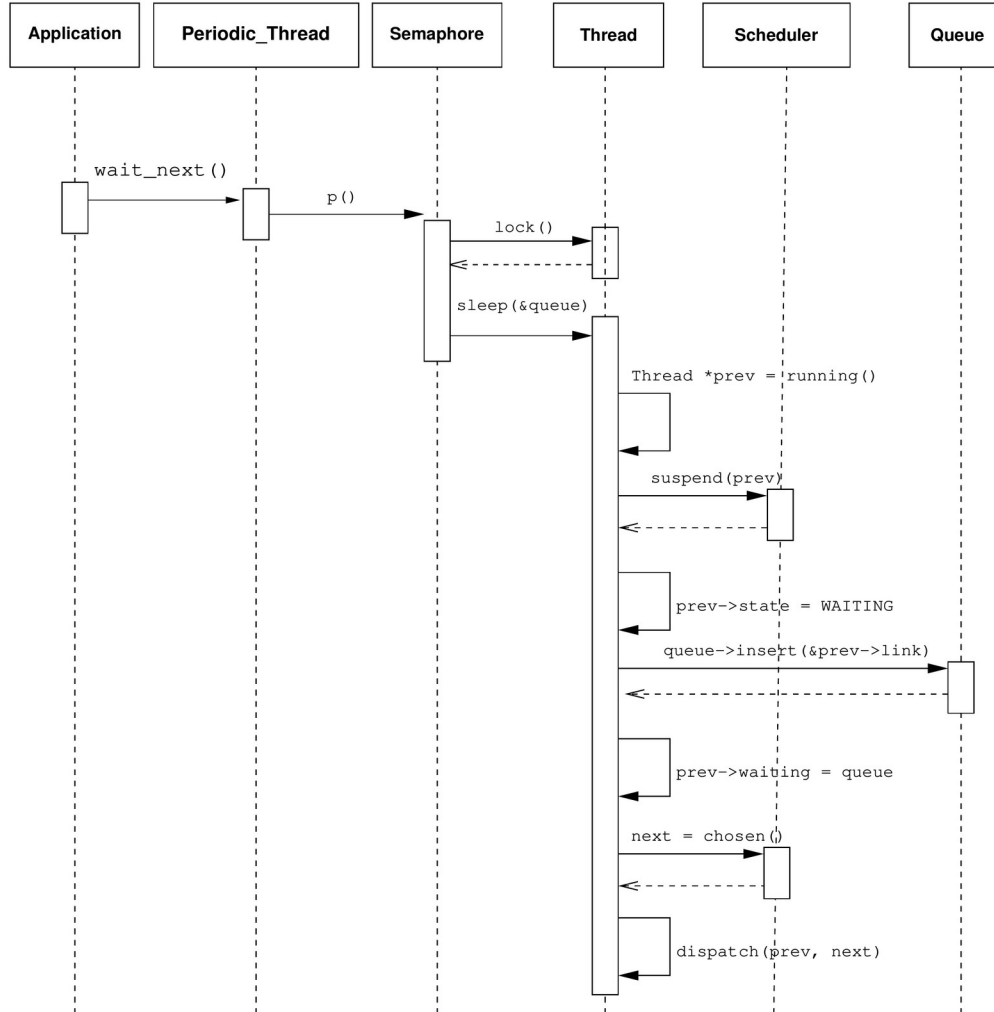

```
{
    if(charge) {
        if(Criterion::timed)
            _timer->reset();
    }
    if(prev != next) {
        if(prev->_state == RUNNING)
            prev->_state = READY;
        next->_state = RUNNING;
        CPU::switch_context(const_cast<Context **>(&prev->_context), next->_context);
    }
    unlock();
}
```


```
class Init_First
```

```
{
public:
    Init_First() {
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }
        Thread::self()->_context->load();
    }
};
```



# sleep() and wakeup()






# Embedded Parallel Operating System

**Coding Journey through OS Design  
–from co-routines to a multicore kernel–**

**Prof. Antônio Augusto Fröhlich, Ph.D.**


**UFSC / LISHA  
September 30, 2020**



# Memory Management

Explore the MMU

- Address spaces
- Segments (not to be confused with segmentation)
- Multiple heaps
- Exported with placement new



```
class Heap: private Grouping_List<char>
{
    static const bool typed = Traits<System>::multiheap;

public:
    using Grouping_List<char>::empty;
    using Grouping_List<char>::size;

    Heap() {}
    Heap(void * addr, unsigned int bytes) { free(addr, bytes); }

    void * alloc(unsigned int bytes) {
        if(!bytes)
            return 0;
        if(!Traits<CPU>::unaligned_memory_access)
            while((bytes % sizeof(void *)))
                ++bytes;
        if(typed)
            bytes += sizeof(void *); // add room for heap pointer
        bytes += sizeof(int); // add room for size
        if(bytes < sizeof(Element))
            bytes = sizeof(Element);
        Element * e = search_decrementing(bytes);
        if(!e) {
            out_of_memory();
            return 0;
        }
        int * addr = reinterpret_cast<int *>(e->object() + e->size());
        if(typed)
            *addr++ = reinterpret_cast<int>(this);
        *addr++ = bytes;
        return addr;
    }
}
```

```

class Heap: private Grouping_List<char>
{
    static const bool typed = Traits<System>::multiheap;


public:
    using Grouping_List<char>::empty;
    using Grouping_List<char>::size;

    Heap() {}
    Heap(void * addr, unsigned int bytes) { free(addr, bytes); }

    void * alloc(unsigned int bytes) {
        if(!bytes)
            return 0;
        if(!Traits<CPU>::unaligned_memory_access)
            while((bytes % sizeof(void *)))
                ++bytes;
        if(typed)
            bytes += sizeof(void *); // add room for heap pointer
        bytes += sizeof(int); // add room for size
        if(bytes < sizeof(Element))
            bytes = sizeof(Element);
        Element * e = search_decrementing(bytes);
        if(!e) {
            out_of_memory();
            return 0;
        }
        int * addr = reinterpret_cast<int *>(e->object() + e->size());
        if(typed)
            *addr++ = reinterpret_cast<int>(this);
        *addr++ = bytes;
        return addr;
    }
}

```

buddy allocator




configurable feature

```
class Heap: private Grouping_List<char>
{
    static const bool typed = Traits<System>::multiheap;

public:
    using Grouping_List<char>::empty;
    using Grouping_List<char>::size;

    Heap() {}
    Heap(void * addr, unsigned int bytes) { free(addr, bytes); }

    void * alloc(unsigned int bytes) {
        if(!bytes)
            return 0;
        if(!Traits<CPU>::unaligned_memory_access)
            while((bytes % sizeof(void *)))
                ++bytes;
        if(typed)
            bytes += sizeof(void *); // add room for heap pointer
        bytes += sizeof(int); // add room for size
        if(bytes < sizeof(Element))
            bytes = sizeof(Element);
        Element * e = search_decrementing(bytes);
        if(!e) {
            out_of_memory();
            return 0;
        }
        int * addr = reinterpret_cast<int *>(e->object() + e->size());
        if(typed)
            *addr++ = reinterpret_cast<int>(this);
        *addr++ = bytes;
        return addr;
    }
}
```




```
class Heap: private Grouping_List<char>
{
    static const bool typed = Traits<System>::multiheap;

public:
    using Grouping_List<char>::empty;
    using Grouping_List<char>::size;

    Heap() {}
    Heap(void * addr, unsigned int bytes) { free(addr, bytes); }

    void * alloc(unsigned int bytes) {
        if(!bytes)
            return 0;
        if(!Traits<CPU>::unaligned_memory_access)
            while((bytes % sizeof(void *)))
                ++bytes;
        if(typed)
            bytes += sizeof(void *); // add room for heap pointer
        bytes += sizeof(int); // add room for size
        if(bytes < sizeof(Element))
            bytes = sizeof(Element);
        Element * e = search_decrementing(bytes);
        if(!e) {
            out_of_memory();
            return 0;
        }
        int * addr = reinterpret_cast<int *>(e->object() + e->size());
        if(typed)
            *addr++ = reinterpret_cast<int>(this);
        *addr++ = bytes;
        return addr;
    }
}
```

heap populated with free()



```
class Heap: private Grouping_List<char>
{
    static const bool typed = Traits<System>::multiheap;


public:
    using Grouping_List<char>::empty;
    using Grouping_List<char>::size;

    Heap() {}
    Heap(void * addr, unsigned int bytes) { free(addr, bytes); }

    void * alloc(unsigned int bytes) {
        if(!bytes)
            return 0;
        if(!Traits<CPU>::unaligned_memory_access)
            while((bytes % sizeof(void *)))
                ++bytes;
        if(typed)
            bytes += sizeof(void *); // add room for heap pointer
        bytes += sizeof(int); // add room for size
        if(bytes < sizeof(Element))
            bytes = sizeof(Element);
        Element * e = search_decrementing(bytes);
        if(!e) {
            out_of_memory();
            return 0;
        }
        int * addr = reinterpret_cast<int *>(e->object() + e->size());
        if(typed)
            *addr++ = reinterpret_cast<int>(this);
        *addr++ = bytes;
        return addr;
    }
}
```

optimization or requirement (e.g. MIPS)






```
class Heap: private Grouping_List<char>
{
    static const bool typed = Traits<System>::multiheap;

public:
    using Grouping_List<char>::empty;
    using Grouping_List<char>::size;

    Heap() {}
    Heap(void * addr, unsigned int bytes) { free(addr, bytes); }

    void * alloc(unsigned int bytes) {
        if(!bytes)
            return 0;
        if(!Traits<CPU>::unaligned_memory_access)
            while((bytes % sizeof(void *)))
                ++bytes;
        if(typed)
            bytes += sizeof(void *); // add room for heap pointer
        bytes += sizeof(int); // add room for size
        if(bytes < sizeof(Element))
            bytes = sizeof(Element);
        Element * e = search_decrementing(bytes);
        if(!e) {
            out_of_memory();
            return 0;
        }
        int * addr = reinterpret_cast<int *>(e->object() + e->size());
        if(typed)
            *addr++ = reinterpret_cast<int>(this);
        *addr++ = bytes;
        return addr;
    }
}
```

metainfo inside allocated unit




```
class Heap: private Grouping_List<char>
{
    static const bool typed = Traits<System>::multiheap;

public:
    using Grouping_List<char>::empty;
    using Grouping_List<char>::size;

    Heap() {}
    Heap(void * addr, unsigned int bytes) { free(addr, bytes); }

    void * alloc(unsigned int bytes) {
        if(!bytes)
            return 0;
        if(!Traits<CPU>::unaligned_memory_access)
            while((bytes % sizeof(void *)))
                ++bytes;
        if(typed)
            bytes += sizeof(void *); // add room for heap pointer
        bytes += sizeof(int); // add room for size
        if(bytes < sizeof(Element))
            bytes = sizeof(Element);
        Element * e = search_decrementing(bytes);
        if(!e) {
            out_of_memory();
            return 0;
        }
        int * addr = reinterpret_cast<int *>(e->object() + e->size());
        if(typed)
            *addr++ = reinterpret_cast<int>(this);
        *addr++ = bytes;
        return addr;
    }
}
```

buddy alloc




```
class Heap: private Grouping_List<char>
{
    static const bool typed = Traits<System>::multiheap;

public:
    using Grouping_List<char>::empty;
    using Grouping_List<char>::size;

    Heap() {}
    Heap(void * addr, unsigned int bytes) { free(addr, bytes); }

    void * alloc(unsigned int bytes) {
        if(!bytes)
            return 0;
        if(!Traits<CPU>::unaligned_memory_access)
            while((bytes % sizeof(void *)))
                ++bytes;
        if(typed)
            bytes += sizeof(void *); // add room for heap pointer
        bytes += sizeof(int); // add room for size
        if(bytes < sizeof(Element))
            bytes = sizeof(Element);
        Element * e = search_decrementing(bytes);
        if(!e) {
            out_of_memory();
            return 0;
        }
        int * addr = reinterpret_cast<int *>(e->object() + e->size());
        if(typed)
            *addr++ = reinterpret_cast<int>(this);
        *addr++ = bytes;
        return addr;
    }
}
```

the famous one!




```
void free(void * ptr, unsigned int bytes) {
    if(ptr && (bytes >= sizeof(Element))) {
        Element * e = new (ptr) Element(reinterpret_cast<char *>(ptr), bytes);
        Element * m1, * m2;
        insert_merging(e, &m1, &m2);
    }
}

static void typed_free(void * ptr) {
    int * addr = reinterpret_cast<int *>(ptr);
    unsigned int bytes = *--addr;
    Heap * heap = reinterpret_cast<Heap *>(*--addr);
    heap->free(addr, bytes);
}

static void untyped_free(Heap * heap, void * ptr) {
    int * addr = reinterpret_cast<int *>(ptr);
    unsigned int bytes = *--addr;
    heap->free(addr, bytes);
}

private:
    void out_of_memory();
};
```




```
void free(void * ptr, unsigned int bytes) {  
    if(ptr && (bytes >= sizeof(Element))) {  
        Element * e = new (ptr) Element(reinterpret_cast<char *>(ptr), bytes);  
        Element * m1, * m2;  
        insert_merging(e, &m1, &m2);  
    }  
}
```

```
static void typed_free(void * ptr) {  
    int * addr = reinterpret_cast<int *>(ptr);  
    unsigned int bytes = *--addr;  
    Heap * heap = reinterpret_cast<Heap *>(*--addr);  
    heap->free(addr, bytes);  
}
```

```
static void untyped_free(Heap * heap, void * ptr) {  
    int * addr = reinterpret_cast<int *>(ptr);  
    unsigned int bytes = *--addr;  
    heap->free(addr, bytes);  
}
```

```
private:  
    void out_of_memory();  
};
```

List::Element in-place  
(size ensured by alloc)




```
void free(void * ptr, unsigned int bytes) {  
    if(ptr && (bytes >= sizeof(Element))) {  
        Element * e = new (ptr) Element(reinterpret_cast<char *>(ptr), bytes);  
        Element * m1, * m2;  
        insert_merging(e, &m1, &m2);  
    }  
}
```

buddy deallocator

```
static void typed_free(void * ptr) {  
    int * addr = reinterpret_cast<int *>(ptr);  
    unsigned int bytes = *--addr;  
    Heap * heap = reinterpret_cast<Heap *>(*--addr);  
    heap->free(addr, bytes);  
}
```

```
static void untyped_free(Heap * heap, void * ptr) {  
    int * addr = reinterpret_cast<int *>(ptr);  
    unsigned int bytes = *--addr;  
    heap->free(addr, bytes);  
}
```

```
private:  
    void out_of_memory();  
};
```




```
void free(void * ptr, unsigned int bytes) {  
    if(ptr && (bytes >= sizeof(Element))) {  
        Element * e = new (ptr) Element(reinterpret_cast<char *>(ptr), bytes);  
        Element * m1, * m2;  
        insert_merging(e, &m1, &m2);  
    }  
}
```

```
static void typed_free(void * ptr) {  
    int * addr = reinterpret_cast<int *>(ptr);  
    unsigned int bytes = *--addr;  
    Heap * heap = reinterpret_cast<Heap *>(*--addr);  
    heap->free(addr, bytes);  
}
```

metainfo inside allocated unit

```
static void untyped_free(Heap * heap, void * ptr) {  
    int * addr = reinterpret_cast<int *>(ptr);  
    unsigned int bytes = *--addr;  
    heap->free(addr, bytes);  
}
```

```
private:  
    void out_of_memory();  
};
```



```
void free(void * ptr, unsigned int bytes) {  
    if(ptr && (bytes >= sizeof(Element))) {  
        Element * e = new (ptr) Element(reinterpret_cast<char *>(ptr), bytes);  
        Element * m1, * m2;  
        insert_merging(e, &m1, &m2);  
    }  
}
```


```
static void typed_free(void * ptr) {  
    int * addr = reinterpret_cast<int *>(ptr);  
    unsigned int bytes = *--addr;  
    Heap * heap = reinterpret_cast<Heap *>(*--addr);  
    heap->free(addr, bytes);  
}
```

```
static void untyped_free(Heap * heap, void * ptr) {  
    int * addr = reinterpret_cast<int *>(ptr);  
    unsigned int bytes = *--addr;  
    heap->free(addr, bytes);  
}
```

```
private:  
    void out_of_memory();  
};
```

typed and untyped free()





```
class Application
{
    friend void * ::malloc(size_t);
    friend void ::free(void *);
    :
private:
    static char _preheap[sizeof(Heap)];
    static Heap * _heap;
};

class System
{
    friend void * ::malloc(size_t);
    friend void ::free(void *);
    friend void * ::operator new(size_t, const EPOS::System_Allocator &);
    friend void * ::operator new[](size_t, const EPOS::System_Allocator &);
    friend void ::operator delete(void *);
    friend void ::operator delete[](void *);
    :
private:
    static char _preheap[(Traits<System>::multiheap ? sizeof(Segment) : 0) + sizeof(Heap)];
    static Segment * _heap_segment;
    static Heap * _heap;
};
```


```
class Application
```

```
{  
    friend void * ::malloc(size_t);  
    friend void ::free(void *);  
:  
private:  
    static char _preheap[sizeof(Heap)];  
    static Heap * _heap;  
};
```

malloc() and free() use the application's heap

```
class System
```


```
{  
    friend void * ::malloc(size_t);  
    friend void ::free(void *);  
    friend void * ::operator new(size_t, const EPOS::System_Allocator &);  
    friend void * ::operator new[](size_t, const EPOS::System_Allocator &);  
    friend void ::operator delete(void *);  
    friend void ::operator delete[](void *);  
:  
private:  
    static char _preheap[(Traits<System>::multiheap ? sizeof(Segment) : 0) + sizeof(Heap)];  
    static Segment * _heap_segment;  
    static Heap * _heap;  
};
```



```
class Application
{
    friend void * ::malloc(size_t);
    friend void ::free(void *);
:
private:
    static char _preheap[sizeof(Heap)];
    static Heap * _heap;
};
```

proto-heap  
(cannot new Heap)  
(in data segment)


```
class System
{
    friend void * ::malloc(size_t);
    friend void ::free(void *);
    friend void * ::operator new(size_t, const EPOS::System_Allocator &);
    friend void * ::operator new[](size_t, const EPOS::System_Allocator &);
    friend void ::operator delete(void *);
    friend void ::operator delete[](void *);
:
private:
    static char _preheap[(Traits<System>::multiheap ? sizeof(Segment) : 0) + sizeof(Heap)];
    static Segment * _heap_segment;
    static Heap * _heap;
};
```



```
class Application
{
    friend void * ::malloc(size_t);
    friend void ::free(void *);
    :
private:
    static char _preheap[sizeof(Heap)];
    static Heap * _heap;
};
```

```
class System
{
    friend void * ::malloc(size_t);
    friend void ::free(void *);
    friend void * ::operator new(size_t, const EPOS::System_Allocator &);
    friend void * ::operator new[](size_t, const EPOS::System_Allocator &);
    friend void ::operator delete(void *);
    friend void ::operator delete[](void *);
    :
private:
    static char _preheap[(Traits<System>::multiheap ? sizeof(Segment) : 0) + sizeof(Heap)];
    static Segment * _heap_segment;
    static Heap * _heap;
};
```

system allocator based on placement  
new operator

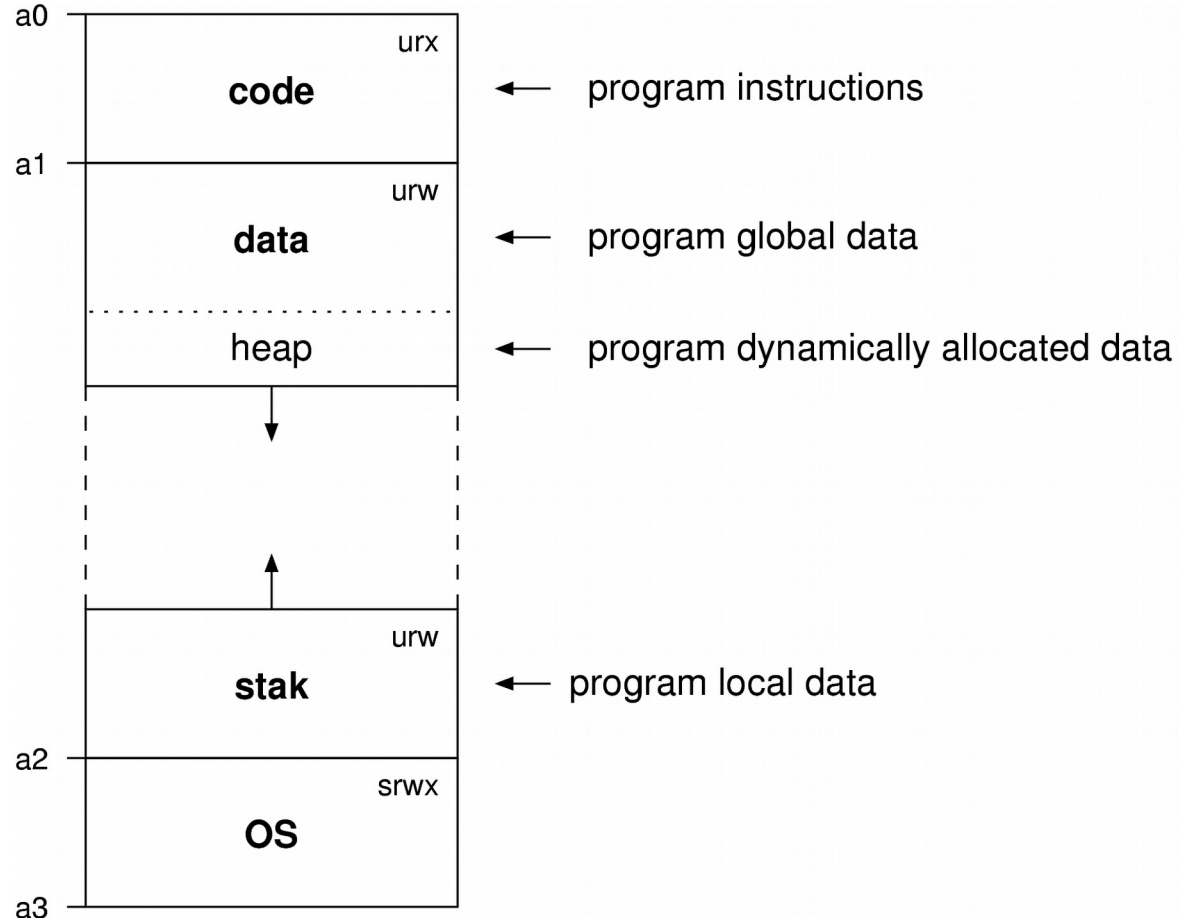




```
class Application
{
    friend void * ::malloc(size_t);
    friend void ::free(void *);
    :
private:
    static char _preheap[sizeof(Heap)];
    static Heap * _heap;
};
```

```
class System
{
    friend void * ::malloc(size_t);
    friend void ::free(void *);
    friend void * ::operator new(size_t, const EPOS::System_Allocator &);
    friend void * ::operator new[](size_t, const EPOS::System_Allocator &);
    friend void ::operator delete(void *);
    friend void ::operator delete[](void *);
    :
private:
    static char _preheap[(Traits<System>::multiheap ? sizeof(Segment) : 0) + sizeof(Heap)];
    static Segment * _heap_segment;
    static Heap * _heap;
};
```

proto-heap  
(heap in separate segment)

# Unix Memory Model






```
class Init_System
{
private:
    static const unsigned int HEAP_SIZE = Traits<System>::HEAP_SIZE;

public:
    Init_System() {
        :
        if(Traits<System>::multiheap) {
            Segment * tmp = reinterpret_cast<Segment *>(&System::_preheap[0]);
            System::_heap_segment = new (tmp) Segment(HEAP_SIZE, WHITE, Segment::Flags::SYS);
            System::_heap = new (&System::_preheap[sizeof(Segment)])
                Heap(Address_Space(MMU::current()).attach(System::_heap_segment, Memory_Map::SYS_HEAP),
                    System::_heap_segment->size());
        } else
            System::_heap = new (&System::_preheap[0]) Heap(MMU::alloc(MMU::pages(HEAP_SIZE)), HEAP_SIZE);
        :
    }
};

// SETUP
:
    if(Traits<System>::multiheap) { // Application heap in data segment
        si->lm.app_data_size = MMU::align_page(si->lm.app_data_size);
        si->lm.app_stack = si->lm.app_data + si->lm.app_data_size;
        si->lm.app_data_size += MMU::align_page(Traits<Application>::STACK_SIZE);
        si->lm.app_heap = si->lm.app_data + si->lm.app_data_size;
        si->lm.app_data_size += MMU::align_page(Traits<Application>::HEAP_SIZE);
    }
:
}
```



```
class Init_System
{
private:
    static const unsigned int HEAP_SIZE = Traits<System>::HEAP_SIZE;


public:
    Init_System() {
        :
        if(Traits<System>::multiheap) {
            Segment * tmp = reinterpret_cast<Segment *>(&System::_preheap[0]);
            System::_heap_segment = new (tmp) Segment(HEAP_SIZE, WHITE, Segment::Flags::SYS);
            System::_heap = new (&System::_preheap[sizeof(Segment)])
                Heap(Address_Space(MMU::current()).attach(System::_heap_segment, Memory_Map::SYS_HEAP),
                    System::_heap_segment->size());
        } else
            System::_heap = new (&System::_preheap[0]) Heap(MMU::alloc(MMU::pages(HEAP_SIZE)), HEAP_SIZE);
        :
    }
};

// SETUP
:
    if(Traits<System>::multiheap) { // Application heap in data segment
        si->lm.app_data_size = MMU::align_page(si->lm.app_data_size);
        si->lm.app_stack = si->lm.app_data + si->lm.app_data_size;
        si->lm.app_data_size += MMU::align_page(Traits<Application>::STACK_SIZE);
        si->lm.app_heap = si->lm.app_data + si->lm.app_data_size;
        si->lm.app_data_size += MMU::align_page(Traits<Application>::HEAP_SIZE);
    }
:

```

system's heap in a private segment






```
class Init_System
{
private:
    static const unsigned int HEAP_SIZE = Traits<System>::HEAP_SIZE;

public:
    Init_System() {
        :
        if(Traits<System>::multiheap) {
            Segment * tmp = reinterpret_cast<Segment *>(&System::_preheap);
            System::_heap_segment = new (tmp) Segment(HEAP_SIZE, WHITE, Segment::Flags::SYS);
            System::_heap = new (&System::_preheap[sizeof(Segment)])
                Heap(Address_Space(MMU::current()).attach(System::_heap_segment, Memory_Map::SYS_HEAP),
                    System::_heap_segment->size());
        } else
            System::_heap = new (&System::_preheap[0]) Heap(MMU::alloc(MMU::pages(HEAP_SIZE)), HEAP_SIZE);
        :
    }
};

// SETUP
:
    if(Traits<System>::multiheap) { // Application heap in data segment
        si->lm.app_data_size = MMU::align_page(si->lm.app_data_size);
        si->lm.app_stack = si->lm.app_data + si->lm.app_data_size;
        si->lm.app_data_size += MMU::align_page(Traits<Application>::STACK_SIZE);
        si->lm.app_heap = si->lm.app_data + si->lm.app_data_size;
        si->lm.app_data_size += MMU::align_page(Traits<Application>::HEAP_SIZE);
    }
:

```

remember: a Segment must be attached to an Address\_Space!




```
class Init_System
{
private:
    static const unsigned int HEAP_SIZE = Traits<System>::HEAP_SIZE;

public:
    Init_System() {
        :
        if(Traits<System>::multiheap) {
            Segment * tmp = reinterpret_cast<Segment *>(&System::_preheap[0]);
            System::_heap_segment = new (tmp) Segment(HEAP_SIZE, WHITE, Segment::Flags::SYS);
            System::_heap = new (&System::_preheap[sizeof(Segment)])
                Heap(Address_Space(MMU::current()).attach(System::_heap_segment, Memory_Map::SYS_HEAP),
                    System::_heap_segment->size());
        } else
            System::_heap = new (&System::_preheap[0]) Heap(MMU::alloc(MMU::pages(HEAP_SIZE)), HEAP_SIZE);
        :
    }
};

// SETUP
:
    if(Traits<System>::multiheap) { // Application heap in data segment
        si->lm.app_data_size = MMU::align_page(si->lm.app_data_size);
        si->lm.app_stack = si->lm.app_data + si->lm.app_data_size;
        si->lm.app_data_size += MMU::align_page(Traits<Application>::STACK_SIZE);
        si->lm.app_heap = si->lm.app_data + si->lm.app_data_size;
        si->lm.app_data_size += MMU::align_page(Traits<Application>::HEAP_SIZE);
    }
    :
```

get a reference to the current address  
space (arranged by SETUP)




```
class Init_System
{
private:
    static const unsigned int HEAP_SIZE = Traits<System>::HEAP_SIZE;

public:
    Init_System() {
        :
        if(Traits<System>::multiheap) {
            Segment * tmp = reinterpret_cast<Segment *>(&System::_preheap[0]);
            System::_heap_segment = new (tmp) Segment(HEAP_SIZE, WHITE, Segment::Flags::SYS);
            System::_heap = new (&System::_preheap[sizeof(Segment)])
                Heap(Address_Space(MMU::current()).attach(System::_heap_segment, Memory_Map::SYS_HEAP),
                    System::_heap_segment->size());
        } else
            System::_heap = new (&System::_preheap[0]) Heap(MMU::al);
        :
    }
};

// SETUP
:
    if(Traits<System>::multiheap) { // Application heap in data segment
        si->lm.app_data_size = MMU::align_page(si->lm.app_data_size);
        si->lm.app_stack = si->lm.app_data + si->lm.app_data_size;
        si->lm.app_data_size += MMU::align_page(Traits<Application>::STACK_SIZE);
        si->lm.app_heap = si->lm.app_data + si->lm.app_data_size;
        si->lm.app_data_size += MMU::align_page(Traits<Application>::HEAP_SIZE);
    }
    :
```

populate the heap with the Segment's memory




```
class Init_System
{
private:
    static const unsigned int HEAP_SIZE = Traits<System>::HEAP_SIZE;

public:
    Init_System() {
        :
        if(Traits<System>::multiheap) {
            Segment * tmp = reinterpret_cast<Segment *>(&System::_preheap[0]);
            System::_heap_segment = new (tmp) Segment(HEAP_SIZE, WHITE, Segment::Flags::SYS);
            System::_heap = new (&System::_preheap[sizeof(Segment)])
                Heap(Address_Space(MMU::current()).attach(System::_heap_segment, Memory_Map::SYS_HEAP),
                    System::_heap_segment->size());
        } else
            System::_heap = new (&System::_preheap[0]) Heap(MMU::alloc(MMU::pages(HEAP_SIZE)), HEAP_SIZE);
        :
    }
};

// SETUP
:
if(Traits<System>::multiheap) { // Application heap in data segment
    si->lm.app_data_size = MMU::align_page(si->lm.app_data_size);
    si->lm.app_stack = si->lm.app_data + si->lm.app_data_size;
    si->lm.app_data_size += MMU::align_page(Traits<Application>::STACK_SIZE);
    si->lm.app_heap = si->lm.app_data + si->lm.app_data_size;
    si->lm.app_data_size += MMU::align_page(Traits<Application>::HEAP_SIZE);
}
:
```

another nice usage for the placement  
new, isn't it?



```
class Init_System
{
private:
    static const unsigned int HEAP_SIZE = Traits<System>::HEAP_SIZE;

public:
    Init_System() {
        :
        if(Traits<System>::multiheap) {
            Segment * tmp = reinterpret_cast<Segment *>(&System::_preheap[0]);
            System::_heap_segment = new (tmp) Segment(HEAP_SIZE, WHITE, Segment::Flags::SYS);
            System::_heap = new (&System::_preheap[sizeof(Segment)])
                Heap(Address_Space(MMU::current()).attach(System::_heap_segment, Memory_Map::SYS_HEAP),
                    System::_heap_segment->size());
        } else
            System::_heap = new (&System::_preheap[0]) Heap(MMU::alloc(MMU::pages(HEAP_SIZE)), HEAP_SIZE);
        :
    }
};

// SETUP
:
if(Traits<System>::multiheap) { // Application heap in data segment
    si->lm.app_data_size = MMU::align_page(si->lm.app_data_size);
    si->lm.app_stack = si->lm.app_data + si->lm.app_data_size;
    si->lm.app_data_size += MMU::align_page(Traits<Application>::STACK_SIZE);
    si->lm.app_heap = si->lm.app_data + si->lm.app_data_size;
    si->lm.app_data_size += MMU::align_page(Traits<Application>::HEAP_SIZE);
}
:
```

SETUP arranges a memory model for single-task configs

```
extern "C"
```

```
{
```

```
    // Standard C Library allocators
```

```
    inline void * malloc(size_t bytes) {
```

```
        __USING_SYS;
```

```
        if(Traits<System>::multiheap)
```

```
            return Application::_heap->alloc(bytes);
```

```
        else
```

```
            return System::_heap->alloc(bytes);
```

```
    }
```

```
    inline void * calloc(size_t n, unsigned int bytes) {
```

```
        void * ptr = malloc(n * bytes);
```

```
        memset(ptr, 0, n * bytes);
```

```
        return ptr;
```

```
    }
```

```
    inline void free(void * ptr) {
```

```
        __USING_SYS;
```

```
        if(Traits<System>::multiheap)
```


```
            Heap::typed_free(ptr);
```


```
        else
```

```
            Heap::untyped_free(System::_heap, ptr);
```

```
    }
```

```
}
```






```
extern "C"
{
    // Standard C Library allocators
    inline void * malloc(size_t bytes) {
        __USING_SYS;
        if(Traits<System>::multiheap)
            return Application::_heap->alloc(bytes);
        else
            return System::_heap->alloc(bytes);
    }

    inline void * calloc(size_t n, unsigned int bytes) {
        void * ptr = malloc(n * bytes);
        memset(ptr, 0, n * bytes);
        return ptr;
    }

    inline void free(void * ptr) {
        __USING_SYS;
        if(Traits<System>::multiheap)
            Heap::typed_free(ptr);
        else
            Heap::untyped_free(System::_heap, ptr);
    }
}
```

single-heap uses System::heap



```
extern "C"
{
    // Standard C Library allocators
    inline void * malloc(size_t bytes) {
        __USING_SYS;
        if(Traits<System>::multiheap)
            return Application::_heap->alloc(bytes);
        else
            return System::_heap->alloc(bytes);
    }

    inline void * calloc(size_t n, unsigned int bytes) {
        void * ptr = malloc(n * bytes);
        memset(ptr, 0, n * bytes);
        return ptr;
    }

    inline void free(void * ptr) {
        __USING_SYS;
        if(Traits<System>::multiheap)
            Heap::typed_free(ptr);
        else
            Heap::untyped_free(System::_heap, ptr);
    }
}
```

originally, delete took only ptr, so the type must be memorized



```
extern "C" { char _end; }
```

```
class Init_Application
```

```
{
```

```
private:
```

```
    static const unsigned int HEAP_SIZE = Traits<Application>::HEAP_SIZE;
```

```
    static const unsigned int STACK_SIZE = Traits<Application>::STACK_SIZE;
```

```
public:
```

```
    Init_Application() {
```

```
        if(Traits<System>::multiheap) {
```

```
            char * heap = MMU::align_page(&_amp;_end);
```

```
            if(Traits<Build>::MODE != Traits<Build>::KERNEL)
```

```
                heap += MMU::align_page(Traits<Application>::STACK_SIZE);
```

```
            Application::_heap = new (&Application::_preheap[0]) Heap(heap, HEAP_SIZE);
```

```
        } else
```


```
            for(unsigned int frames = MMU::allocable(); frames; frames = MMU::allocable())
```

```
                System::_heap->free(MMU::alloc(frames), frames * sizeof(MMU::Page));
```

```
        db<Init>(INF) << "done!" << endl;
```

```
    }
```

```
};
```



```
extern "C" { char _end; }
```

symbol defined by the compiler to indicate  
the end of the data segment

```
class Init_Application
```

```
{
```

```
private:
```

```
    static const unsigned int HEAP_SIZE = Traits<Application>::HEAP_SIZE;
```

```
    static const unsigned int STACK_SIZE = Traits<Application>::STACK_SIZE;
```

```
public:
```

```
    Init_Application() {
```

```
        if(Traits<System>::multiheap) {
```

```
            char * heap = MMU::align_page(&_end);
```

```
            if(Traits<Build>::MODE != Traits<Build>::KERNEL)
```

```
                heap += MMU::align_page(Traits<Application>::STACK_SIZE);
```

```
            Application::_heap = new (&Application::_preheap[0]) Heap(heap, HEAP_SIZE);
```

```
        } else
```

```
            for(unsigned int frames = MMU::allocable(); frames; frames = MMU::allocable())
```

```
                System::_heap->free(MMU::alloc(frames), frames * sizeof(MMU::Page));
```

```
        db<Init>(INF) << "done!" << endl;
```

```
    }
```

```
};
```

```
extern "C" { char _end; }
```

```
class Init_Application
```

```
{
```

```
private:
```

```
    static const unsigned int HEAP_SIZE = Traits<Application>::HEAP_SIZE;
```

```
    static const unsigned int STACK_SIZE = Traits<Application>::STACK_SIZE;
```

```
public:
```

```
    Init_Application() {
```

```
        if(Traits<System>::multiheap) {
```

```
            char * heap = MMU::align_page(&_amp_end);
```

```
            if(Traits<Build>::MODE != Traits<Build>::KERNEL)
```

```
                heap += MMU::align_page(Traits<Application>::STACK_SIZE);
```

```
            Application::_heap = new (&Application::_preheap[0]) Heap(heap, HEAP_SIZE);
```

```
        } else
```


```
            for(unsigned int frames = MMU::allocable(); frames; frames = MMU::allocable())
```

```
                System::_heap->free(MMU::alloc(frames), frames * sizeof(MMU::Page));
```


```
        db<Init>(INF) << "done!" << endl;
```

```
    }
```

```
};
```



heap in data segment arranged by  
SETUP



```
extern "C" { char _end; }
```

```
class Init_Application
```

```
{
```

```
private:
```

```
    static const unsigned int HEAP_SIZE = Traits<Application>::HEAP_SIZE;
```

```
    static const unsigned int STACK_SIZE = Traits<Application>::STACK_SIZE;
```

```
public:
```

```
    Init_Application() {
```

```
        if(Traits<System>::multiheap) {
```

```
            char * heap = MMU::align_page(&_end);
```

```
            if(Traits<Build>::MODE != Traits<Build>::KERNEL)
```

```
                heap += MMU::align_page(Traits<Application>::STACK_SIZE);
```

```
            Application::_heap = new (&Application::_preheap[0]) Heap(heap, HEAP_SIZE);
```

```
        } else
```

```
            for(unsigned int frames = MMU::allocable(); frames; frames = MMU::allocable())
```


```
                System::_heap->free(MMU::alloc(frames), frames * sizeof(MMU::Page));
```

```
        db<Init>(INF) << "done!" << endl;
```

```
    }
```

```
};
```

if not a kernel, then use the stack allocated by  
SETUP, otherwise make it part of the heap



```
extern "C" { char _end; }
```

```
class Init_Application
```

```
{
```

```
private:
```

```
    static const unsigned int HEAP_SIZE = Traits<Application>::HEAP_SIZE;
```

```
    static const unsigned int STACK_SIZE = Traits<Application>::STACK_SIZE;
```

```
public:
```

```
    Init_Application() {
```

```
        if(Traits<System>::multiheap) {
```

```
            char * heap = MMU::align_page(&_end);
```

```
            if(Traits<Build>::MODE != Traits<Build>::KERNEL)
```

```
                heap += MMU::align_page(Traits<Application>::STACK_SIZE);
```

```
            Application::_heap = new (&Application::_preheap[0]) Heap(heap, HEAP_SIZE);
```

```
        } else
```

```
            for(unsigned int frames = MMU::allocable(); frames; frames = MMU::allocable())
```


```
                System::_heap->free(MMU::alloc(frames), frames * sizeof(MMU::Page));
```

```
        db<Init>(INF) << "done!" << endl;
```

```
    }
```

```
};
```

if !multiheap, put all free memory in the application's heap



```
// new allocators in types.h
void * operator new(size_t, const EPOS::System_Allocator &);
void * operator new[](size_t, const EPOS::System_Allocator &);

void * operator new(size_t, const EPOS::Scratchpad_Allocator &);
void * operator new[](size_t, const EPOS::Scratchpad_Allocator &);

// C++ dynamic memory allocators and deallocators
inline void * operator new(size_t bytes) {
    return malloc(bytes);
}

inline void * operator new[](size_t bytes) {
    return malloc(bytes);
}

inline void * operator new(size_t bytes, const EPOS::System_Allocator & allocator) {
    return _SYS::System::_heap->alloc(bytes);
}

inline void * operator new[](size_t bytes, const EPOS::System_Allocator & allocator) {
    return _SYS::System::_heap->alloc(bytes);
}

void operator delete(void * ptr);
void operator delete[](void * ptr);
void operator delete(void * ptr, size_t bytes);
void operator delete[](void * ptr, size_t bytes);
```

```
// new allocators in types.h
```

```
void * operator new(size_t, const EPOS::System_Allocator &);
```

```
void * operator new[](size_t, const EPOS::System_Allocator &);
```

```
void * operator new(size_t, const EPOS::Scratchpad_Allocator &);
```

```
void * operator new[](size_t, const EPOS::Scratchpad_Allocator &);
```

```
// C++ dynamic memory allocators and deallocators
```

```
inline void * operator new(size_t bytes) {
```

```
    return malloc(bytes);
```

```
}
```

```
inline void * operator new[](size_t bytes) {
```

```
    return malloc(bytes);
```

```
}
```

```
inline void * operator new(size_t bytes, const EPOS::System_Allocator & allocator) {
```

```
    return _SYS::System::_heap->alloc(bytes);
```

```
}
```

```
inline void * operator new[](size_t bytes, const EPOS::System_Allocator & allocator) {
```

```
    return _SYS::System::_heap->alloc(bytes);
```

```
}
```


```
void operator delete(void * ptr);
```

```
void operator delete[](void * ptr);
```

```
void operator delete(void * ptr, size_t bytes);
```

```
void operator delete[](void * ptr, size_t bytes);
```

placement new's second parameter is free  
(yet): overloading



```
// new allocators in types.h
void * operator new(size_t, const EPOS::System_Allocator &);
void * operator new[](size_t, const EPOS::System_Allocator &);

void * operator new(size_t, const EPOS::Scratchpad_Allocator &);
void * operator new[](size_t, const EPOS::Scratchpad_Allocator &);

// C++ dynamic memory allocators and deallocators
inline void * operator new(size_t bytes) {
    return malloc(bytes);
}

inline void * operator new[](size_t bytes) {
    return malloc(bytes);
}


inline void * operator new(size_t bytes, const EPOS::System_Allocator & allocator) {
    return _SYS::System::_heap->alloc(bytes);
}

inline void * operator new[](size_t bytes, const EPOS::System_Allocator & allocator) {
    return _SYS::System::_heap->alloc(bytes);
}

void operator delete(void * ptr);
void operator delete[](void * ptr);
void operator delete(void * ptr, size_t bytes);
void operator delete[](void * ptr, size_t bytes);
```

applications get the plain new operator





```
// new allocators in types.h
void * operator new(size_t, const EPOS::System_Allocator &);
void * operator new[](size_t, const EPOS::System_Allocator &);

void * operator new(size_t, const EPOS::Scratchpad_Allocator &);
void * operator new[](size_t, const EPOS::Scratchpad_Allocator &);
```

```
// C++ dynamic memory allocators and deallocators
```

```
inline void * operator new(size_t bytes) {
    return malloc(bytes);
}
```


```
inline void * operator new[](size_t bytes) {
    return malloc(bytes);
}
```

```
inline void * operator new(size_t bytes, const EPOS::System_Allocator & allocator) {
    return _SYS::System::_heap->alloc(bytes);
}
```

```
inline void * operator new[](size_t bytes, const EPOS::System_Allocator & allocator) {
    return _SYS::System::_heap->alloc(bytes);
}
```

```
void operator delete(void * ptr);
void operator delete[](void * ptr);
void operator delete(void * ptr, size_t bytes);
void operator delete[](void * ptr, size_t bytes);
```

system uses a typed one (SYSTEM)



```
// new allocators in types.h
void * operator new(size_t, const EPOS::System_Allocator &);
void * operator new[](size_t, const EPOS::System_Allocator &);

void * operator new(size_t, const EPOS::Scratchpad_Allocator &);
void * operator new[](size_t, const EPOS::Scratchpad_Allocator &);

// C++ dynamic memory allocators and deallocators
inline void * operator new(size_t bytes) {
    return malloc(bytes);
}


inline void * operator new[](size_t bytes) {
    return malloc(bytes);
}

inline void * operator new(size_t bytes, const EPOS::System_Allocator & allocator) {
    return _SYS::System::_heap->alloc(bytes);
}

inline void * operator new[](size_t bytes, const EPOS::System_Allocator & allocator) {
    return _SYS::System::_heap->alloc(bytes);
}

void operator delete(void * ptr);
void operator delete[](void * ptr);
void operator delete(void * ptr, size_t bytes);
void operator delete[](void * ptr, size_t bytes);
```

ordinary delete



```
// new allocators in types.h
void * operator new(size_t, const EPOS::System_Allocator &);
void * operator new[](size_t, const EPOS::System_Allocator &);

void * operator new(size_t, const EPOS::Scratchpad_Allocator &);
void * operator new[](size_t, const EPOS::Scratchpad_Allocator &);

// C++ dynamic memory allocators and deallocators
inline void * operator new(size_t bytes) {
    return malloc(bytes);
}


inline void * operator new[](size_t bytes) {
    return malloc(bytes);
}

inline void * operator new(size_t bytes, const EPOS::System_Allocator & allocator) {
    return _SYS::System::_heap->alloc(bytes);
}

inline void * operator new[](size_t bytes, const EPOS::System_Allocator & allocator) {
    return _SYS::System::_heap->alloc(bytes);
}

void operator delete(void * ptr);
void operator delete[](void * ptr);
void operator delete(void * ptr, size_t bytes);
void operator delete[](void * ptr, size_t bytes);
```

delete cannot be declared inline  
because of virtual destructors



```
// new allocators in types.h
void * operator new(size_t, const EPOS::System_Allocator &);
void * operator new[](size_t, const EPOS::System_Allocator &);

void * operator new(size_t, const EPOS::Scratchpad_Allocator &);
void * operator new[](size_t, const EPOS::Scratchpad_Allocator &);

// C++ dynamic memory allocators and deallocators
inline void * operator new(size_t bytes) {
    return malloc(bytes);
}

inline void * operator new[](size_t bytes) {
    return malloc(bytes);
}

inline void * operator new(size_t bytes, const EPOS::System_Allocator & allocator) {
    return _SYS::System::_heap->alloc(bytes);
}

inline void * operator new[](size_t bytes, const EPOS::System_Allocator & allocator) {
    return _SYS::System::_heap->alloc(bytes);
}

void operator delete(void * ptr);
void operator delete[](void * ptr);
void operator delete(void * ptr, size_t bytes);
void operator delete[](void * ptr, size_t bytes);
```

delete cannot be declared inline  
because of virtual destructors

```
// no more kmalloc!
```

```
_timer = new (SYSTEM) Scheduler_Timer(QUANTUM, time_slicer);
```


```
_stack = new (SYSTEM) char[stack_size];
```

```
delete _stack;
```

```
Thread::_running = new (SYSTEM) Thread(Thread::Configuration(Thread::RUNNING,  
                                                                    Thread::MAIN),
```

```
                                reinterpret_cast<int (*)()>(__epos_app_entry));
```

```
new (SYSTEM) Thread(Thread::Configuration(Thread::READY, Thread::IDLE), &Thread::idle);
```



```
// no more kmalloc!
```

```
_timer = new (SYSTEM) Scheduler_Timer(QUANTUM, time_slicer);
```


```
_stack = new (SYSTEM) char[stack_size];
```

```
delete _stack;
```

```
Thread::_running = new (SYSTEM) Thread(Thread::Configuration(Thread::RUNNING,  
                                                                    Thread::MAIN),
```

```
                                                                    reinterpret_cast<int (*)(>(__epos_app_entry));
```

```
new (SYSTEM) Thread(Thread::Configuration(Thread::READY, Thread::IDLE), &Thread::idle);
```



```
// no more kmalloc!
```

```
_timer = new (SYSTEM) Scheduler_Timer(QUANTUM, time_slicer);
```


```
_stack = new (SYSTEM) char[stack_size];
```

```
delete _stack;
```

```
Thread::_running = new (SYSTEM) Thread(Thread::Configuration(Thread::RUNNING,  
                                                                    Thread::MAIN),
```

```
                                                                    reinterpret_cast<int (*)(>(__epos_app_entry));
```

```
new (SYSTEM) Thread(Thread::Configuration(Thread::READY, Thread::IDLE), &Thread::idle);
```



```
// no more kmalloc!
```

```
_timer = new (SYSTEM) Scheduler_Timer(QUANTUM, time_slicer);
```

```
_stack = new (SYSTEM) char[stack_size];
```


```
delete _stack;
```

```
Thread::_running = new (SYSTEM) Thread(Thread::Configuration(Thread::RUNNING,  
                                                                    Thread::MAIN),
```


```
                                reinterpret_cast<int (*)()>(__epos_app_entry));
```

```
new (SYSTEM) Thread(Thread::Configuration(Thread::READY, Thread::IDLE), &Thread::idle);
```

no more kmalloc()  
typed new() instead!








# Embedded Parallel Operating System

**Coding Journey through OS Design  
–from co-routines to a multicore kernel–**

**Prof. Antônio Augusto Fröhlich, Ph.D.**

**UFSC / LISHA  
September 30, 2020**




# True Idle Waiting: Time Management

Calling `Thread::yield()` whenever a condition is not met isn't actually "busy-waiting", but we now can easily block and unblock threads to achieve a "full idle-waiting" system

Our timer tick handler in Alarm is not re-entrant

- If a handler executes for longer than a tick, events may be lost



```
class Alarm
{
private:
    typedef Timer::Tick Tick;
    typedef Relative_Queue<Alarm, Tick> Queue;

public:
    typedef TSC::Hertz Hertz;
    typedef RTC::Microsecond Microsecond;
    enum { INFINITE = RTC::INFINITE };

public:
    Alarm(const Microsecond & time, Handler * handler, int times = 1);
    ~Alarm();

    const Microsecond & period() const { return _time; }
    static Hertz frequency() { return _timer->frequency(); }

    static void delay(const Microsecond & time);

private:
    static Microsecond timer_period() {
        return 1000000 / frequency();
    }
    static Tick ticks(const Microsecond & time) {
        return (time + timer_period() / 2) / timer_period();
    }

    static void lock() { Thread::lock(); }
    static void unlock() { Thread::unlock(); }

    static void handler(const IC::Interrupt_Id & i);
```

```

class Alarm
{
private:
    typedef Timer::Tick Tick;
    typedef Relative_Queue<Alarm, Tick> Queue;

public:
    typedef TSC::Hertz Hertz;
    typedef RTC::Microsecond Microsecond;
    enum { INFINITE = RTC::INFINITE };

public:
    Alarm(const Microsecond & time, Handler * handler, int times = 1);
    ~Alarm();

    const Microsecond & period() const { return _time; }
    static Hertz frequency() { return _timer->frequency(); }

    static void delay(const Microsecond & time);


private:
    static Microsecond timer_period() {
        return 1000000 / frequency();
    }
    static Tick ticks(const Microsecond & time) {
        return (time + timer_period() / 2) / timer_period();
    }


    static void lock() { Thread::lock(); }
    static void unlock() { Thread::unlock(); }

    static void handler(const IC::Interrupt_Id & i);

```

tick-based timer





```
class Alarm
{
private:
    typedef Timer::Tick Tick;
    typedef Relative_Queue<Alarm, Tick> Queue;

public:
    typedef TSC::Hertz Hertz;
    typedef RTC::Microsecond Microsecond;
    enum { INFINITE = RTC::INFINITE };

public:
    Alarm(const Microsecond & time, Handler * handler, int times = 1);
    ~Alarm();

    const Microsecond & period() const { return _time; }
    static Hertz frequency() { return _timer->frequency(); }


    static void delay(const Microsecond & time);

private:
    static Microsecond timer_period() {
        return 1000000 / frequency();
    }
    static Tick ticks(const Microsecond & time) {
        return (time + timer_period() / 2) / timer_period();
    }

    static void lock() { Thread::lock(); }
    static void unlock() { Thread::unlock(); }

    static void handler(const IC::Interrupt_Id & i);
```

Relative\_Queue is an ordered queue in which each element's rank is an offset to the previous



```
class Alarm
{
private:
    typedef Timer::Tick Tick;
    typedef Relative_Queue<Alarm, Tick> Queue;

public:
    typedef TSC::Hertz Hertz;
    typedef RTC::Microsecond Microsecond;
    enum { INFINITE = RTC::INFINITE };

public:
    Alarm(const Microsecond & time, Handler * handler, int times = 1);
    ~Alarm();

    const Microsecond & period() const { return _time; }
    static Hertz frequency() { return _timer->frequency(); }


    static void delay(const Microsecond & time);

private:
    static Microsecond timer_period() {
        return 1000000 / frequency();
    }
    static Tick ticks(const Microsecond & time) {
        return (time + timer_period() / 2) / timer_period();
    }

    static void lock() { Thread::lock(); }
    static void unlock() { Thread::unlock(); }

    static void handler(const IC::Interrupt_Id & i);
```

import, don't redefine!



```
class Alarm
{
private:
    typedef Timer::Tick Tick;
    typedef Relative_Queue<Alarm, Tick> Queue;

public:
    typedef TSC::Hertz Hertz;
    typedef RTC::Microsecond Microsecond;
    enum { INFINITE = RTC::INFINITE };

public:
    Alarm(const Microsecond & time, Handler * handler, int times = 1),
    ~Alarm();

    const Microsecond & period() const { return _time; }
    static Hertz frequency() { return _timer->frequency(); }


    static void delay(const Microsecond & time);

private:
    static Microsecond timer_period() {
        return 1000000 / frequency();
    }
    static Tick ticks(const Microsecond & time) {
        return (time + timer_period() / 2) / timer_period();
    }

    static void lock() { Thread::lock(); }
    static void unlock() { Thread::unlock(); }

    static void handler(const IC::Interrupt_Id & i);
```

periodically invoked handler



```
class Alarm
{
private:
    typedef Timer::Tick Tick;
    typedef Relative_Queue<Alarm, Tick> Queue;

public:
    typedef TSC::Hertz Hertz;
    typedef RTC::Microsecond Microsecond;
    enum { INFINITE = RTC::INFINITE };

public:
    Alarm(const Microsecond & time, Handler * handler, int times = 1);
    ~Alarm();

    const Microsecond & period() const { return _time; }
    static Hertz frequency() { return _timer->frequency(); }

    static void delay(const Microsecond & time);


private:
    static Microsecond timer_period() {
        return 1000000 / frequency();
    }
    static Tick ticks(const Microsecond & time) {
        return (time + timer_period() / 2) / timer_period();
    }

    static void lock() { Thread::lock(); }
    static void unlock() { Thread::unlock(); }

    static void handler(const IC::Interrupt_Id & i);
```

function to wait for a given time





```
class Alarm
{
private:
    typedef Timer::Tick Tick;
    typedef Relative_Queue<Alarm, Tick> Queue;

public:
    typedef TSC::Hertz Hertz;
    typedef RTC::Microsecond Microsecond;
    enum { INFINITE = RTC::INFINITE };

public:
    Alarm(const Microsecond & time, Handler * handler, int times = 1);
    ~Alarm();

    const Microsecond & period() const { return _time; }
    static Hertz frequency() { return _timer->frequency(); }


    static void delay(const Microsecond & time);

private:
    static Microsecond timer_period() {
        return 1000000 / frequency();
    }
    static Tick ticks(const Microsecond & time) {
        return (time + timer_period() / 2) / timer_period();
    }

    static void lock() { Thread::lock(); }
    static void unlock() { Thread::unlock(); }

    static void handler(const IC::Interrupt_Id & i);
```

preemption control



```
class Alarm
{
private:
    typedef Timer::Tick Tick;
    typedef Relative_Queue<Alarm, Tick> Queue;

public:
    typedef TSC::Hertz Hertz;
    typedef RTC::Microsecond Microsecond;
    enum { INFINITE = RTC::INFINITE };

public:
    Alarm(const Microsecond & time, Handler * handler, int times = 1);
    ~Alarm();

    const Microsecond & period() const { return _time; }
    static Hertz frequency() { return _timer->frequency(); }


    static void delay(const Microsecond & time);

private:
    static Microsecond timer_period() {
        return 1000000 / frequency();
    }
    static Tick ticks(const Microsecond & time) {
        return (time + timer_period() / 2) / timer_period();
    }

    static void lock() { Thread::lock(); }
    static void unlock() { Thread::unlock(); }

    static void handler(const IC::Interrupt_Id & i);
```

(peseudo) interrupt handler




```
private:
    Microsecond _time;
    Handler * _handler;
    int _times;
    Tick _ticks;
    Queue::Element _link;

    static Alarm_Timer * _timer;
    static volatile Tick _elapsed;
    static Queue _request;
};

Alarm::Alarm(const Microsecond & time, Handler * handler, unsigned int times)
: _time(time), _handler(handler), _times(times), _ticks(ticks(time)), _link(this, _ticks)
{
    lock();

    db<Alarm>(TRC) << "Alarm(t=" << time << ",tk=" << _ticks
        << ",h=" << reinterpret_cast<void *>(handler)
        << ",x=" << times << ") => " << this << endl;

    if(_ticks) {
        _request.insert(&_link);
        unlock();
    } else {
        unlock();
        (*handler)();
    }
}
```



private:

```
Microsecond _time;  
Handler * _handler;  
int _times;  
Tick _ticks;  
Queue::Element _link;
```

```
static Alarm_Timer * _timer;  
static volatile Tick _elapsed;  
static Queue _request;
```

};

```
Alarm::Alarm(const Microsecond & time, Handler * handler, unsigned int times)  
: _time(time), _handler(handler), _times(times), _ticks(ticks(time)), _link(this, _ticks)  
{
```


```
    lock();
```

```
    db<Alarm>(TRC) << "Alarm(t=" << time << ",tk=" << _ticks  
        << ",h=" << reinterpret_cast<void *>(handler)  
        << ",x=" << times << ") => " << this << endl;
```

```
    if(_ticks) {  
        _request.insert(&_link);  
        unlock();  
    } else {  
        unlock();  
        (*handler)();  
    }  
}
```

}

tick counter (overflow?)



```
private:
    Microsecond _time;
    Handler * _handler;
    int _times;
    Tick _ticks;
    Queue::Element _link;


    static Alarm_Timer * _timer;
    static volatile Tick _elapsed;
    static Queue _request;
};

Alarm::Alarm(const Microsecond & time, Handler * handler, unsigned int times)
: _time(time), _handler(handler), _times(times), _ticks(ticks(time)), _link(this, _ticks)
{
    lock();

    db<Alarm>(TRC) << "Alarm(t=" << time << ",tk=" << _ticks
        << ",h=" << reinterpret_cast<void *>(handler)
        << ",x=" << times << ") => " << this << endl;

    if(_ticks) {
        _request.insert(&_link);
        unlock();
    } else {
        unlock();
        (*handler)();
    }
}
```

insert pending event in the queue if >  
tick



```
private:
    Microsecond _time;
    Handler * _handler;
    int _times;
    Tick _ticks;
    Queue::Element _link;

    static Alarm_Timer * _timer;
    static volatile Tick _elapsed;
    static Queue _request;
};

Alarm::Alarm(const Microsecond & time, Handler * handler, unsigned int times)
: _time(time), _handler(handler), _times(times), _ticks(ticks(time)), _link(this, _ticks)
{
    lock();


    db<Alarm>(TRC) << "Alarm(t=" << time << ",tk=" << _ticks
        << ",h=" << reinterpret_cast<void *>(handler)
        << ",x=" << times << ") => " << this << endl;

    if(_ticks) {
        _request.insert(&_link);
        unlock();
    } else {
        unlock();
        (*handler)();
    }
}
```

dispatch handler immediately otherwise

```
Alarm::~Alarm()  
{  
    lock();  
    _request.remove(this);  
    unlock();  
}
```


```
void Timer::int_handler(const Interrupt_Id & i)  
{  
    if(_channels[ALARM]) {  
        _channels[ALARM]->_current[0] = _channels[ALARM]->_initial;  
        _channels[ALARM]->_handler(i);  
    }  
    :  
}
```



```
Alarm::~Alarm()
```

```
{  
    lock();  
    _request.remove(this);  
    unlock();  
}
```


destructor just remove the alarm from  
the queue



```
void Timer::int_handler(const Interrupt_Id & i)
```

```
{  
    if(_channels[ALARM]) {  
        _channels[ALARM]->_current[0] = _channels[ALARM]->_initial;  
        _channels[ALARM]->_handler(i);  
    }  
    :  
}
```






```
Alarm::~Alarm()  
{  
    lock();  
    _request.remove(this);  
    unlock();  
}
```

```
void Timer::int_handler(const Interrupt_Id & i)  
{  
    if(_channels[ALARM]) {  
        _channels[ALARM]->_current[0] = _channels[ALARM]->_initial;  
        _channels[ALARM]->_handler(i);  
    }  
    :  
}
```

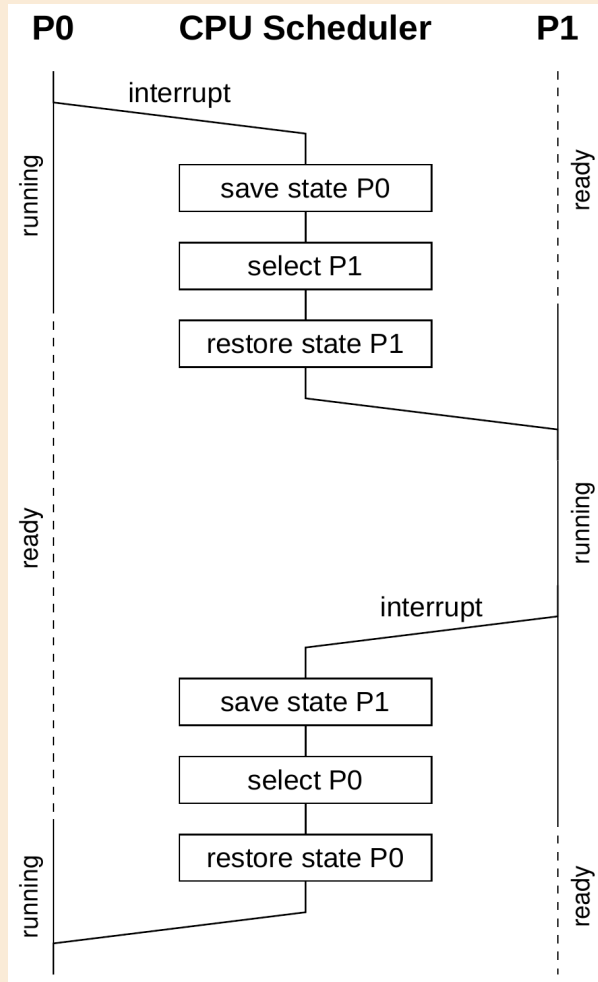
real interrupt handler for PC, which  
multiplex logical timers

```
Alarm::~Alarm()  
{  
    lock();  
    _request.remove(this);  
    unlock();  
}
```

```
void Timer::int_handler(const Interrupt_Id & i)  
{  
    if(_channels[ALARM]) {  
        _channels[ALARM]->_current[0] = _channels[ALARM]->_initial;  
        _channels[ALARM]->_handler(i);  
    }  
    :  
}
```





## Time-triggered Reschedule



`Id & i)`

`_channels[ALARM]->_initial;`





```
// Current (bad) handler
void Alarm::handler(const IC::Interrupt_Id & i)
{
    static Tick next_tick;
    static Handler * next_handler;

    lock();

    _elapsed++;

    if(next_tick)
        next_tick--;
    if(!next_tick) {
        if(next_handler)
            (*next_handler)();
        if(_request.empty())
            next_handler = 0;
        else {
            Queue::Element * e = _request.remove();
            Alarm * alarm = e->object();
            next_tick = alarm->_ticks;
            next_handler = alarm->_handler;
            if(alarm->_times != INFINITE)
                alarm->_times--;
            if(alarm->_times) {
                e->rank(alarm->_ticks);
                _request.insert(e);
            }
        }
    }
    unlock();
}
```

```

// Current (bad) handler
void Alarm::handler(const IC::Interrupt_Id & i)
{
    static Tick next_tick;
    static Handler * next_handler;


    lock();


    _elapsed++;

    if(next_tick)
        next_tick--;
    if(!next_tick) {
        if(next_handler)
            (*next_handler)();
        if(_request.empty())
            next_handler = 0;
        else {
            Queue::Element * e = _request.remove();
            Alarm * alarm = e->object();
            next_tick = alarm->_ticks;
            next_handler = alarm->_handler;
            if(alarm->_times != INFINITE)
                alarm->_times--;
            if(alarm->_times) {
                e->rank(alarm->_ticks);
                _request.insert(e);
            }
        }
    }
    unlock();
}

```

static variables!






```
// Current (bad) handler
void Alarm::handler(const IC::Interrupt_Id & i)
{
    static Tick next_tick;
    static Handler * next_handler;

    lock();

    _elapsed++;
    if(next_tick)
        next_tick--;
    if(!next_tick) {
        if(next_handler)
            (*next_handler)();
        if(_request.empty())
            next_handler = 0;
        else {
            Queue::Element * e = _request.remove();
            Alarm * alarm = e->object();
            next_tick = alarm->_ticks;
            next_handler = alarm->_handler;
            if(alarm->_times != INFINITE)
                alarm->_times--;
            if(alarm->_times) {
                e->rank(alarm->_ticks);
                _request.insert(e);
            }
        }
    }
    unlock();
}
```

tick counter increment




```
// Current (bad) handler
void Alarm::handler(const IC::Interrupt_Id & i)
{
    static Tick next_tick;
    static Handler * next_handler;

    lock();

    _elapsed++;

    if(next_tick)
        next_tick--;
    if(!next_tick) {
        if(next_handler)
            (*next_handler)();
        if(_request.empty())
            next_handler = 0;
        else {
            Queue::Element * e = _request.remove();
            Alarm * alarm = e->object();
            next_tick = alarm->_ticks;
            next_handler = alarm->_handler;
            if(alarm->_times != INFINITE)
                alarm->_times--;
            if(alarm->_times) {
                e->rank(alarm->_ticks);
                _request.insert(e);
            }
        }
    }
    unlock();
}
```

down stepping until handler is called  
what happens if next\_handler has an infinite  
loop?



```
// Current (bad) handler
void Alarm::handler(const IC::Interrupt_Id & i)
{
    static Tick next_tick;
    static Handler * next_handler;


    lock();

    _elapsed++;

    if(next_tick)
        next_tick--;
    if(!next_tick) {
        if(next_handler)
            (*next_handler)();
        if(_request.empty())
            next_handler = 0;
        else {
            Queue::Element * e = _request.remove();
            Alarm * alarm = e->object();
            next_tick = alarm->_ticks;
            next_handler = alarm->_handler;
            if(alarm->_times != INFINITE)
                alarm->_times--;
            if(alarm->_times) {
                e->rank(alarm->_ticks);
                _request.insert(e);
            }
        }
    }
    unlock();
}
```

next alarm handling






```
// Current (bad) handler
void Alarm::handler(const IC::Interrupt_Id & i)
{
    static Tick next_tick;
    static Handler * next_handler;

    lock();

    _elapsed++;

    if(next_tick)
        next_tick--;
    if(!next_tick) {
        if(next_handler)
            (*next_handler)();
        if(_request.empty())
            next_handler = 0;
        else {
            Queue::Element * e = _request.remove();
            Alarm * alarm = e->object();
            next_tick = alarm->_ticks;
            next_handler = alarm->_handler;
            if(alarm->_times != INFINITE)
                alarm->_times--;
            if(alarm->_times) {
                e->rank(alarm->_ticks);
                _request.insert(e);
            }
        }
    }
    unlock();
}
```

periodic alarm handling



```
// new handler
void Alarm::handler(const IC::Interrupt_Id & i)
{
    lock();

    _elapsed++;

    Alarm * alarm = 0;

    if(!_request.empty()) {
        if(_request.head()->promote() <= 0) {
            Queue::Element * e = _request.remove();
            alarm = e->object();
            if(alarm->_times != INFINITE)
                alarm->_times--;
            if(alarm->_times) {
                e->rank(alarm->_ticks);
                _request.insert(e);
            }
        }
    }

    unlock();

    if(alarm)
        (*alarm->_handler)();
}
```

```
// new handler
```

```
void Alarm::handler(const IC::Interrupt_Id & i)
```

```
{
```

```
    lock();
```

```
    _elapsed++;
```

```
    Alarm * alarm = 0;
```

```
    if(!_request.empty()) {
```

```
        if(_request.head()->promote() <= 0) {
```

```
            Queue::Element * e = _request.remove();
```

```
            alarm = e->object();
```

```
            if(alarm->_times != INFINITE)
```

```
                alarm->_times--;
```

```
            if(alarm->_times) {
```

```
                e->rank(alarm->_ticks);
```

```
                _request.insert(e);
```

```
            }
```

```
        }
```

```
    }
```


```
    unlock();
```


```
    if(alarm)
```

```
        (*alarm->_handler)();
```

```
}
```

first reentrance rule of thumb: local context only!





```
// new handler
void Alarm::handler(const IC::Interrupt_Id & i)
{
    lock();

    _elapsed++;


    Alarm * alarm = 0;

    if(!_request.empty()) {
        if(_request.head()->promote() <= 0) {
            Queue::Element * e = _request.remove();
            alarm = e->object();
            if(alarm->_times != INFINITE)
                alarm->_times--;
            if(alarm->_times) {
                e->rank(alarm->_ticks);
                _request.insert(e);
            }
        }
    }

    unlock();

    if(alarm)
        (*alarm->_handler)();
}
```

replacing the "if" by a "while" is tempting, but recovering the lock and dispatching the handler is troublesome if the Alarm gets destroyed in between



```
// new handler
void Alarm::handler(const IC::Interrupt_Id & i)
{
    lock();

    _elapsed++;


    Alarm * alarm = 0;

    if(!_request.empty()) {
        if(_request.head()->promote() <= 0) {
            Queue::Element * e = _request.remove();
            alarm = e->object();
            if(alarm->_times != INFINITE)
                alarm->_times--;
            if(alarm->_times) {
                e->rank(alarm->_ticks);
                _request.insert(e);
            }
        }
    }

    unlock();

    if(alarm)
        (*alarm->_handler)();
}
```

rank can be negative whenever multiple  
handlers get created for the same time  
tick



```
// new handler
void Alarm::handler(const IC::Interrupt_Id & i)
{
    lock();

    _elapsed++;


    Alarm * alarm = 0;

    if(!_request.empty()) {
        if(_request.head()->promote() <= 0) {
            Queue::Element * e = _request.remove();
            alarm = e->object();
            if(alarm->_times != INFINITE)
                alarm->_times--;
            if(alarm->_times) {
                e->rank(alarm->_ticks);
                _request.insert(e);
            }
        }
    }

    unlock();

    if(alarm)
        (*alarm->_handler)();
}
```

handling periodic alarms



```
// new handler
void Alarm::handler(const IC::Interrupt_Id & i)
{
    lock();

    _elapsed++;


    Alarm * alarm = 0;

    if(!_request.empty()) {
        if(_request.head()->promote() <= 0) {
            Queue::Element * e = _request.remove();
            alarm = e->object();
            if(alarm->_times != INFINITE)
                alarm->_times--;
            if(alarm->_times) {
                e->rank(alarm->_ticks);
                _request.insert(e);
            }
        }
    }

    unlock();

    if(alarm)
        (*alarm->_handler)();
}
```

second reentrance rule of thumb:  
release locks before dispatching!



```
// new handler
void Alarm::handler(const IC::Interrupt_Id & i)
{
    lock();

    _elapsed++;


    Alarm * alarm = 0;

    if(!_request.empty()) {
        if(_request.head()->promote() <= 0) {
            Queue::Element * e = _request.remove();
            alarm = e->object();
            if(alarm->_times != INFINITE)
                alarm->_times--;
            if(alarm->_times) {
                e->rank(alarm->_ticks);
                _request.insert(e);
            }
        }
    }

    unlock();

    if(alarm)
        (*alarm->_handler)();
}
```





```
class Handler
{
public:
    // A handler function
    typedef void (Function)();

public:
    Handler() {}
    virtual ~Handler() {}

    virtual void operator()() = 0;
};

class Function_Handler: public Handler
{
public:
    Function_Handler(Function * h): _handler(h) {}
    ~Function_Handler() {}

    void operator()() { _handler(); }

private:
    Function * _handler;
};
```

```
template<typename T>
class Functor_Handler: public Handler
{
public:
    typedef void (Functor)(T *);

public:
    Functor_Handler(Functor * h, T * p):
        _handler(h), _ptr(p) {}
    ~Functor_Handler() {}


    void operator()() { _handler(_ptr); }

private:
    Functor * _handler;
    T * _ptr;
};

class Thread_Handler : public Handler
{
public:
    Thread_Handler(Thread * h) : _handler(h) {}
    ~Thread_Handler() {}

    void operator()() { _handler->resume(); }

private:
    Thread * _handler;
};
```



```
class Handler
{
public:
    // A handler function
    typedef void (Function)();
```

```
public:
    Handler() {}
    virtual ~Handler() {}

    virtual void operator()() = 0;
};
```

```
class Function_Handler: public Handler
{
public:
    Function_Handler(Function * h): _handler(h) {}
    ~Function_Handler() {}

    void operator()() { _handler(); }

private:
    Function * _handler;
};
```

first polymorphic class in our design (abstract, indeed)  
(don't try to do yourself what the compiler can do better)

```
Functor_Handler(Functor * h, T * p):
    _handler(h), _ptr(p) {}
~Functor_Handler() {}


void operator()() { _handler(_ptr); }

private:
    Functor * _handler;
    T * _ptr;
};

class Thread_Handler : public Handler
{
public:
    Thread_Handler(Thread * h) : _handler(h) {}
    ~Thread_Handler() {}

    void operator()() { _handler->resume(); }

private:
    Thread * _handler;
};
```



```
class Handler
{
public:
    // A handler function
    typedef void (Function)();
```

```
public:
    Handler() {}
    virtual ~Handler() {}

    virtual void operator>()() = 0;
};
```

```
class Function_Handler: public Handler
{
public:
    Function_Handler(Function * h): _handler(h) {}
    ~Function_Handler() {}

    void operator>()() { _handler(); }

private:
    Function * _handler;
};
```

```
template<typename T>
class Functor_Handler: public Handler
{
public:
    typedef void (Functor)(T *);
```

```
public:
    Functor_Handler(Functor * h, T * p):
```

overload of the call operator


```
void operator>()() { _handler(_ptr); }
```

```
private:
    Functor * _handler;
    T * _ptr;
};
```

```
class Thread_Handler : public Handler
{
public:
    Thread_Handler(Thread * h) : _handler(h) {}
    ~Thread_Handler() {}

    void operator>()() { _handler->resume(); }
```

```
private:
    Thread * _handler;
};
```



```
class Handler
{
public:
    // A handler function
    typedef void (Function)();

public:
    Handler() {}
    virtual ~Handler() {}

    virtual void operator()() = 0;
};
```

```
class Function_Handler: public Handler
{
public:
    Function_Handler(Function * h): _handler(h) {}
    ~Function_Handler() {}

    void operator()() { _handler(); }

private:
    Function * _handler;
};
```

```
template<typename T>
class Functor_Handler: public Handler
{
public:
    typedef void (Functor)(T *);

public:
    Functor_Handler(Functor * h, T * p):
        _handler(h), _ptr(p) {}
    ~Functor_Handler() {}

    void operator()() { _handler(_ptr); }
```

priv

a simple indirect function call

```
class Thread_Handler : public Handler
{
public:
    Thread_Handler(Thread * h) : _handler(h) {}
    ~Thread_Handler() {}

    void operator()() { _handler->resume(); }

private:
    Thread * _handler;
};
```

a traditional Functor  
(should be variadic?)

```
typedef void (Function)();
```

```
public:
    Handler() {}
    virtual ~Handler() {}

    virtual void operator()() = 0;
};

class Function_Handler: public Handler
{
public:
    Function_Handler(Function * h): _handler(h) {}
    ~Function_Handler() {}

    void operator()() { _handler(); }

private:
    Function * _handler;
};
```

```
template<typename T>
class Functor_Handler: public Handler
{
public:
    typedef void (Functor)(T *);

public:
    Functor_Handler(Functor * h, T * p):
        _handler(h), _ptr(p) {}
    ~Functor_Handler() {}


    void operator()() { _handler(_ptr); }


private:
    Functor * _handler;
    T * _ptr;
};

class Thread_Handler : public Handler
{
public:
    Thread_Handler(Thread * h) : _handler(h) {}
    ~Thread_Handler() {}

    void operator()() { _handler->resume(); }

private:
    Thread * _handler;
};
```





```
class Handler
{
public:
    // A handler function
    typedef void (Function)();

public:
    Handler() {}
    virtual ~Handler() {}

    virtual void operator()() = 0;
};

class Function_Handler: public Handler
{
public:
    Function_Handler(Function * h): _handler(h) {}
    ~Function_Handler() {}

    void operator()() { _handler(); }
```

handler calls Thread::resume()

```
template<typename T>
class Functor_Handler: public Handler
{
public:
    typedef void (Functor)(T *);

public:
    Functor_Handler(Functor * h, T * p):
        _handler(h), _ptr(p) {}
    ~Functor_Handler() {}


    void operator()() { _handler(_ptr); }

private:
    Functor * _handler;
    T * _ptr;
};

class Thread_Handler : public Handler
{
public:
    Thread_Handler(Thread * h) : _handler(h) {}
    ~Thread_Handler() {}

    void operator()() { _handler->resume(); }

private:
    Thread * _handler;
};
```



```
class Handler
{
public:
    // A handler function
    typedef void (Function)();

public:
    Handler() {}
    virtual ~Handler() {}

    virtual void operator()() = 0;
};

class Function_Handler: public Handler
{
public:
    Function_Handler(Function * h): _handler(h) {}
    ~Function_Handler() {}

    void operator()() { _handler(); }

private:
    Function * _handler;
};
```

programmer calls Thread::suspend()

```
template<typename T>
class Functor_Handler: public Handler
{
public:
    typedef void (Functor)(T *);

public:
    Functor_Handler(Functor * h, T * p):
        _handler(h), _ptr(p) {}
    ~Functor_Handler() {}

    void operator()() { _handler(_ptr); }

private:
    Functor * _handler;
    T * _ptr;
};

class Thread_Handler : public Handler
{
public:
    Thread_Handler(Thread * h) : _handler(h) {}
    ~Thread_Handler() {}

    void operator()() { _handler->resume(); }

private:
    Thread * _handler;
};
```

```
class Mutex_Handler: public Handler
{
public:
    Mutex_Handler(Mutex * h) : _handler(h) {}
    ~Mutex_Handler() {}

    void operator()() { _handler->unlock(); }


private:
    Mutex * _handler;
};
```

handler calls Mutex::unlock()


```
class Semaphore_Handler: public Handler
{
public:
    Semaphore_Handler(Semaphore * h) : _handler(h)
    {}
    ~Semaphore_Handler() {}

    void operator()() { _handler->v(); }

private:
    Semaphore * _handler;
};
```







```
class Mutex_Handler: public Handler
{
public:
    Mutex_Handler(Mutex * h) : _handler(h) {}
    ~Mutex_Handler() {}

    void operator()() { _handler->unlock(); }


private:
    Mutex * _handler;
};
```

programmer calls Mutex::lock()

```
class Semaphore_Handler: public Handler
{
public:
    Semaphore_Handler(Semaphore * h) : _handler(h)
    {}
    ~Semaphore_Handler() {}

    void operator()() { _handler->v(); }

private:
    Semaphore * _handler;
};
```



```
class Mutex_Handler: public Handler
{
public:
    Mutex_Handler(Mutex * h) : _handler(h) {}
    ~Mutex_Handler() {}

    void operator()() { _handler->unlock(); }


private:
    Mutex * _handler;
};
```

```
class Semaphore_Handler: public Handler
{
public:
    Semaphore_Handler(Semaphore * h) : _handler(h)
    {}
    ~Semaphore_Handler() {}

    void operator()() { _handler->v(); }

private:
    Semaphore * _handler;
};
```

handler calls Semaphore::v()



```
class Mutex_Handler: public Handler
{
public:
    Mutex_Handler(Mutex * h) : _handler(h) {}
    ~Mutex_Handler() {}

    void operator()() { _handler->unlock(); }


private:
    Mutex * _handler;
};
```

```
class Semaphore_Handler: public Handler
{
public:
    Semaphore_Handler(Semaphore * h) : _handler(h)
    {}
    ~Semaphore_Handler() {}

    void operator()() { _handler->v(); }

private:
    Semaphore * _handler;
};
```

programmer calls Semaphore::p()



```
class Mutex_Handler: public Handler
{
public:
    Mutex_Handler(Mutex * h) : _handler(h) {}
    ~Mutex_Handler() {}


    void operator()() { _handler->unlock(); }

private:
    Mutex * _handler;
};
```

```
class Semaphore_Handler: public Handler
{
public:
    Semaphore_Handler(Semaphore * h) : _handler(h)
    {}
    ~Semaphore_Handler() {}

    void operator()() { _handler->v(); }

private:
    Semaphore * _handler;
};
```




```
// current handler
void Alarm::delay(const Microsecond & time)
{
    db<Alarm>(TRC) << "Alarm::delay(time="
                << time << ")" << endl;

    Tick t = _elapsed + ticks(time);

    while(_elapsed < t);
}
```

```
// new handler
void Alarm::delay(const Microsecond & time)
{
    db<Alarm>(TRC) << "Alarm::delay(time="
                << time << ")" << endl;

    Semaphore semaphore(0);
    Semaphore_Handler handler(&semaphore);
    Alarm alarm(time, &handler, 1);
    semaphore.p();
}
```



```
// current handler
void Alarm::delay(const Microsecond & time)
{
    db<Alarm>(TRC) << "Alarm::delay(time="
                << time << ")" << endl;


    Tick t = _elapsed + ticks(time);

    while(_elapsed < t);
}
```

busy-waiting!

```
// new handler
void Alarm::delay(const Microsecond & time)
{
    db<Alarm>(TRC) << "Alarm::delay(time="
                << time << ")" << endl;

    Semaphore semaphore(0);
    Semaphore_Handler handler(&semaphore);
    Alarm alarm(time, &handler, 1);
    semaphore.p();
}
```



```
// current handler
void Alarm::delay(const Microsecond & time)
{
    db<Alarm>(TRC) << "Alarm::delay(time="
                << time << ")" << endl;


    Tick t = _elapsed + ticks(time);

    while(_elapsed < t);
}
```

```
// new handler
void Alarm::delay(const Microsecond & time)
{
    db<Alarm>(TRC) << "Alarm::delay(time="
                << time << ")" << endl;

    Semaphore semaphore(0);
    Semaphore_Handler handler(&semaphore);
    Alarm alarm(time, &handler, 1);
    semaphore.p();
}
```

Semaphore has memory!!!



```
// current handler
void Alarm::delay(const Microsecond & time)
{
    db<Alarm>(TRC) << "Alarm::delay(time="
                << time << ")" << endl;

    Tick t = _elapsed + ticks(time);


    while(_elapsed < t);
}
```

```
// new handler
void Alarm::delay(const Microsecond & time)
{
    db<Alarm>(TRC) << "Alarm::delay(time="
                << time << ")" << endl;

    Semaphore semaphore(0);
    Semaphore_Handler handler(&semaphore);
    Alarm alarm(time, &handler, 1);
    semaphore.p();
}
```

if time < tick, then call v()






# Embedded Parallel Operating System

**Coding Journey through OS Design  
–from co-routines to a multicore kernel–**

**Prof. Antônio Augusto Fröhlich, Ph.D.**

**UFSC / LISHA  
September 22, 2020**




# What about preemption?

We have now an OS that does blocking synchronization

- But Threads only leave the CPU when they block or finish
- Many scheduling policies imply in preempting the CPU

Let's make our system preemptive!




```
template<> struct Traits<Thread>: public Traits<void>
{
    static const bool preemptive = true;
    static const bool trace_idle = hysterically_debugged;
    static const unsigned int QUANTUM = 10000; // us
};

void Thread::constructor_epilogue(const Log_Addr & entry, unsigned int stack_size)
{
    assert((_state != WAITING) && (_state != FINISHING)); // Invalid states

    switch(_state) {
        case RUNNING: assert(entry == __epos_app_entry); break;
        case READY: _ready.insert(&_link); break;
        case SUSPENDED: _suspended.insert(&_link); break;
        case WAITING: break; // Invalid state, for switch completion only
        case FINISHING: break; // Invalid state, for switch completion only
    }

    if(preemptive && (_state == READY) && (_link.rank() != IDLE))
        reschedule();
    else
        unlock();
}
```




```
template<> struct Traits<Thread>: public Traits<void>
{
    static const bool preemptive = true;
    static const bool trace_idle = hysterically_debugged;
    static const unsigned int QUANTUM = 10000; // us
};
```

preemption is now a feature controlled  
by programmers

```
void Thread::constructor_epilogue(const Log_Addr & entry, unsigned int stack_size)
{
    assert((_state != WAITING) && (_state != FINISHING)); // Invalid states

    switch(_state) {
        case RUNNING: assert(entry == __epos_app_entry); break;
        case READY: _ready.insert(&_amp;link); break;
        case SUSPENDED: _suspended.insert(&_amp;link); break;
        case WAITING: break; // Invalid state, for switch completion only
        case FINISHING: break; // Invalid state, for switch completion only
    }

    if(preemptive && (_state == READY) && (_link.rank() != IDLE))
        reschedule();
    else
        unlock();
}
```



```
template<> struct Traits<Thread>: public Traits<void>
{
    static const bool preemptive = true;
    static const bool trace_idle = hysterically_debugged;
    static const unsigned int QUANTUM = 10000; // us
};

void Thread::constructor_epilogue(const Log_Addr & entry, unsigned int stack_size)
{
    assert((_state != WAITING) && (_state != FINISHING)); // Invalid state
    switch(_state) {
        case RUNNING: assert(entry == __epos_app_entry); break;
        case READY: _ready.insert(&_link); break;
        case SUSPENDED: _suspended.insert(&_link); break;
        case WAITING: break; // Invalid state, for switch completion only
        case FINISHING: break; // Invalid state, for switch completion only
    }

    if(preemptive && (_state == READY) && (_link.rank() != IDLE))
        reschedule();
    else
        unlock();
}
```

reschedule at Thread creation,  
observing  
state == READY and priority != IDLE

```

void Thread::resume(bool unpreemptive)
{
    lock();

    if(_state == SUSPENDED) {
        _suspended.remove(this);
        _state = READY;
        _ready.insert(&_link);

        if(preemptive && !unpreemptive)
            reschedule();
        else
            unlock();
    } else {
        db<Thread>(WRN) << "Resume called for
unsuspended object!" << endl;

        unlock();
    }
}

```

```

void Thread::exit(int status)
{
    :
    if(prev->_joining) {
        Thread * joining = prev->_joining;
        prev->_joining = 0;
        joining->resume(true);
        lock();
    }
    :
}

```


```

void Thread::wakeup(Queue * q)
{
    // lock() must be called before entering this method
    assert(locked());

    if(!q->empty()) {
        Thread * t = q->remove()->object();
        t->_state = READY;
        t->_waiting = 0;
        _ready.insert(&t->_link);

        if(preemptive)
            reschedule();
        else
            unlock();
    } else
        unlock();
}

```



```

void Thread::resume(bool unpreemptive)
{
    lock();

    if(_state == SUSPENDED) {
        _suspended.remove(this);
        _state = READY;
        _ready.insert(&_link);

        if(preemptive && !unpreemptive)
            reschedule();
        else
            unlock();
    } else {
        db<Thread>(WRN) << "Resume called for
unsuspended object!" << endl;

        unlock();
    }
}

```

```

void Thread::exit(int status)
{
    :
    if(prev->_joining) {
        Thread * joining = prev->_joining;
        prev->_joining = 0;
        joining->resume(true);
        lock();
    }
    :
}

```


```

void Thread::wakeup(Queue * q)
{
    // lock() must be called before entering this method
    assert(locked());

    if(!q->empty()) {
        Thread * t = q->remove()->object();
        t->_state = READY;
        t->_waiting = 0;
        _ready.insert(&t->_link);

        if(preemptive)
            reschedule();
        else
            unlock();
    } else
        unlock();
}

```



reschedule at resume()

```

void Thread::resume(bool unpreemptive)
{
    lock();

    if(_state == SUSPENDED) {
        _suspended.remove(this);
        _state = READY;
        _ready.insert(&_link);

        if(preemptive && !unpreemptive)
            reschedule();
        else
            unlock();
    } else {
        db<Thread>(WRN) << "Resume called for
unsuspended object!" << endl;

        unlock();
    }
}

```

```

void Thread::exit(int status)
{
    :
    if(prev->_joining) {
        Thread * joining = prev->_joining;
        prev->_joining = 0;
        joining->resume(true);
        lock();
    }
    :
}

```


```

void Thread::wakeup(Queue * q)
{
    // lock() must be called before entering this method
    assert(locked());

    if(!q->empty()) {
        Thread * t = q->remove()->object();
        t->_state = READY;
        t->_waiting = 0;
        _ready.insert(&t->_link);


        if(preemptive)
            reschedule();
        else
            unlock();
    } else
        unlock();
}

```



prevent preemption here, since exit will always cause a reschedule





```
void Thread::resume(bool unpreemptive)
{
    lock();

    if(_state == SUSPENDED) {
        suspended_remove(this);

        reschedule at wakeup()

        if(preemptive && !unpreemptive)
            reschedule();
        else
            unlock();
    } else {
        db<Thread>(WRN) << "Resume called for
unsuspended object!" << endl;

        unlock();
    }
}
```

```
void Thread::exit(int status)
{
    :
    if(prev->_joining) {
        Thread * joining = prev->_joining;
        prev->_joining = 0;
        joining->resume(true);
        lock();
    }
    :
}
```

```
void Thread::wakeup(Queue * q)
{
    // lock() must be called before entering this method
    assert(locked());

    if(!q->empty()) {
        Thread * t = q->remove()->object();
        t->_state = READY;
        t->_waiting = 0;
        _ready.insert(&t->_link);


        if(preemptive)
            reschedule();
        else
            unlock();
    } else
        unlock();
}
```

```
void Thread::wakeup_all(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup_all(running=" << running() << ",q=" << q << ")" << endl;

    // lock() must be called before entering this method
    assert(locked());

    if(!q->empty()) {
        while(!q->empty()) {
            Thread * t = q->remove()->object();
            t->_state = READY;
            t->_waiting = 0;
            _ready.insert(&t->_link);
        }

        if(preemptive)
            reschedule();
        else
            unlock();
    } else
        unlock();
}
```



```
void Thread::wakeup_all(Queue * q)
```

```
{
```

```
    db<Thread>(TRC) << "Thread::wakeup_all(running=" << running() << ",q=" << q << ")" << endl;
```

```
    // lock() must be called before entering this method
```

```
    assert(locked());
```

```
    if(!q->empty()) {
```

```
        while(!q->empty()) {
```

```
            Thread * t = q->remove()->object();
```

```
            t->_state = READY;
```

```
            t->_waiting = 0;
```

```
            _ready.insert(&t->_link);
```

```
        }
```

```
        if(preemptive)
```

```
            reschedule();
```

```
        else
```


```
            unlock();
```

```
    } else
```


```
        unlock();
```

```
}
```

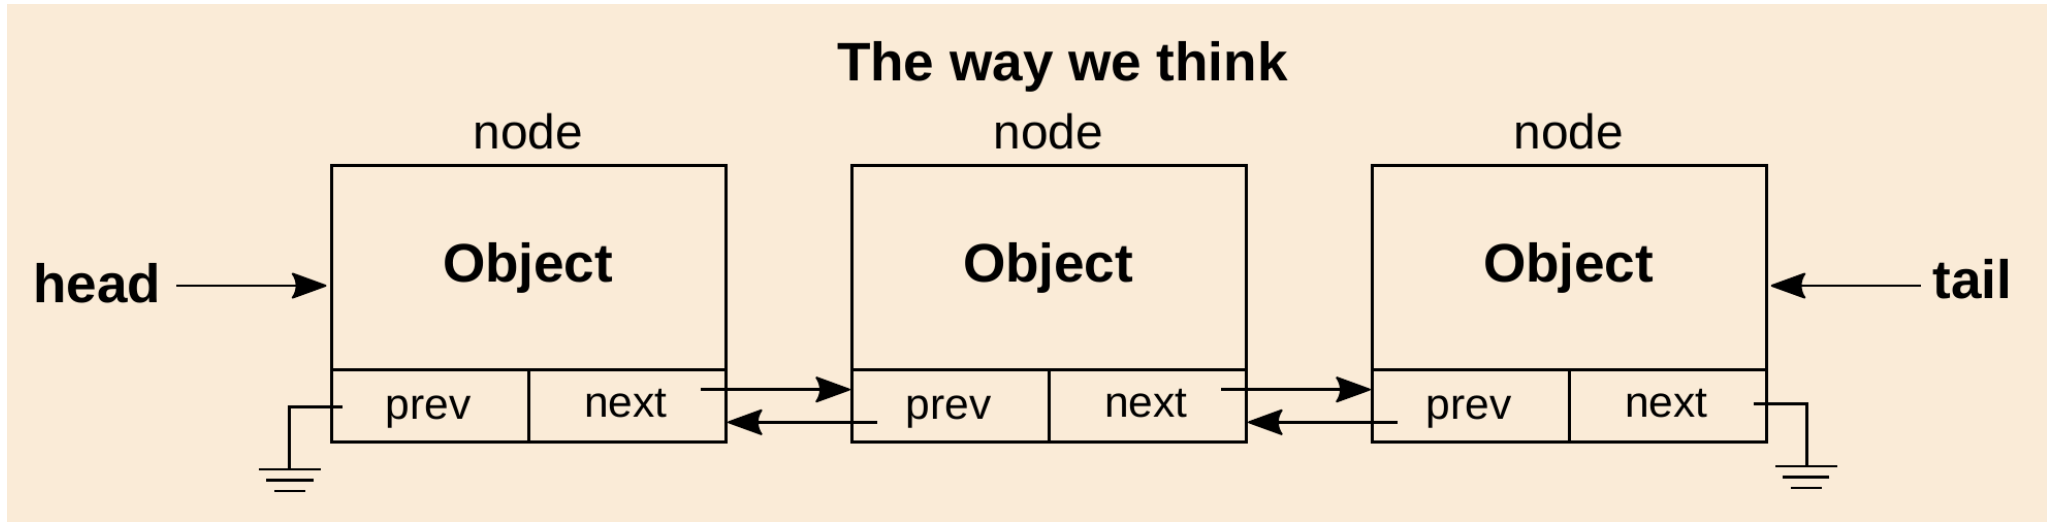

wake up everybody before calling  
reschedule()




# About Lists at System-level




# About Lists at System-level




# About Lists at System-level



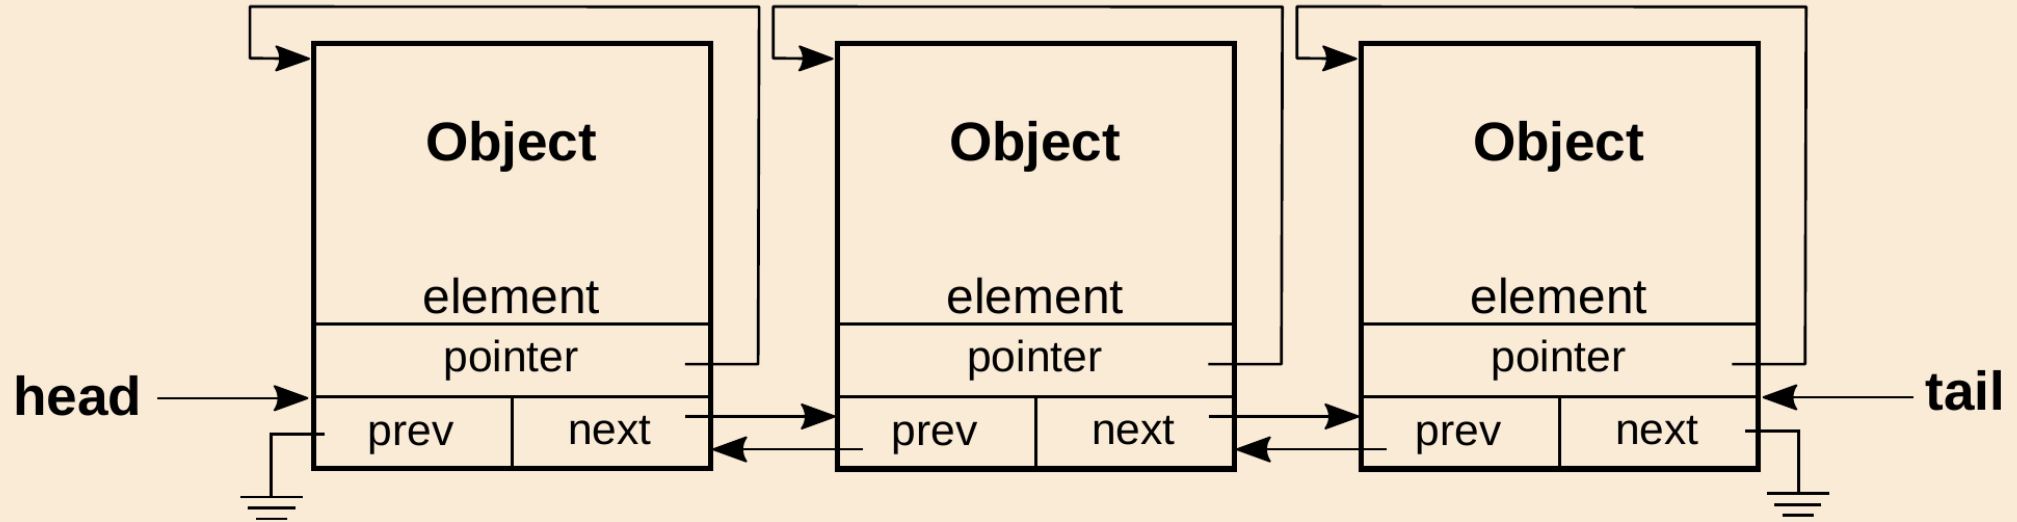
The way it usually is




# About Lists at System-level



The way it is in EPOS






# Embedded Parallel Operating System

**Coding Journey through OS Design  
–from co-routines to a multicore kernel–**

**Prof. Antônio Augusto Fröhlich, Ph.D.**

**UFSC / LISHA  
September 30, 2020**





# Doing nothing well!

Our idle() function can be improved!

It is being explicitly called all around

Making it a thread, would also eliminate many test for READY being empty

We need a idle Thread that **does nothing well!**

```

class Thread
{
public:
    // Thread Priority
    typedef unsigned int Priority;
    enum {
        MAIN    = 0,
        HIGH    = 1,
        NORMAL  =(unsigned(1)<<(sizeof(int)*8-1))-4,
        LOW     =(unsigned(1)<<(sizeof(int)*8-1))-3,
        IDLE    =(unsigned(1)<<(sizeof(int)*8-1))-2
    };

protected:
    static volatile unsigned int _thread_count;
};


void Thread::constructor_prologue(unsigned int
stack_size)
{
    lock();
    _thread_count++;
    _stack = reinterpret_cast<char *>(
        kmalloc(stack_size));
}

```

```

Thread::~Thread()
{
    lock();
    // The running thread cannot delete itself!
    assert(_state != RUNNING);
    switch(_state) {
        case RUNNING: // For switch completion only:
            the running thread would have deleted itself! Stack
            wouldn't have been released!
            exit(-1);
            break;
        case READY:
            _ready.remove(this);
            _thread_count--;
            break;
        case SUSPENDED:
            _suspended.remove(this);
            _thread_count--;
            break;
        case WAITING:
            _waiting->remove(this);
            _thread_count--;
            break;
        case FINISHING: // Already called exit()
            break;
    }
    if(_joining)
        _joining->resume();
    unlock();
    kfree(_stack);
}

```



```

class Thread
{
public:
    // Thread Priority
    typedef unsigned int Priority;
    enum {
        MAIN    = 0,
        HIGH     = 1,
        NORMAL   = (unsigned(1)<<(sizeof(int)*8-1))-4,
        LOW      = (unsigned(1)<<(sizeof(int)*8-1))-3,
        IDLE     = (unsigned(1)<<(sizeof(int)*8-1))-2
    };

protected:
    static volatile unsigned int _thread_count;
};


void Thread::constructor_prologue(unsigned int
stack_size)
{
    lock();
    _thread_count++;
    _stack = reinterpret_cast<char *>(
        kmalloc(stack_size));
}

```

Th

```
{
```

rearranging priorities and including IDLE



```

// The running thread cannot delete itself!
assert(_state != RUNNING);
switch(_state) {
case RUNNING: // For switch completion only:
the running thread would have deleted itself! Stack
wouldn't have been released!
    exit(-1);
    break;
case READY:
    _ready.remove(this);
    _thread_count--;
    break;
case SUSPENDED:
    _suspended.remove(this);
    _thread_count--;
    break;
case WAITING:
    _waiting->remove(this);
    _thread_count--;
    break;
case FINISHING: // Already called exit()
    break;
}
if(_joining)
    _joining->resume();
unlock();
kfree(_stack);
}

```

```

class Thread
{
public:
    // Thread Priority
    typedef unsigned int Priority;
    enum {
        MAIN    = 0,
        HIGH    = 1,
        NORMAL  = (unsigned(1)<<(sizeof(int)*8-1))-4,
        LOW     = (unsigned(1)<<(sizeof(int)*8-1))-3,
        IDLE    = (unsigned(1)<<(sizeof(int)*8-1))-2
    };

protected:
    static volatile unsigned int _thread_count;
};

void Thread::constructor_prologue(unsigned int
stack_size)
{
    lock();
    _thread_count++;
    _stack = reinterpret_cast<char *>(
        kmalloc(stack_size));
}

```

```

Thread::~Thread()
{
    lock();
    // The running thread cannot delete itself!
    assert( state != RUNNING);


```

Thread counting to decide when it's time to reboot or halt the system

```

    the
    wouldn't have been released!
    exit(-1);
    break;
case READY:
    _ready.remove(this);
    _thread_count--;
    break;
case SUSPENDED:
    _suspended.remove(this);
    _thread_count--;
    break;
case WAITING:
    _waiting->remove(this);
    _thread_count--;
    break;
case FINISHING: // Already called exit()
    break;
}
if(_joining)
    _joining->resume();
unlock();
kfree(_stack);
}

```



a running Thread cannot delete itself!  
how could it finish the delete operation?


```
typedef unsigned int Priority;
enum {
    MAIN    = 0,
    HIGH    = 1,
    NORMAL  =(unsigned(1)<<(sizeof(int)*8-1))-4,
    LOW     =(unsigned(1)<<(sizeof(int)*8-1))-3,
    IDLE    =(unsigned(1)<<(sizeof(int)*8-1))-2
};
```

```
protected:
    static volatile unsigned int _thread_count;
};
```

```
void Thread::constructor_prologue(unsigned int
stack_size)
{
    lock();
    _thread_count++;
    _stack = reinterpret_cast<char*>(
        kmalloc(stack_size));
}
```

Thread::~Thread()

```
{
    lock();
    // The running thread cannot delete itself!
    assert(_state != RUNNING);
    switch(_state) {
        case RUNNING: // For switch completion only:
            the running thread would have deleted itself! Stack
            wouldn't have been released!
            exit(-1);
            break;
        case READY:
            _ready.remove(this);
            _thread_count--;
            break;
        case SUSPENDED:
            _suspended.remove(this);
            _thread_count--;
            break;
        case WAITING:
            _waiting->remove(this);
            _thread_count--;
            break;
        case FINISHING: // Already called exit()
            break;
    }
    if(_joining)
        _joining->resume();
    unlock();
    kfree(_stack);
}
```



```

class Thread
{
public:
    // Thread Priority
    typedef unsigned int Priority;
    enum {
        MAIN    = 0,
        HIGH    = 1,
        NORMAL  = (unsigned(1)<<(sizeof(int)*8-1))-4,
        LOW     = (unsigned(1)<<(sizeof(int)*8-1))-3,
        IDLE    = (unsigned(1)<<(sizeof(int)*8-1))-2
    };
    static volatile unsigned int _thread_count;

    void Thread::constructor_prologue(unsigned int
    stack_size)
    {
        lock();
        _thread_count++;
        _stack = reinterpret_cast<char*>(
            kmalloc(stack_size));
    }
}


```

the "expected" case

```

Thread::~Thread()
{
    lock();
    // The running thread cannot delete itself!
    assert(_state != RUNNING);
    switch(_state) {
        case RUNNING: // For switch completion only:
            the running thread would have deleted itself! Stack
            wouldn't have been released!
            exit(-1);
            break;
        case READY:
            _ready.remove(this);
            _thread_count--;
            break;
        case SUSPENDED:
            _suspended.remove(this);
            _thread_count--;
            break;
        case WAITING:
            _waiting->remove(this);
            _thread_count--;
            break;
        case FINISHING: // Already called exit()
            break;
    }
    if(_joining)
        _joining->resume();
    unlock();
    kfree(_stack);
}

```



```

void Thread::suspend()
{
    lock();

    if(_running != this)
        _ready.remove(this);

    _state = SUSPENDED;
    _suspended.insert(&_link);

    if(_running == this) {
        _running = _ready.remove()->object();
        _running->_state = RUNNING;
        dispatch(this, _running);
    }

    unlock();
}

```

```

void Thread::yield()
{
    lock();

    Thread * prev = _running;
    prev->_state = READY;
    _ready.insert(&prev->_link);
    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(prev, _running);

    unlock();
}

```

```

void Thread::exit(int status)
{
    lock();

    Thread * prev = _running;
    prev->_state = FINISHING;
    *reinterpret_cast<int *>(prev->_stack)
        = status;
    _thread_count--;


    if(prev->_joining) {
        Thread * joining = prev->_joining;
        prev->_joining = 0;
        joining->resume(); // implicit unlock()
        lock();
    }

    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(prev, _running);

    unlock();
}

```



```

void Thread::suspend()
{
    lock();
    if(_running != this)
        _ready.remove(this);

    _state = SUSPENDED;
    _suspended.insert(&_link);

    if(_running == this) {
        _running = _ready.remove()->object();
        _running->_state = RUNNING;
        dispatch(this, _running);
    }
    unlock();
}

```

```

void Thread::yield()
{
    lock();

    Thread * prev = _running;
    prev->_state = READY;
    _ready.insert(&prev->_link);
    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(prev, _running);
    unlock();
}

```

no more \_ready.empty() tests!


```

Thread * prev = _running;
prev->_state = FINISHING;
*reinterpret_cast<int *>(prev->_stack)
    = status;
_thread_count--;

if(prev->_joining) {
    Thread * joining = prev->_joining;
    prev->_joining = 0;
    joining->resume(); // implicit unlock()
    lock();
}

_running = _ready.remove()->object();
_running->_state = RUNNING;
dispatch(prev, _running);
unlock();
}

```





```

void Thread::suspend()
{
    lock();

    if(_running != this)
        _ready.remove(this);

    _state = SUSPENDED;
    _suspended.insert(&_link);

    if(_running == this) {
        _running = _ready.remove()->object();
        _running->_state = RUNNING;
        dispatch(this, _running);
    }

    unlock();
}

```

```

void Thread::yield()
{
    lock();

    Thread * prev = _running;
    prev->_state = READY;
    _ready.insert(&prev->_link);
    _running = _ready.remove()->object();
    _running->_state = RUNNING;
    dispatch(prev, _running);
    unlock();
}

```

```

void Thread::exit(int status)
{
    lock();


    Thread * prev = _running;
    prev->_state = FINISHING;
    *reinterpret_cast<int *>(prev->_stack)
        = status;
    _thread_count--;

    if(prev->_joining) {
        Thread * joining = prev->_joining;
        prev->_joining = 0;
        joining->resume(); // implicit unlock()
        lock();
    }

    t();

    unlock();
}

```



no more \_ready.empty() tests!

```

void Thread::suspend()
{
    lock();
    if(_running != this)
        _ready.remove(this);

    _state = SUSPENDED;
    _suspended.insert(&_link);
    if (
        no more _ready.empty() tests!
        last Thread now handled by idle
    );
    dispatch(this, _running);
}
unlock();
}

```

```


void Thread::yield()
{
    lock();
    Thread * prev = _running;
    prev->_state = READY;
    _ready.insert(&prev->_link);
    _running = _ready.remove()->object();
    _running->_state = RUNNING;
    dispatch(prev, _running);
    unlock();
}

```

```

void Thread::exit(int status)
{
    lock();
    Thread * prev = _running;
    prev->_state = FINISHING;
    *reinterpret_cast<int *>(prev->_stack)
        = status;
    _thread_count--;
    if(prev->_joining) {
        Thread * joining = prev->_joining;
        prev->_joining = 0;
        joining->resume(); // implicit unlock()
        lock();
    }
    _running = _ready.remove()->object();
    _running->_state = RUNNING;
    dispatch(prev, _running);
    unlock();
}

```



```

int Thread::idle()
{
    while(_thread_count > 1) {
        if(Traits<Thread>::trace_idle)
            db<Thread>(TRC) << "Thread::idle(this="
                << running() << ")" << endl;
        CPU::int_enable();
        CPU::halt();
    }

    CPU::int_disable();
    db<Thread>(WRN)
        << "The last thread has exited!" << endl;
    if(reboot) {
        db<Thread>(WRN)
            << "Rebooting the machine ..." << endl;
        Machine::reboot();
    } else {
        db<Thread>(WRN)
            << "Halting the machine ..." << endl;
        CPU::halt();
    }

    return 0;
}

extern "C" { void __epos_app_entry(); }

```

```

void Thread::init()
{
    db<Init,Thread>(TRC) << "Thread::init()" << endl;

    Thread::_running =
        new (kmallocc(sizeof(Thread))) Thread(
            Thread::Configuration(Thread::RUNNING,
                Thread::MAIN), __epos_app_entry));

    new (kmallocc(sizeof(Thread))) Thread(
        Thread::Configuration(Thread::READY,
            Thread::IDLE), &Thread::idle);

    _timer = new (kmallocc(sizeof(Scheduler_Timer)))
        Scheduler_Timer(QUANTUM, time_slicer);

    This_Thread::not_booting();
}

class Init_First
{
public:
    Init_First() {
        db<Init>(TRC) << "Init_First()" << endl;
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::running()->_context->load();
    }
};

```

```

int Thread::idle()
{
    while(_thread_count > 1) {
        if(Traits<Thread>::trace_idle)
            db<Thread>(TRC) << "Thread::idle(this="
                << running() << ")" << endl;
        CPU::int_enable();
        CPU::halt();
    }

    CPU::int_disable();
    db<Thread>(WRN)
        << "The last thread has exited!" << endl;
    if(reboot) {
        db<Thread>(WRN)
            << "Rebooting the machine ..." << endl;
        Machine::reboot();
    } else {
        db<Thread>(WRN)
            << "Halting the machine ..." << endl;
        CPU::halt();
    }

    return 0;
}

extern "C" { void __epos_app_entry(); }

```

```

void Thread::init()
{
    db<Init>(TRC) << "Thread::init()" << endl;

    Thread::_running =
        new (kmalloc(sizeof(Thread))) Thread(
            Thread::Configuration(Thread::RUNNING,
                Thread::MAIN), __epos_app_entry));

    new (kmalloc(sizeof(Thread))) Thread(
        Thread::Configuration(Thread::READY,
            Thread::IDLE), &Thread::idle);

    _timer = new (kmalloc(sizeof(Scheduler_Timer)))
        Scheduler_Timer(QUANTUM, time_slicer);

    This_Thread::not_booting();
}

class Init_First
{
public:
    Init_First() {
        db<Init>(TRC) << "Init_First()" << endl;
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::running()->_context->load();
    }
};

```

idle loop until there is only one thread in the system (i.e. idle)

```

int Thread::idle()
{
    while(_thread_count > 1) {
        if(Traits<Thread>::trace_idle)
            db<Thread>(TRC) << "Thread::idle(this="
                << running() << ")" << endl;
        CPU::int_enable();
        CPU::halt();
    }

    CPU::int_disable();
    db<Thread>(WRN)
        << "The last thread has exited!" << endl;
    if(reboot) {
        db<Thread>(WRN)
            << "Rebooting the machine ..." << endl;
        Machine::reboot();
    } else {
        db<Thread>(WRN)
            << "Halting the machine ..." << endl;
        CPU::halt();
    }

    return 0;
}

extern "C" { void __epos_app_entry(); }

```

```

void Thread::init()
{
    db<Init,Thread>(TRC) << "Thread::init()" << endl;

    Thread::_running =
        new (kmallocc(sizeof(Thread))) Thread(
            Thread::Configuration(Thread::RUNNING,
                last Thread exit handling
            ));

    new (kmallocc(sizeof(Thread))) Thread(
        Thread::Configuration(Thread::READY,
            Thread::IDLE), &Thread::idle);

    _timer = new (kmallocc(sizeof(Scheduler_Timer)))
        Scheduler_Timer(QUANTUM, time_slicer);

    This_Thread::not_booting();
}

class Init_First
{
public:
    Init_First() {
        db<Init>(TRC) << "Init_First()" << endl;
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::running()->_context->load();
    }
};

```

```

int Thread::idle()
{
    while(_thread_count > 1) {
        if(Traits<Thread>::trace_idle)
            db<Thread>(TRC) << "Thread::idle(this="
                << running() << ")" << endl;
        CPU::int_enable();
        CPU::halt();
    }

    CPU::int_disable();
    db<Thread>(WRN)
        << "The last thread has exited!" << endl;
    if(reboot) {
        db<Thread>(WRN)
            << "Rebooting the machine ..." << endl;
        Machine::reboot();
    } else {
        db<Thread>(WRN)
            << "Halting the machine ..." << endl;
        CPU::halt();
    }

    return 0;
}

extern "C" { void __epos_app_entry(); }

```

```

void Thread::init()
{
    db<Init,Thread>(TRC) << "Thread::init()" << endl;

    Thread::_running =
        new (kmallocc(sizeof(Thread))) Thread(
            Thread::Configuration(Thread::RUNNING,
                Thread::MAIN), __epos_app_entry));

    new (kmallocc(sizeof(Thread))) Thread(
        Thread::Configuration(Thread::READY,
            Thread::IDLE), &Thread::idle);

    _timer = new (kmallocc(sizeof(Scheduler_Timer)))
        Scheduler_Timer(QUANTUM, time_slicer);

    This_Thread::not_booting();
}

class Init_First
{
public:
    1
};

Thread::running()->_context->load();
};

```

If EPOS is a library, then adjust the application entry point to \_\_epos\_app\_entry, which will directly call main(). In this case, \_init will have already been called before Init\_Application to construct MAIN's global objects.

```

int Thread::idle()
{
    while(_thread_count > 1) {
        if(Traits<Thread>::trace_idle)
            Thread::idle(this="
            create the main Thread with
            state=RUNNING and priority = MAIN
            " << endl;
    }

    CPU::int_disable();
    db<Thread>(WRN)
        << "The last thread has exited!" << endl;
    if(reboot) {
        db<Thread>(WRN)
            << "Rebooting the machine ..." << endl;
        Machine::reboot();
    } else {
        db<Thread>(WRN)
            << "Halting the machine ..." << endl;
        CPU::halt();
    }

    return 0;
}

extern "C" { void __epos_app_entry(); }

```

```

void Thread::init()
{
    db<Init,Thread>(TRC) << "Thread::init()" << endl;

    Thread::_running =
        new (kmalloc(sizeof(Thread))) Thread(
            Thread::Configuration(Thread::RUNNING,
                Thread::MAIN), __epos_app_entry));

    new (kmalloc(sizeof(Thread))) Thread(
        Thread::Configuration(Thread::READY,
            Thread::IDLE), &Thread::idle);

    _timer = new (kmalloc(sizeof(Scheduler_Timer)))
        Scheduler_Timer(QUANTUM, time_slicer);

    This_Thread::not_booting();
}

class Init_First
{
public:
    Init_First() {
        db<Init>(TRC) << "Init_First()" << endl;
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::running()->_context->load();
    }
};

```

```

int Thread::idle()
{
    while(_thread_count > 1) {
        if(Traits<Thread>::trace_idle)
            db<Thread>(TRC) << "Thread::idle(this="
                << running() << ")" << endl;
        CPU::int_enable();
        CPU::halt();
    }
}

```

create the idle Thread with  
state = READY and priority = IDLE

```

        << "The last thread has exited!" << endl;
    if(reboot) {
        db<Thread>(WRN)
            << "Rebooting the machine ..." << endl;
        Machine::reboot();
    } else {
        db<Thread>(WRN)
            << "Halting the machine ..." << endl;
        CPU::halt();
    }

    return 0;
}

extern "C" { void __epos_app_entry(); }

```

```

void Thread::init()
{
    db<Init,Thread>(TRC) << "Thread::init()" << endl;

    Thread::_running =
        new (kmallocc(sizeof(Thread))) Thread(
            Thread::Configuration(Thread::RUNNING,
                Thread::MAIN), __epos_app_entry));

    new (kmallocc(sizeof(Thread))) Thread(
        Thread::Configuration(Thread::READY,
            Thread::IDLE), &Thread::idle);

    _timer = new (kmallocc(sizeof(Scheduler_Timer)))
        Scheduler_Timer(QUANTUM, time_slicer);

    This_Thread::not_booting();
}

class Init_First
{
public:
    Init_First() {
        db<Init>(TRC) << "Init_First()" << endl;
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::running()->_context->load();
    }
};

```



```

int Thread::idle()
{
    while(_thread_count > 1) {
        if(Traits<Thread>::trace_idle)
            db<Thread>(TRC) << "Thread::idle(this="
                << running() << ")" << endl;
        CPU::int_enable();
        CPU::halt();
    }

    CPU::int_disable();

    install a time slicer for time-sharing algorithms

    << "Rebooting the machine ..." << endl;
    Machine::reboot();
} else {
    db<Thread>(WRN)
        << "Halting the machine ..." << endl;
    CPU::halt();
}

return 0;
}

extern "C" { void __epos_app_entry(); }

```

```

void Thread::init()
{
    db<Init,Thread>(TRC) << "Thread::init()" << endl;

    Thread::_running =
        new (kmallocc(sizeof(Thread))) Thread(
            Thread::Configuration(Thread::RUNNING,
                Thread::MAIN), __epos_app_entry));

    new (kmallocc(sizeof(Thread))) Thread(
        Thread::Configuration(Thread::READY,
            Thread::IDLE), &Thread::idle);

    _timer = new (kmallocc(sizeof(Scheduler_Timer)))
        Scheduler_Timer(QUANTUM, time_slicer);

    This_Thread::not_booting();
}

class Init_First
{
public:
    Init_First() {
        db<Init>(TRC) << "Init_First()" << endl;
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::running()->_context->load();
    }
};

```

```

int Thread::idle()
{
    while(_thread_count > 1) {
        if(Traits<Thread>::trace_idle)
            db<Thread>(TRC) << "Thread::idle(this="
                << running() << ")" << endl;
        CPU::int_enable();
        CPU::halt();
    }

    CPU::int_disable();
    db<Thread>(WRN)
        << "The last thread has exited!" << endl;
    if(reboot) {
        transition from CPU-based locking to
        thread-based locking
        db<Thread>(WRN)
            << "Halting the machine ..." << endl;
        CPU::halt();
    }

    return 0;
}

extern "C" { void __epos_app_entry(); }

```

```

void Thread::init()
{
    db<Init,Thread>(TRC) << "Thread::init()" << endl;

    Thread::_running =
        new (kmallocc(sizeof(Thread))) Thread(
            Thread::Configuration(Thread::RUNNING,
                Thread::MAIN), __epos_app_entry));

    new (kmallocc(sizeof(Thread))) Thread(
        Thread::Configuration(Thread::READY,
            Thread::IDLE), &Thread::idle);

    _timer = new (kmallocc(sizeof(Scheduler_Timer)))
        Scheduler_Timer(QUANTUM, time_slicer);

    This_Thread::not_booting();
}

class Init_First
{
public:
    Init_First() {
        db<Init>(TRC) << "Init_First()" << endl;
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::running()->_context->load();
    }
};

```

```

int Thread::idle()
{
    while(_thread_count > 1) {
        if(Traits<Thread>::trace_idle)
            db<Thread>(TRC) << "Thread::idle(this="
                << running() << ")" << endl;
        CPU::int_enable();
        CPU::halt();
    }

    CPU::int_disable();
    db<Thread>(WRN)
        << "The last thread has exited!" << endl;
    if(reboot) {
        db<Thread>(WRN)
            << "Rebooting the machine ..." << endl;
        Machine::reboot();
    } else {
        db<Thread>(WRN)
            << "Halting the machine ..." << endl;
        CPU::halt();
    }
}

```

simple, single-core scenarios might not even need multithreading!

```
extern "C" { void __epos_app_entry(); }
```

```

void Thread::init()
{
    db<Init,Thread>(TRC) << "Thread::init()" << endl;

    Thread::_running =
        new (kmallocc(sizeof(Thread))) Thread(
            Thread::Configuration(Thread::RUNNING,
                Thread::MAIN), __epos_app_entry));

    new (kmallocc(sizeof(Thread))) Thread(
        Thread::Configuration(Thread::READY,
            Thread::IDLE), &Thread::idle);

    _timer = new (kmallocc(sizeof(Scheduler_Timer)))
        Scheduler_Timer(QUANTUM, time_slicer);

    This_Thread::not_booting();
}

class Init_First
{
public:
    Init_First() {
        db<Init>(TRC) << "Init_First()" << endl;
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::running()->_context->load();
    }
};

```

```

int Thread::idle()
{
    while(_thread_count > 1) {
        if(Traits<Thread>::trace_idle)
            db<Thread>(TRC) << "Thread::idle(this="
                << running() << ")" << endl;
        CPU::int_enable();
        CPU::halt();
    }

    CPU::int_disable();
    db<Thread>(WRN)
        << "The last thread has exited!" << endl;
    if(reboot) {
        db<Thread>(WRN)
            << "Rebooting the machine ..." << endl;
        Machine::reboot();
    } else {
        db<Thread>(WRN)
            << "Halting the machine ..." << endl;
    }
}

```

loading the first (main) Thread's context, causing it to start executing

```

extern "C" { void __epos_app_entry(); }

```

```

void Thread::init()
{
    db<Init,Thread>(TRC) << "Thread::init()" << endl;

    Thread::_running =
        new (kmallocc(sizeof(Thread))) Thread(
            Thread::Configuration(Thread::RUNNING,
                Thread::MAIN), __epos_app_entry));

    new (kmallocc(sizeof(Thread))) Thread(
        Thread::Configuration(Thread::READY,
            Thread::IDLE), &Thread::idle);

    _timer = new (kmallocc(sizeof(Scheduler_Timer)))
        Scheduler_Timer(QUANTUM, time_slicer);

    This_Thread::not_booting();
}

class Init_First
{
public:
    Init_First() {
        db<Init>(TRC) << "Init_First()" << endl;
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }
    }

    Thread::running()->_context->load();
};

```

```
int Thread::idle()
{
    while(_thread_count > 1) {
        if(Traits<Thread>::trace_idle)
            db<Thread>(TRC) << "Thread::idle(this="
                << running() << ")" << endl;
        CPU::int_enable();
        CPU::halt();
    }
}
```

```
CPU::int_disable();
db<Thread>(WRN)
    << "The last thread has exited!" << endl;
if(reboot) {
```

### Context::load()

From the perspective of the CPU, a process is already running on the CPU since it was switched on. Short after, the stack pointer was set and the process, which already had an implicit address space, code and data, got a stack. The OS, however, is not aware of this. So we cannot “switch” context. We must load the first context.

```
    e ..." << endl;
```

```
    ..." << endl;
```

```
; }
```

```
void Thread::init()
{
    db<Init,Thread>(TRC) << "Thread::init()" <<endl;

    Thread::_running =
        new (kmallocc(sizeof(Thread))) Thread(
            Thread::Configuration(Thread::RUNNING,
                Thread::MAIN), __epos_app_entry));

    new (kmallocc(sizeof(Thread))) Thread(
        Thread::Configuration(Thread::READY,
            Thread::IDLE), &Thread::idle);

    _timer = new (kmallocc(sizeof(Scheduler_Timer)))
        Scheduler_Timer(QUANTUM, time_slicer);

    This_Thread::not_booting();
}

class Init_First
{
public:
    Init_First() {
        db<Init>(TRC) << "Init_First()" << endl;
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }
        Thread::running()->_context->load();
    }
};
```

```

int Thread::idle()
{
    while(_thread_count > 1) {
        if(Traits<Thread>::trace_idle)
            db<Thread>(TRC) << "Thread::idle(this="
                << running() << ")" << endl;
        CPU::int_enable();
        CPU::halt();
    }

    CPU::int_disable();
    db<Thread>(WRN)
        << "The last thread has exited!" << endl;
    if(reboot) {
        db<Thread>(WRN)
            << "Rebooting the machine ..." << endl;
        Machine::reboot();
    } else {
        db<Thread>(WRN)
            << "Halting the machine ..." << endl;
        CPU::halt();
    }

    return 0;
}

extern "C" { void __epos_app_entry(); }

```

```

void Thread::init()
{
    db<Init,Thread>(TRC) << "Thread::init()" << endl;

    Thread::_running =
        new (kmallocc(sizeof(Thread))) Thread(
            Thread::Configuration(Thread::RUNNING,
                Thread::MAIN), __epos_app_entry));

    new (kmallocc(sizeof(Thread))) Thread(
        Thread::Configuration(Thread::READY,
            Thread::IDLE), &Thread::idle);

    _timer = new (kmallocc(sizeof(Scheduler_Timer)))
        Scheduler_Timer(QUANTUM, time_slicer);

    This_Thread::not_booting();
}

class Init_First
{
public:
    Init_First() {
        db<Init>(TRC) << "Init_First()" << endl;
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::running()->_context->load();
    }
};

```

```

int Thread::idle()
{
    while(_thread_count > 1) {
        if(Traits<Thread>::trace_idle)
            db<Thread>(TRC) << "Thread::idle(this="
                << running() << ")" << endl;
        CPU::int_enable();
        CPU::halt();
    }
}

```

Who holds a pointer to idle?

```

        << "The last thread has exited!" << endl;
    if(reboot) {
        db<Thread>(WRN)
            << "Rebooting the machine ..." << endl;
        Machine::reboot();
    } else {
        db<Thread>(WRN)
            << "Halting the machine ..." << endl;
        CPU::halt();
    }

    return 0;
}

extern "C" { void __epos_app_entry(); }

```

```

void Thread::init()
{
    db<Init,Thread>(TRC) << "Thread::init()" << endl;

    Thread::_running =
        new (kmallocc(sizeof(Thread))) Thread(
            Thread::Configuration(Thread::RUNNING,
                Thread::MAIN), __epos_app_entry));

    new (kmallocc(sizeof(Thread))) Thread(
        Thread::Configuration(Thread::READY,
            Thread::IDLE), &Thread::idle);


    _timer = new (kmallocc(sizeof(Scheduler_Timer)))
        Scheduler_Timer(QUANTUM, time_slicer);

    This_Thread::not_booting();
}

class Init_First
{
public:
    Init_First() {
        db<Init>(TRC) << "Init_First()" << endl;
        if(!Traits<System>::multithread) {
            CPU::int_enable();
            return;
        }

        Thread::running()->_context->load();
    }
};

```




# Embedded Parallel Operating System

**Coding Journey through OS Design  
–from co-routines to a multicore kernel–**

**Prof. Antônio Augusto Fröhlich, Ph.D.**

**UFSC / LISHA  
September 30, 2020**





# Blocking Synchronization

We will now start to change the OS so it does **blocking synchronization** instead of co-routines.

Let's replace `yield()` calls for waiting queues operations.

Let's first focus on the Synchronizers and then on `Thread::join()`.

Let's leave `Alarm::delay()` for later on!

```

class Synchronizer_Common
{
protected:
    typedef Thread::Queue Queue;

protected:
    Synchronizer_Common() {}
    ~Synchronizer_Common() {
        begin_atomic();
        wakeup_all();
    }

    // Atomic operations
    bool tsl(volatile bool & lock) {
        return CPU::tsl(lock); }
    int finc(volatile int & number) {
        return CPU::finc(number); }
    int fdec(volatile int & number) {
        return CPU::fdec(number); }

    // Thread operations
    void begin_atomic() { Thread::lock(); }
    void end_atomic() { Thread::unlock(); }

    void sleep() { Thread::sleep(&_queue); }
    void wakeup() { Thread::wakeup(&_queue); }
    void wakeup_all() {
        Thread::wakeup_all(&_queue); }

protected:
    Queue _queue;
};

```

```

class Thread
{
    :
protected:
    static void sleep(Queue * q);
    static void wakeup(Queue * q);
    static void wakeup_all(Queue * q);


    :


protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Queue::Element _link;

    :

};

```





```
class Synchronizer_Common
{
protected:
    typedef Thread::Queue Queue;

protected:
    Synchronizer_Common() {}
    ~Synchronizer_Common() {
        begin_atomic();
        wakeup_all();
    }

    // Atomic operations
    bool tsl(volatile bool & lock) {
        return CPU::tsl(lock); }
    int finc(volatile int & number) {
        return CPU::finc(number); }
    int fdec(volatile int & number) {
        return CPU::fdec(number); }

    // Thread operations
    void begin_atomic() { Thread::lock(); }
    void end_atomic() { Thread::unlock(); }

    void sleep() { Thread::sleep(&_amp;queue); }
    void wakeup() { Thread::wakeup(&_amp;queue); }
    void wakeup_all() {
        Thread::wakeup_all(&_amp;queue); }

protected:
    Queue _queue;
};
```

```
import Queue from Thread,
    honoring ordering
```


```
protected:
    static void sleep(Queue * q);
    static void wakeup(Queue * q);
    static void wakeup_all(Queue * q);

    :

protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Queue::Element _link;

    :

};
```



```
class Synchronizer_Common
{
protected:
    typedef Thread::Queue Queue;

protected:
    Synchronizer_Common() {}
    ~Synchronizer_Common() {
        begin_atomic();
        wakeup_all();
    }

    // Atomic operations
    bool tsl(volatile bool & lock) {
        return CPU::tsl(lock); }
    int finc(volatile int & number) {
        return CPU::finc(number); }
    int fdec(volatile int & number) {
        return CPU::fdec(number); }

    // Thread operations
    void begin_atomic() { Thread::lock(); }
    void end_atomic() { Thread::unlock(); }

    void sleep() { Thread::sleep(&_queue); }
    void wakeup() { Thread::wakeup(&_queue); }
    void wakeup_all() {
        Thread::wakeup_all(&_queue); }


protected:
    Queue _queue;
};
```

```
class Thread
{
    :
protected:
    static void sleep(Queue * q);
    static void wakeup(Queue * q);
    static void wakeup_all(Queue * q);

protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Queue::Element _link;

    :
};
```

release blocked threads if deleted



```
class Synchronizer_Common
{
protected:
    typedef Thread::Queue Queue;

protected:
    Synchronizer_Common() {}
    ~Synchronizer_Common() {
        begin_atomic();
        wakeup_all();
    }

    // Atomic operations
    bool tsl(volatile bool & lock) {
        return CPU::tsl(lock); }
    int finc(volatile int & number) {
        return CPU::finc(number); }
    int fdec(volatile int & number) {
        return CPU::fdec(number); }

    // Thread operations
    void begin_atomic() { Thread::lock(); }
    void end_atomic() { Thread::unlock(); }

    void sleep() { Thread::sleep(&_queue); }
    void wakeup() { Thread::wakeup(&_queue); }
    void wakeup_all() {
        Thread::wakeup_all(&_queue); }

protected:
    Queue _queue;
};
```


```
class Thread
{
    :
protected:
    static void sleep(Queue * q);
    static void wakeup(Queue * q);
    static void wakeup_all(Queue * q);

    :

protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Queue::Element _link;

    :
};
```

forward blocking operations to  
Thread



```
class Synchronizer_Common
{
protected:
    typedef Thread::Queue Queue;

protected:
    Synchronizer_Common() {}
    ~Synchronizer_Common() {
        begin_atomic();
        wakeup_all();
    }

    // Atomic operations
    bool tsl(volatile bool & lock) {
        return CPU::tsl(lock); }
    int finc(volatile int & number) {
        return CPU::finc(number); }
    int fdec(volatile int & number) {
        return CPU::fdec(number); }

    // Thread operations
    void begin_atomic() { Thread::lock(); }
    void end_atomic() { Thread::unlock(); }

    void sleep() { Thread::sleep(&_queue); }
    void wakeup() { Thread::wakeup(&_queue); }
    void wakeup_all() {
        Thread::wakeup_all(&_queue); }

protected:
    Queue _queue;
};
```


```
class Thread
{
    :
protected:
    static void sleep(Queue * q);
    static void wakeup(Queue * q);
    static void wakeup_all(Queue * q);

    :

protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Queue::Element _link;

    :
};
```

a queue of blocked Threads



```
class Synchronizer_Common
{
protected:
    typedef Thread::Queue Queue;

    protected:
    Synchronizer_Common() {
        begin_atomic();
        wakeup_all();
    }

    // Atomic operations
    bool tsl(volatile bool & lock) {
        return CPU::tsl(lock); }
    int finc(volatile int & number) {
        return CPU::finc(number); }
    int fdec(volatile int & number) {
        return CPU::fdec(number); }

    // Thread operations
    void begin_atomic() { Thread::lock(); }
    void end_atomic() { Thread::unlock(); }

    void sleep() { Thread::sleep(&_queue); }
    void wakeup() { Thread::wakeup(&_queue); }
    void wakeup_all() {
        Thread::wakeup_all(&_queue); }

protected:
    Queue _queue;
};
```

takes a reference to a  
waiting queue


```
class Thread
{
    :
protected:
    static void sleep(Queue * q);
    static void wakeup(Queue * q);
    static void wakeup_all(Queue * q);

    :

protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Queue::Element _link;

    :

};
```



```
class Synchronizer_Common
{
protected:
    typedef Thread::Queue Queue;

protected:
    Synchronizer_Common() {}
    ~Synchronizer_Common() {
        begin_atomic();
        wakeup_all();
    }

    // Atomic operations
    bool finc(volatile int & number) {
        return CPU::finc(number); }
    int fdec(volatile int & number) {
        return CPU::fdec(number); }

    // Thread operations
    void begin_atomic() { Thread::lock(); }
    void end_atomic() { Thread::unlock(); }

    void sleep() { Thread::sleep(&_queue); }
    void wakeup() { Thread::wakeup(&_queue); }
    void wakeup_all() {
        Thread::wakeup_all(&_queue); }

protected:
    Queue _queue;
};
```

a reference to the queue currently  
waiting on

```
class Thread
{
    :
protected:
    static void sleep(Queue * q);
    static void wakeup(Queue * q);
    static void wakeup_all(Queue * q);

    :

protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Queue::Element _link;

    :

};
```



```

void Thread::sleep(Queue * q)
{
    db<Thread>(TRC) << "Thread::sleep(running="
                << running() << ",q="
                << q << ")" << endl;

    assert(locked());

    while(_ready.empty())
        idle();

    Thread * prev = running();
    prev->_state = WAITING;
    prev->_waiting = q;
    q->insert(&prev->_link);

    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(prev, _running);
    unlock();
}

void Thread::wakeup(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup(running="
                << running() << ",q="
                << q << ")" << endl;

```

```

assert(locked());

    if(!q->empty()) {
        Thread * t = q->remove()->object();
        t->_state = READY;
        t->_waiting = 0;
        _ready.insert(&t->_link);
    }
    unlock();
}

```

```


void Thread::wakeup_all(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup_all(running="
                << running()
                << ",q=" << q << ")" << endl;


    assert(locked());

    while(!q->empty()) {
        Thread * t = q->remove()->object();
        t->_state = READY;
        t->_waiting = 0;
        _ready.insert(&t->_link);
    }

    unlock();
}

```





```
void Thread::sleep(Queue * q)
{
    db<Thread>(TRC) << "Thread::sleep(running="
                  << running() << ",q="
                  << q << ")" << endl;

    assert(locked());

    while(_ready.empty())
        idle();

    Thread * prev = running();
    prev->_state = WAITING;
    prev->_waiting = q;
    q->insert(&prev->_link);

    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(prev, _running);
    unlock();
}

void Thread::wakeup(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup(running="
                  << running() << ",q="
                  << q << ")" << endl;
```

callers must have locked before, thus  
ensuring atomicity

```
Thread * t = q->remove()->object();
t->_state = READY;
t->_waiting = 0;
_ready.insert(&t->_link);
}
unlock();
}

void Thread::wakeup_all(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup_all(running="
                  << running()
                  << ",q=" << q << ")" << endl;

    assert(locked());

    while(!q->empty()) {
        Thread * t = q->remove()->object();
        t->_state = READY;
        t->_waiting = 0;
        _ready.insert(&t->_link);
    }

    unlock();
}
```

```

void Thread::sleep(Queue * q)
{
    db<Thread>(TRC) << "Thread::sleep(running="
                << running() << ",q="
                << q << ")" << endl;

    assert(locked());

    while(_ready.empty())
        idle();

    Thread * prev = running();
    prev->_state = WAITING;
    prev->_waiting = q;
    q->insert(&prev->_link);

    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(prev, _running);
    unlock();
}

void Thread::wakeup(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup(running="
                << running() << ",q="
                << q << ")" << endl;

```

## #assert

```

#ifdef NDEBUG
#define assert(condition) ((void)0)
#else
#define assert(condition) /*implementation defined*/
#endif

```

If NDEBUG is defined then assert is a **model checking construct**

If NDEBUG is not defined, then assert checks the condition and outputs implementation-specific diagnostic information if it does not hold. In POSIX, it prints on the standard error output and calls std::abort.

```

    assert(locked());

    while(!q->empty()) {
        Thread * t = q->remove()->object();
        t->_state = READY;
        t->_waiting = 0;
        _ready.insert(&t->_link);
    }

    unlock();
}

```

```

void Thread::sleep(Queue * q)
{
    db<Thread>(TRC) << "Thread::sleep(running="
                << running() << ",q="
                << q << ")" << endl;

    assert(locked());

    while(!_ready.empty())
        idle();

    Thread * prev = running();
    prev->_state = WAITING;
    prev->_waiting = q;
    q->insert(&prev->_link);

    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(prev, _running);
    unlock();
}

void Thread::wakeup(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup(running="
                << running() << ",q="
                << q << ")" << endl;

```

```

    assert(locked());

    if(!q->empty()) {
        Thread * t = q->remove()->object();
        t->_state = READY;
        t->_waiting = 0;
    }

```

similar to suspend(), but holding a  
reference to the queue

```


void Thread::wakeup_all(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup_all(running="
                << running()
                << ",q=" << q << ")" << endl;

    assert(locked());

    while(!q->empty()) {
        Thread * t = q->remove()->object();
        t->_state = READY;
        t->_waiting = 0;
        _ready.insert(&t->_link);
    }

    unlock();
}

```



```

void Thread::wakeup_all(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup_all(running="
    << running() << ",q="
    << q << ")" << endl;

    assert(locked());

    while(!_ready.empty())
        idle();

    Thread * prev = running();
    prev->_state = WAITING;
    prev->_waiting = q;
    q->insert(&prev->_link);

    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(prev, _running);
    unlock();
}

void Thread::wakeup(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup(running="
    << running() << ",q="
    << q << ")" << endl;

```

similar to resume(), inserting  
threads in the READY queue

```

assert(locked());

if(!q->empty()) {
    Thread * t = q->remove()->object();
    t->_state = READY;
    t->_waiting = 0;
    _ready.insert(&t->_link);
}
unlock();
}


void Thread::wakeup_all(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup_all(running="
    << running()
    << ",q=" << q << ")" << endl;

    assert(locked());

    while(!q->empty()) {
        Thread * t = q->remove()->object();
        t->_state = READY;
        t->_waiting = 0;
        _ready.insert(&t->_link);
    }

    unlock();
}

```



```
void Thread::sleep(Queue * q)
{
    db<Thread>(TRC) << "Thread::sleep(running="
                << running() << ",q="
                << q << ")" << endl;
```

```
    assert(locked());
```

```
    while(!_ready.empty())
        idle();
```

```
    Thread * prev = running();
```

releases several threads at once (for  
Condition::broadcast())

```
    _running = _ready.remove()->object();
    _running->_state = RUNNING;
```

```
    dispatch(prev, _running);
    unlock();
}
```

```
void Thread::wakeup(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup(running="
                << running() << ",q="
                << q << ")" << endl;
```

```
    assert(locked());
```


```
    if(!q->empty()) {
        Thread * t = q->remove()->object();
        t->_state = READY;
        t->_waiting = 0;
        _ready.insert(&t->_link);
    }
    unlock();
}
```


```
void Thread::wakeup_all(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup_all(running="
                << running()
                << ",q=" << q << ")" << endl;
```

assert(locked());

```
    while(!q->empty()) {
        Thread * t = q->remove()->object();
        t->_state = READY;
        t->_waiting = 0;
        _ready.insert(&t->_link);
    }

    unlock();
}
```





```
void Thread::sleep(Queue * q)
{
    db<Thread>(TRC) << "Thread::sleep(running="
                << running() << ",q="
                << q << ")" << endl;

    assert(locked());

    while(_ready.empty())
        idle();

    Thread * prev = running();
    prev->_state = WAITING;
    prev->_waiting = q;
    q->insert(&prev->_link);

    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(prev, _running);
    unlock();
}

void Thread::wakeup(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup(running="
                << running() << ",q="
                << q << ")" << endl;
```

```
assert(locked());

if(!q->empty()) {
    Thread * t = q->remove()->object();
    t->_state = READY;
    t->_waiting = 0;
    _ready.insert(&t->_link);
}
unlock();
}
```

```
void Thread::wakeup_all(Queue * q)
{
    db<Thread>(TRC) << "Thread::wakeup_all(running="
                << running() << ",q="
                << q << ")" << endl;

    assert(locked());

    while(!q->empty()) {
        Thread * t = q->remove()->object();
        t->_state = READY;
        t->_waiting = 0;
        _ready.insert(&t->_link);
    }

    unlock();
}
```

there are many queue operations here that  
can get involved in race conditions!

```

void Mutex::lock()
{
    db<Synchronizer>(TRC) << "Mutex::lock(this="
                      << this << ")" << endl;

    begin_atomic();
    if(tsl(_locked))
        sleep(); // implicit end_atomic()
    else
        end_atomic();
}

```

```

void Mutex::unlock()
{
    db<Synchronizer>(TRC) << "Mutex::unlock(this="
                      << this << ")" << endl;

    begin_atomic();
    if(_queue.empty()) {
        _locked = false;
        end_atomic();
    } else
        wakeup(); // implicit end_atomic()
}

```

```

void Semaphore::p()
{
    db<Synchronizer>(TRC) << "Semaphore::p(this="
                      << this << ",value="
                      << _value << ")" << endl;

    begin_atomic();
    if(fdec(_value) < 1)
        sleep(); // implicit end_atomic()
    else
        end_atomic();
}


```

```

void Semaphore::v()
{
    db<Synchronizer>(TRC) << "Semaphore::v(this="
                      << this << ",value="
                      << _value << ")" << endl;

    begin_atomic();
    if(finc(_value) < 0)
        wakeup(); // implicit end_atomic()
    else
        end_atomic();
}

```





```

void Mutex::lock()
{
    db<Synchronizer>(TRC) << "Mutex::lock(this="
                      << this << ")" << endl;

    begin_atomic();
    if(tsl(_locked))
        sleep(); // implicit end_atomic()
    else
        end_atomic();
}

```

```

void Mutex::unlock()
{
    db<Synchronizer>(TRC) << "Mutex::unlock(this="
                      << this << ")" << endl;

    begin_atomic();
    if(_queue.empty()) {
        _locked = false;
        end_atomic();
    } else
        wakeup(); // implicit end_atomic()
}

```

loop removed

```

    re::p(this
    << this << ",value="
    << _value << ")" << endl;

    begin_atomic();
    if(fdec(_value) < 1)
        sleep(); // implicit end_atomic()
    else
        end_atomic();
}


```

```

void Semaphore::v()
{
    db<Synchronizer>(TRC) << "Semaphore::v(this="
                      << this << ",value="
                      << _value << ")" << endl;

    begin_atomic();
    if(finc(_value) < 0)
        wakeup(); // implicit end_atomic()
    else
        end_atomic();
}

```



```

void Mutex::lock()
{
    db<Synchronizer>(TRC) << "Mutex::lock(this="
                      << this << ")" << endl;

    begin_atomic();
    if(tsl(_locked))
        sleep(); // implicit end_atomic()
    else
        end_atomic();
}

```

```

void Mutex::unlock()
{
    db<Synchronizer>(TRC) << "Mutex::unlock(this="
                      << this << ")" << endl;

    begin_atomic();
    if(_queue.empty()) {
        _locked = false;
        end_atomic();
    } else
        wakeup(); // implicit end_atomic()
}


```

```

void Semaphore::p()
{
    db<Synchronizer>(TRC) << "Semaphore::p(this="
                      << this << ",value="
                      << _value << ")" << endl;

    begin_atomic();
    if(fdec(_value) < 1)
        sleep(); // implicit end_atomic()
    else
        end_atomic();
}

```



only unlock when the queue is empty,  
thus preventing race conditions

```

{
    db<Synchronizer>(TRC) << "Semaphore::v(this="
                      << this << ",value="
                      << _value << ")" << endl;

    begin_atomic();
    if(finc(_value) < 0)
        wakeup(); // implicit end_atomic()
    else
        end_atomic();
}

```

```

void Mutex::lock()
{
    db<Synchronizer>(TRC) << "Mutex::lock(this="
    loop removed
    ;

    begin_atomic();
    if(tsl(_locked))
        sleep(); // implicit end_atomic()
    else
        end_atomic();
}

```

```

void Mutex::unlock()
{
    db<Synchronizer>(TRC) << "Mutex::unlock(this="
    << this << ")" << endl;

    begin_atomic();
    if(_queue.empty()) {
        _locked = false;
        end_atomic();
    } else
        wakeup(); // implicit end_atomic()
}

```

```

void Semaphore::p()
{
    db<Synchronizer>(TRC) << "Semaphore::p(this="
    << this << ",value="
    << _value << ")" << endl;

    begin_atomic();
    if(fdec(_value) < 1)
        sleep(); // implicit end_atomic()
    else
        end_atomic();
}


```

```

void Semaphore::v()
{
    db<Synchronizer>(TRC) << "Semaphore::v(this="
    << this << ",value="
    << _value << ")" << endl;

    begin_atomic();
    if(finc(_value) < 0)
        wakeup(); // implicit end_atomic()
    else
        end_atomic();
}

```



```

void Mutex::lock()
{
    db<Synchronizer>(TRC) << "Mutex::lock(this="
                      << this << ")" << endl;

    begin_atomic();
    if(tsl(_locked))
        )
}

```

ensures ordering

```

void Mutex::unlock()
{
    db<Synchronizer>(TRC) << "Mutex::unlock(this="
                      << this << ")" << endl;

    begin_atomic();
    if(_queue.empty()) {
        _locked = false;
        end_atomic();
    } else
        wakeup(); // implicit end_atomic()
}


```

```

void Semaphore::p()
{
    db<Synchronizer>(TRC) << "Semaphore::p(this="
                      << this << ",value="
                      << _value << ")" << endl;

    begin_atomic();
    if(fdec(_value) < 1)
        sleep(); // implicit end_atomic()
    else
        end_atomic();
}

```



```

void Semaphore::v()
{
    db<Synchronizer>(TRC) << "Semaphore::v(this="
                      << this << ",value="
                      << _value << ")" << endl;

    begin_atomic();
    if(finc(_value) < 0)
        wakeup(); // implicit end_atomic()
    else
        end_atomic();
}

```

```

void Mutex::lock()
{
    db<Synchronizer>(TRC) << "Mutex::lock(this="
                        << this << ")" << endl;

    begin_atomic();
    if(tsl(_locked))
        sleep(); // implicit end_atomic()
    else
        end_atomic();
}

```

```

void Mutex::unlock()
{
    db<Synchroniz
    compares finc()'s result
    (not _value)

    begin_atomic();
    if(_queue.empty()) {
        _locked = false;
        end_atomic();
    } else
        wakeup(); // implicit end_atomic()
}

```

```

void Semaphore::p()
{
    db<Synchronizer>(TRC) << "Semaphore::p(this="
                        << this << ",value="
                        << _value << ")" << endl;

    begin_atomic();
    if(fdec(_value) < 1)
        sleep(); // implicit end_atomic()
    else
        end_atomic();
}


```

```

void Semaphore::v()
{
    db<Synchronizer>(TRC) << "Semaphore::v(this="
                        << this << ",value="
                        << _value << ")" << endl;

    begin_atomic();
    if(finc(_value) < 0)
        wakeup(); // implicit end_atomic()
    else
        end_atomic();
}

```



```

void Mutex::lock()
{
    db<Synchronizer>(TRC) << "Mutex::lock(this="
                    << this << ")" << endl;

    begin_atomic();
    if(tsl(_locked))
        sleep(); // implicit end_atomic()
    else
        end_atomic();
}

```

```

void Mutex::unlock()
{
    db<Synchronizer>(TRC) << "Mutex::unlock(this="
                    << this << ")" << endl;

    begin_atomic();
    if(_queue.empty()) {
        _locked = false;
        end_atomic();
    } else
        w
}

```

are CPU::finc() and Thread:lock() enough to ensure queues won't be corrupted?

```

void Semaphore::p()
{
    db<Synchronizer>(TRC) << "Semaphore::p(this="
                    << this << ",value="
                    << _value << ")" << endl;

    begin_atomic();
    if(fdec(_value) < 1)
        sleep(); // implicit end_atomic()
    else
        end_atomic();
}


```


```

void Semaphore::v()
{
    db<Synchronizer>(TRC) << "Semaphore::v(this="
                    << this << ",value="
                    << _value << ")" << endl;

    begin_atomic();
    if(finc(_value) < 0)
        wakeup(); // implicit end_atomic()
    else
        end_atomic();
}

```





```
class Thread
{
    :

protected:
    static void sleep(Queue * q);
    static void wakeup(Queue * q);
    static void wakeup_all(Queue * q);

    :

protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Thread * volatile _joining;
    Queue::Element _link;

    :

};

template<typename ... Tn>
inline Thread::Thread(int (* entry)(Tn ...), Tn ... an)
: _state(READY), _waiting(0), _joining(0),
  _link(this, NORMAL) {
    constructor_prologue(STACK_SIZE);
    _context = CPU::init_stack(0, _stack + STACK_SIZE,
                               &__exit, entry, an ...);

    constructor_epilogue(entry, STACK_SIZE);
}
```

```
Thread::~Thread()
{
    lock();

    db<Thread>(TRC) << "~Thread(this=" << this
                   << ",state=" << _state
                   << ",priority="
                   << _link.rank()
                   << ",stack={b="
                   << reinterpret_cast<void *>
                       (_stack)
                   << ",context={b=" << _context
                   << ", " << *_context << "})"
                   << endl;


    _ready.remove(this);
    _suspended.remove(this);

    if(_waiting)
        _waiting->remove(this);

    if(_joining)
        _joining->resume();

    unlock();

    kfree(_stack);
}
```



```
class Thread
{
    :

protected:
    static void sleep(Queue * q);
    static void wakeup(Queue * q);
    static void wakeup_all(Queue * q);

    :

protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Thread * volatile _joining;
    Queue::Element _link;

    :

};

template<typename ... Tn>
inline Thread::Thread(int (* entry)(Tn ...), Tn ... an)
: _state(READY), _waiting(0), _joining(0),
  _link(this, NORMAL) {
    constructor_prologue(STACK_SIZE);
    _context = CPU::init_stack(0, _stack + STACK_SIZE,
                               &__exit, entry, an ...);

    constructor_epilogue(entry, STACK_SIZE);
}
```

```
Thread::~~Thread()
{
```

```
    lock();
```

adding a pointer to an eventual joiner (a single one; no queue)

```
    this=" << this
    << _state
    y="
    << _link.rank()
    << ",stack={b="
    << reinterpret_cast<void *>
        (_stack)
    << ",context={b=" << _context
    << ", " << *_context << "})"
    << endl;

    _ready.remove(this);
    _suspended.remove(this);


    if(_waiting)
        _waiting->remove(this);

    if(_joining)
        _joining->resume();

    unlock();

    kfree(_stack);
}
```





```
class Thread
{
    :

protected:
    static void sleep(Queue * q);
    static void wakeup(Queue * q);
    static void wakeup_all(Queue * q);

    :

protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Thread * volatile _joining;
    Queue::Element _link;

    :

};

template<typename ... Tn>
inline Thread::Thread(int (* entry)(Tn ...), Tn ... an)
: _state(READY), _waiting(0), _joining(0),
  _link(this, NORMAL) {
    constructor_prologue(STACK_SIZE);
    _context = CPU::init_stack(0, _stack + STACK_SIZE,
                              &__exit, entry, an ...);

    constructor_epilogue(entry, STACK_SIZE);
}
```

```
Thread::~Thread()
{
    lock();

    db<Thread>(TRC) << "~Thread(this=" << this
                  << ",state=" << _state
                  << ",priority="
                  << _link.rank()
                  << ",stack={b="
                  << cast<void *>
                  << _context
                  << ",_context=" << _context << "})"
                  << endl;

    _ready.remove(this);
    _suspended.remove(this);


    if(_waiting)
        _waiting->remove(this);

    if(_joining)
        _joining->resume();

    unlock();

    kfree(_stack);
}
```

initialize pointers with null (0)



```
class Thread
{
    :
protected:
    static void sleep(Queue * q);
    static void wakeup(Queue * q);
    static void wakeup_all(Queue * q);
    :
```

```
protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
```

if the deleted thread was waiting in a Synchronizer's queue, remove it

```
};

template<typename ... Tn>
inline Thread::Thread(int (* entry)(Tn ...), Tn ... an)
: _state(READY), _waiting(0), _joining(0),
  _link(this, NORMAL) {
    constructor_prologue(STACK_SIZE);
    _context = CPU::init_stack(0, _stack + STACK_SIZE,
                              &__exit, entry, an ...);

    constructor_epilogue(entry, STACK_SIZE);
}
```

```
Thread::~Thread()
{
    lock();

    db<Thread>(TRC) << "~Thread(this=" << this
                   << ",state=" << _state
                   << ",priority="
                   << _link.rank()
                   << ",stack={b="
                   << reinterpret_cast<void *>
                       (_stack)
                   << ",context={b=" << _context
                   << ", " << *_context << "})"
                   << endl;


    _ready.remove(this);
    _suspended.remove(this);

    if(_waiting)
        _waiting->remove(this);

    if(_joining)
        _joining->resume();

    unlock();

    kfree(_stack);
}
```



```
class Thread
{
    :

protected:
    static void sleep(Queue * q);
    static void wakeup(Queue * q);
    static void wakeup_all(Queue * q);

    :

protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Thread * volatile _joining;
    Queue::Element _link;

    :

};
```

```
template<...>
inline void Thread::... (
    ...), Tn ... an)
{
    _link(this, NORMAL) {
        constructor_prologue(STACK_SIZE);
        _context = CPU::init_stack(0, _stack + STACK_SIZE,
                                   &__exit, entry, an ...);

        constructor_epilogue(entry, STACK_SIZE);
    }
}
```

if someone was waiting for this thread termination, resume it

```
Thread::~~Thread()
{
    lock();

    db<Thread>(TRC) << "~Thread(this=" << this
                   << ",state=" << _state
                   << ",priority="
                   << _link.rank()
                   << ",stack={b="
                   << reinterpret_cast<void *>
                       (_stack)
                   << ",context={b=" << _context
                   << ", " << *_context << "}")
                   << endl;


    _ready.remove(this);
    _suspended.remove(this);

    if(_waiting)
        _waiting->remove(this);

    if(_joining)
        _joining->resume();

    unlock();

    kfree(_stack);
}
```



```
class Thread
{
    :

protected:
    static void sleep(Queue * q);
    static void wakeup(Queue * q);
    static void wakeup_all(Queue * q);

    :

protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Thread * volatile _joining;
    Queue::Element _link;

    :

};
```

```
template<...>
inline void Thread::Thread(...) {
    if someone was waiting for this
    thread termination, resume it
    ...), Tn ... an)
    0),
    _link(this, NORMAL) {
        constructor_prologue(STACK_SIZE);
        _context = CPU::init_stack(0, _stack + STACK_SIZE,
                                   &__exit, entry, an ...);

        constructor_epilogue(entry, STACK_SIZE);
    }
```

```
Thread::~Thread()
{
    lock();

    db<Thread>(TRC) << "~Thread(this=" << this
                   << ",state=" << _state
                   << ",priority="
                   << _link.rank()
                   << ",stack={b="
                   << reinterpret_cast<void *>
                       (_stack)
                   << ",context={b=" << _context
                   << ", " << *_context << "}")
                   << endl;

    _ready.remove(this);
    _suspended.remove(this);

    if(_waiting)
        _waiting->remove(this);

    if(_joining)
        _joining->resume();

    unlock();

    kfree(_stack);
}
```

```
class Thread
```

```
{
    :
protected:
    static void sleep(Queue * q);
    static void wakeup(Queue * q);
    static void wakeup_all(Queue * q);
    :
protected:
    char * _stack;
    Context * volatile _context;
    volatile State _state;
    Queue * _waiting;
    Thread * volatile _joining;
    Queue::Element _link;
```

```
    :
};

template<...>
inline ... (Tn ... an)
: _sta ...
_link(this, NORMAL) {
    constructor_prologue(STACK_SIZE);
    _context = CPU::init_stack(0, _stack + STACK_SIZE,
                               &__exit, entry, an ...);
    constructor_epilogue(entry, STACK_SIZE);
}
```

if someone was waiting for this thread termination, resume it

## Process, Task and Thread

A **process** is a program in execution. It has an **address space**, mapping the memory where the program's **code** and **data** are stored, one or more **threads**, and a set of OS **resources**.

A **thread** is an execution flow of a process. It has a **state**, a **context**, and a **stack**. It also has some sort of **id**.

A task can be:

- The passive part of a process (i.e. address space, code and data);
- A synonym for thread.

```
_suspended.remove(this);
```

```
if(_waiting)
    _waiting->remove(this);
```

```
if(_joining)
    _joining->resume();
```

```
unlock();
```

```
kfree(_stack);
```

```
}
```

```

int Thread::join()
{
    lock();

    db<Thread>(TRC) << "Thread::join(this="
                    << this << ",state="
                    << _state << ")" << endl;

    // Precondition: no Thread::self()->join()
    assert(running() != this);

    // Precondition: a single joiner
    assert(!_joining);

    if(_state != FINISHING) {
        _joining = running();
        _joining->suspend();
    } else
        unlock();

    return *reinterpret_cast<int *>(_stack);
}

```

```

void Thread::suspend()
{
    lock();

    db<Thread>(TRC) << "Thread::suspend(this="
                    << this << ")" << endl;

    if(_running != this)
        _ready.remove(this);

    _state = SUSPENDED;
    _suspended.insert(&_link);


    if(_running == this) {
        while(_ready.empty())
            idle();

        _running = _ready.remove()->object();
        _running->_state = RUNNING;

        dispatch(this, _running);
    }

    unlock();
}

```



```

int Thread::join()
{
    lock();

    db<Thread>(TRC) << "Thread::join(this="
                    << this << ",state="
                    << _state << ")" << endl;

    // Precondition: no Thread::self()->join()
    assert(running() != this);


    // Precondition: a single joiner
    assert(!_joining);

    if(_state != FINISHING) {
        _joining = running();
        _joining->suspend();
    } else
        unlock();

    return *reinterpret_cast<int *>(_stack);
}

```

precondition: a Thread cannot join itself



```

    db<Thread>(TRC) << "Thread::suspend(this="
                    << this << ")" << endl;

    if(_running != this)
        _ready.remove(this);

    _state = SUSPENDED;
    _suspended.insert(&_amp;link);

    if(_running == this) {
        while(_ready.empty())
            idle();

        _running = _ready.remove()->object();
        _running->_state = RUNNING;

        dispatch(this, _running);
    }

    unlock();
}

```

```

int Thread::join()
{
    lock();

    db<Thread>(TRC) << "Thread::join(this="
                    << this << ",state="
                    << _state << ")" << endl;

    // Precondition: no Thread::self()->join()
    assert(running() != this);

    // Precondition: a single joiner
    assert(!_joining);

    if(_state != FINISHING) {
        _joining = running();
        _joining->suspend();
    } else
        unlock();

    return *reinterpret_cast<int *>(_stack);
}

```

```

void Thread::suspend()
{
    lock();

    db<Thread>(TRC) << "Thread::suspend(this="
                    << this << ",state="
                    << _state << ")" << endl;

    precondition: only one joiner a time

    _ready.remove(this);

    _state = SUSPENDED;
    _suspended.insert(&_link);


    if(_running == this) {
        while(_ready.empty())
            idle();

        _running = _ready.remove()->object();
        _running->_state = RUNNING;

        dispatch(this, _running);
    }

    unlock();
}

```





```

int Thread::join()
{
    lock();

    db<Thread>(TRC) << "Thread::join(this="
                    << this << ",state="
                    << _state << ")" << endl;

    // Precondition: no Thread::self()->join()
    assert(running() != this);

    // Precondition: a single joiner
    assert(!_joining);

    if(_state != FINISHING) {
        _joining = running();
        _joining->suspend();
    } else
        unlock();

    return *reinterpret_cast<int *>(_stack);
}

```

```

void Thread::suspend()
{
    lock();

    db<Thread>(TRC) << "Thread::suspend(this="
                    << this << ")" << endl;

    if(_running != this)
        _ready.remove(this);

    state = SUSPENDED;

    if the Thread hasn't finished yet,
    suspend the caller


    while(_ready.empty())
        idle();

    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(this, _running);
}

unlock();
}

```



```

int Thread::join()
{
    lock();

    db<Thread>(TRC) << "Thread::join(this="
                    << this << ",state="
                    << _state << ")" << endl;

    // Precondition: no Thread::self()->join()
    assert(running() != this);

    // Precondition: a single joiner
    assert(!_joining);

    if(_state != FINISHING) {
        _joining = running();
        _joining->suspend();
    } else
        unlock();

    return *reinterpret_cast<int *>(_stack);
}

```

```

void Thread::suspend()
{
    lock();

    db<Thread>(TRC) << "Thread::suspend(this="
                    << this << ")" << endl;

    if(_running != this)
        _ready.remove(this);


    _state = SUSPENDED;
    _suspended.insert(&_link);

    if(_running == this) {
        while(_ready.empty())
            idle();
    }

    dispatch(this, _running);

    unlock();
}

```



returns to the caller what the Thread  
returned when finished

```

int Thread::join()
{
    lock();

    db<Thread>(TRC) << "Thread::join(this="
                    << this << ",state="
                    << _state << ")" << endl;

    // Precondition: no Thread::self()->join()
    assert(running() != this);

    // Precondition: a single joiner
    assert(!_joining);
    if(_state == RUNNING)
        _joining = running();
    _joining->suspend();
} else
    unlock();

return *reinterpret_cast<int *>(_stack);
}

```

wait for a Thread to become  
READY before actually suspending

```

void Thread::suspend()
{
    lock();

    db<Thread>(TRC) << "Thread::suspend(this="
                    << this << ")" << endl;

    if(_running != this)
        _ready.remove(this);

    _state = SUSPENDED;
    _suspended.insert(&_link);


    if(_running == this) {
        while(_ready.empty())
            idle();

        _running = _ready.remove()->object();
        _running->_state = RUNNING;

        dispatch(this, _running);
    }

    unlock();
}

```



```

void Thread::exit(int status)
{
    lock();

    db<Thread>(TRC) << "Thread::exit(status="
                    << status << ") [running="
                    << running() << "]" << endl;

    Thread * prev = _running;
    prev->_state = FINISHING;
    *reinterpret_cast<int *>(prev->_stack) =
        status;

    if(prev->_joining) {
        Thread * joining = prev->_joining;
        prev->_joining = 0;
        joining->resume(); // implicit unlock()
        lock();
    }
}

```


```

if(!_ready.empty()) {
    if(!_suspended.empty()) {
        while(_ready.empty())
            idle(); // implicit unlock();
        lock();
    } else {
        db<Thread>(WRN) << "The last thread in"
                        << "the system has exited!"
                        << endl;
        if(reboot) {
            db<Thread>(WRN)
                << "Rebooting the machine ..."
                << endl;
            Machine::reboot();
        } else {
            db<Thread>(WRN)
                << "Halting the CPU ..." << endl;
            CPU::halt();
        }
    }
} else {
    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(prev, _running);
}

unlock();
}

```



```

void Thread::exit(int status)
{
    lock();

    db<Thread>(TRC) << "Thread::exit(status="
                    << status << ") [running="
                    << running() << "]" << endl;

    Thread * prev = _running;
    prev->_state = FINISHING;
    *reinterpret_cast<int *>(prev->_stack) =
        status;

    if(prev->_joining) {
        Thread * joining = prev->_joining;
        prev->_joining = 0;
        joining->resume(); // implicit unlock()
        lock();
    }
}

```

set state to FINISHING and save the return value

```


if(!_ready.empty()) {
    if(!_suspended.empty()) {
        while(_ready.empty())
            it unlock();

        db<Thread>(WRN) << "The last thread in"
                        << "the system has exited!"
                        << endl;
        if(reboot) {
            db<Thread>(WRN)
                << "Rebooting the machine ..."
                << endl;
            Machine::reboot();
        } else {
            db<Thread>(WRN)
                << "Halting the CPU ..." << endl;
            CPU::halt();
        }
    }
} else {
    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(prev, _running);
}

unlock();
}

```



```

void Thread::exit(int status)
{
    lock();

    db<Thread>(TRC) << "Thread::exit(status="
                    << status << ") [running="
                    << running() << "]" << endl;

    Thread * prev = _running;
    prev->_state = FINISHING;
    *reinterpret_cast<int *>(prev->_stack) =
        status;

    if(prev->_joining) {
        Thread * joining = prev->_joining;
        prev->_joining = 0;
        joining->resume(); // implicit unlock()
        lock();
    }
}

```

if someone was waiting for this Thread  
to exit, wake him up



```

if(!_ready.empty()) {
    if(!_suspended.empty()) {
        while(_ready.empty())
            idle(); // implicit unlock();
        lock();
    } else {
        db<Thread>(WRN) << "The last thread in"
                        << "the system has exited!"
                        << endl;
        Machine::reboot();
    } else {
        db<Thread>(WRN)
            << "Halting the CPU ..." << endl;
        CPU::halt();
    }
} else {
    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(prev, _running);
}

unlock();
}

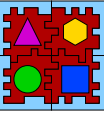
```

wait for a Thread to become  
READY before actually exiting

```
db<Thread>(TRC) << "Thread::exit(status="
    << status << ") [running="
    << running() << "]" << endl;
```

```
Thread * prev = _running;
prev->_state = FINISHING;
*reinterpret_cast<int *>(prev->_stack) =
    status;
```

```
if(prev->_joining) {
    Thread * joining = prev->_joining;
    prev->_joining = 0;
    joining->resume(); // implicit unlock()
    lock();
}
```



```
if(!_ready.empty()) {
    if(!_suspended.empty()) {
        while(_ready.empty())
            idle(); // implicit unlock();
        lock();
    } else {
        db<Thread>(WRN) << "The last thread in"
            << "the system has exited!"
            << endl;
        if(reboot) {
            db<Thread>(WRN)
                << "Rebooting the machine ..."
                << endl;
            Machine::reboot();
        } else {
            db<Thread>(WRN)
                << "Halting the CPU ..." << endl;
            CPU::halt();
        }
    }
} else {
    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(prev, _running);
}

unlock();
}
```

```
void Thread::exit(int status)
{
```

```
    lock();
```

```
    db<Thread>(TRC) << "Thread::exit(status="
                  << status << ") [running="
                  << endl;
```

if both READY and SUSPENDED queues  
are empty, then ...


```
    prev->_state = FINISHING;
    *reinterpret_cast<int *>(prev->_stack) =
        status;

    if(prev->_joining) {
        Thread * joining = prev->_joining;
        prev->_joining = 0;
        joining->resume(); // implicit unlock()
        lock();
    }
```


```
    if(!_ready.empty()) {
        if(!_suspended.empty()) {
            while(_ready.empty())
                idle(); // implicit unlock();
            lock();
        } else {
            db<Thread>(WRN) << "The last thread in"
                          << "the system has exited!"
                          << endl;
            if(reboot) {
                db<Thread>(WRN)
                    << "Rebooting the machine ..."
                    << endl;
                Machine::reboot();
            } else {
                db<Thread>(WRN)
                    << "Halting the CPU ..." << endl;
                CPU::halt();
            }
        }
    } else {
        _running = _ready.remove()->object();
        _running->_state = RUNNING;

        dispatch(prev, _running);
    }

    unlock();
}
```








# Embedded Parallel Operating System

**Coding Journey through OS Design  
–from co-routines to a multicore kernel–**

**Prof. Antônio Augusto Fröhlich, Ph.D.**

**UFSC / LISHA  
September 30, 2020**



# General Pointers

- Site

<https://epos.lisha.ufsc.br>

- Documentation


<https://epos.lisha.ufsc.br/EPOS+2+User+Guide>

- Git

<https://gitlab.lisha.ufsc.br/epos/epos/tree/master>

- Cross-compilers

<https://epos.lisha.ufsc.br/EPOS+Software>



# Co-routines: the minimal OS

“Co-routines are computer program components that generalize subroutines for non-preemptive multitasking, by allowing execution to be suspended and resumed. Co-routines are well-suited for implementing familiar program components such as cooperative tasks, exceptions, event loops, iterators, infinite lists and pipes.

According to Donald Knuth, Melvin Conway coined the term co-routine in 1958 when he applied it to construction of an assembly program.”

(Wikipedia)

```
// EPOS Semaphore Component Test Program
```

```
#include <machine/display.h>
#include <time.h>
#include <synchronizer.h>
#include <process.h>
```

```
using namespace EPOS;
```

```
const int iterations = 10;
```

```
Mutex table;
```

```
Thread * phil[5];
Semaphore * chopstick[5];
```

```
OStream cout;
```

```
int philosopher(int n, int l, int c)
{
    int first = (n < 4)? n : 0;
    int second = (n < 4)? n + 1 : 4;

    for(int i = iterations; i > 0; i--) {

        table.lock();
        Display::position(l, c);
        cout << "thinking";
        table.unlock();
```

```
        Delay thinking(2000000);
```

```
        // get first chopstick
        chopstick[first]->p();
        // get second chopstick
        chopstick[second]->p();
```

```
        table.lock();
        Display::position(l, c);
        cout << " eating ";
        table.unlock();
```

```
        Delay eating(1000000);
```


```
        // release first chopstick
        chopstick[first]->v();
        // release second chopstick
        chopstick[second]->v();
```

```
    }
```

```
    table.lock();
    Display::position(l, c);
    cout << " done ";
    table.unlock();
```

```
    return iterations;
```

```
}
```



## // EPOS Semaphore Component Test Program

```
#include <machine/display.h>
#include <time.h>
#include <synchronizer.h>
#include <process.h>

using namespace EPOS;

const int iterations = 10;

Mutex table;

Thread * phil[5];
Semaphore * chopstick[5];

OStream cout;

int philosopher(int n, int l, int c)
{
    int first = (n < 4)? n : 0;
    int second = (n < 4)? n + 1 : 4;

    for(int i = iterations; i > 0; i--) {

        table.lock();
        Display::position(l, c);
        cout << "thinking";
        table.unlock();
```

API inclusion

```
        delay_eating(2000000);

        chopstick[first]->p();
        // get second chopstick
        chopstick[second]->p();


        table.lock();
        Display::position(l, c);
        cout << "eating ";
        table.unlock();


        Delay eating(1000000);

        // release first chopstick
        chopstick[first]->v();
        // release second chopstick
        chopstick[second]->v();
    }

    table.lock();
    Display::position(l, c);
    cout << "done ";
    table.unlock();

    return iterations;
}
```





```
// EPOS Semaphore Component Test Program
```

```
#include <machine/display.h>
#include <time.h>
#include <synchronizer.h>
#include <process.h>
```

```
using namespace EPOS;
```

```
const int iterations = 10;
```

```
Mutex table;
```

```
Thread * phil[5];
Semaphore * chopstick[5];
```

```
OStream cout;
```

```
int philosopher(int n, int l, int c)
{
```

```
    int first = (n < 4)? n : 0;
    int second = (n < 4)? n + 1 : 4;
```

```
    for(int i = iterations; i > 0; i--) {
```

```
        table.lock();
        Display::position(l, c);
        cout << "thinking";
        table.unlock();
```

```
        Delay thinking(2000000);
```

```
        // get first chopstick
```

```
        chopstick[first]->p();
        // get second chopstick
        chopstick[second]->p();
```

```
        table.lock();
        Display::position(l, c);
        cout << " eating ";
        table.unlock();
```

```
        Delay eating(1000000);
```

```
        // release first chopstick
        chopstick[first]->v();
        // release second chopstick
        chopstick[second]->v();
```


```
    }
```

```
    table.lock();
    Display::position(l, c);
    cout << " done ";
    table.unlock();
```

```
    return iterations;
```

```
}
```

namespace import



```
// EPOS Semaphore Component Test Program
```

```
#include <machine/display.h>
#include <time.h>
#include <synchronizer.h>
#include <process.h>
```

```
using namespace EPOS;
```

```
const int iterations = 10;
```

```
Mutex table;
```

```
Thread * phil[5];
Semaphore * chopstick[5];
```

```
OStream cout;
```

```
int philosopher(int n, int l, int c)
{
    int first = (n < 4)? n : 0;
    int second = (n < 4)? n + 1 : 4;

    for(int i = iterations; i > 0; i--) {

        table.lock();
        Display::position(l, c);
        cout << "thinking";
        table.unlock();
```

```
        Delay thinking(2000000);
```

```
        // get first chopstick
        chopstick[first]->p();
        // get second chopstick
        chopstick[second]->p();
```

system objects can be statically allocated  
(constructor called before main())

```
        Delay eating(1000000);
```


```
        // release first chopstick
        chopstick[first]->v();
        // release second chopstick
        chopstick[second]->v();
```

```
    }
```

```
    table.lock();
    Display::position(l, c);
    cout << " done ";
    table.unlock();
```

```
    return iterations;
```

```
}
```



```
// EPOS Semaphore Component Test Program
```

```
#include <machine/display.h>
#include <time.h>
#include <synchronizer.h>
#include <process.h>
```

```
using namespace EPOS;
```

```
const int iterations = 10;
```

```
Mutex table;
```

```
Thread * phil[5];
Semaphore * chopstick[5];
```

```
OStream cout;
```

```
int philosopher(int n, int l, int c)
{
    int first = (n < 4)? n : 0;
    int second = (n < 4)? n + 1 : 4;

    for(int i = iterations; i > 0; i--) {

        table.lock();
        Display::position(l, c);
        cout << "thinking";
        table.unlock();
```

or dynamically

```
Delay thinking(2000000);
```

```
// get first chopstick
chopstick[first]->p();
// get second chopstick
chopstick[second]->p();
```

```
table.lock();
Display::position(l, c);
cout << " eating ";
table.unlock();
```

```
Delay eating(1000000);
```

```
// release first chopstick
chopstick[first]->v();
// release second chopstick
chopstick[second]->v();
```


```
}
```

```
table.lock();
Display::position(l, c);
cout << " done ";
table.unlock();
```

```
return iterations;
```

```
}
```





```
// EPOS Semaphore Component Test Program
```

```
#include <machine/display.h>
#include <time.h>
#include <synchronizer.h>
#include <process.h>
```

```
using namespace EPOS;
```

```
const int iterations = 10;
```

```
Mutex table;
```

```
Thread * phil[5];
Semaphore * chopstick[5];
```

```
OStream cout;
```

```
int philosopher(int n, int l, int c)
{
    int first = (n < 4)? n : 0;
    int second = (n < 4)? n + 1 : 4;

    for(int i = iterations; i > 0; i--) {

        table.lock();
        Display::position(l, c);
        cout << "thinking";
        table.unlock();
```

```
        Delay thinking(2000000);
```

```
        // get first chopstick
        chopstick[first]->p();
        // get second chopstick
        chopstick[second]->p();
```

```
        table.lock();
        Display::position(l, c);
        cout << " eating ";
        table.unlock();
```

```
        Delay eating(1000000);
```

```
        // release first chopstick
        chopstick[first]->v();
        // release second chopstick
        chopstick[second]->v();
```


```
    }
```

```
    table.lock();
    Display::position(l, c);
    cout << " done ";
    table.unlock();
```

```
    return iterations;
```

```
}
```

cout not created by default



```
// EPOS Semaphore Component Test Program
```

```
#include <machine/display.h>
#include <time.h>
#include <synchronizer.h>
#include <process.h>
```

```
using namespace EPOS;
```

```
const int iterations = 10;
```

```
Mutex table;
```

```
Thread * phil[5];
Semaphore * chopstick[5];
```

```
OStream cout;
```

```
int philosopher(int n, int l, int c)
{
    int first = (n < 4)? n : 0;
    int second = (n < 4)? n + 1 : 4;

    for(int i = iterations; i > 0; i--) {

        table.lock();
        Display::position(l, c);
        cout << "thinking";
        table.unlock();
```

```
        Delay thinking(2000000);
```

```
        // get first chopstick
        chopstick[first]->p();
        // get second chopstick
        chopstick[second]->p();
```

```
        table.lock();
        Display::position(l, c);
        cout << " eating ";
        table.unlock();
```

```
        Delay eating(1000000);
```

```
        // release first chopstick
        chopstick[first]->v();
```

any function can become a thread

```
        table.lock();
        Display::position(l, c);
        cout << " done ";
        table.unlock();
```

```
        return iterations;
```

```
    }
```

## Lamport's deadlock-free solution to the Dining Philosophers Problem




Figure by Dr. C.-K. Shene

```
int philosopher(int n, int l, int c)
{
    int first = (n < 4)? n : 0;
    int second = (n < 4)? n + 1 : 4;

    for(int i = iterations; i > 0; i--) {

        table.lock();
        Display::position(l, c);
        cout << "thinking";
        table.unlock();
```

```
Delay thinking(2000000);
```

```
// get first chopstick
chopstick[first]->p();
// get second chopstick
chopstick[second]->p();
```

```
table.lock();
Display::position(l, c);
cout << " eating ";
table.unlock();
```

```
Delay eating(1000000);
```


```
// release first chopstick
chopstick[first]->v();
// release second chopstick
chopstick[second]->v();
```

```
}
```

```
table.lock();
Display::position(l, c);
cout << " done ";
table.unlock();
```

```
return iterations;
```

```
}
```



## // EPOS Semaphore Component Test Program

```
#include <machine/display.h>
#include <time.h>
#include <synchronizer.h>
#include <process.h>
```

```
using namespace EPOS;
```

```
const int iterations = 10;
```

```
Mutex table;
```

```
Thread * phil[5];
Semaphore * chopstick[5];
```

```
OStream cout;
```

```
int philosopher(int n, int l, int c)
{
```

```
    int first = (n < 4)? n : 0;
    int second = (n < 4)? n + 1 : 4;
```

```
    for(int i = iterations; i > 0; i--) {
```

```
        table.lock();
        Display::position(l, c);
        cout << "thinking";
        table.unlock();
```

```
        Delay thinking(2000000);
```

```
        // get first chopstick
        chopstick[first]->p();
        // get second chopstick
        chopstick[second]->p();
```

```
        table.lock();
        Display::position(l, c);
        cout << " eating ";
        table.unlock();
```

```
        Delay eating(1000000);
```

```
        // release first chopstick
        chopstick[first]->v();
        // release second chopstick
        chopstick[second]->v();
```


```
    }
```

```
    table.lock();
    Display::position(l, c);
    cout << " done ";
    table.unlock();
```

```
    cout << " iterations";
```

```
}
```

atomic printing



## // EPOS Semaphore Component Test Program

```
#include <wait (in μs)>
#include <wait (in μs)>
#include <synchronizer.h>
#include <process.h>

using namespace EPOS;

const int iterations = 10;

Mutex table;

Thread * phil[5];
Semaphore * chopstick[5];

OStream cout;

int philosopher(int n, int l, int c)
{
    int first = (n < 4)? n : 0;
    int second = (n < 4)? n + 1 : 4;

    for(int i = iterations; i > 0; i--) {

        table.lock();
        Display::position(l, c);
        cout << "thinking";
        table.unlock();
```

```
Delay thinking(2000000);
```

```
// get first chopstick
chopstick[first]->p();
// get second chopstick
chopstick[second]->p();
```

```
table.lock();
Display::position(l, c);
cout << " eating ";
table.unlock();
```

```
Delay eating(1000000);
```


```
// release first chopstick
chopstick[first]->v();
// release second chopstick
chopstick[second]->v();
```


```
}
```

```
table.lock();
Display::position(l, c);
cout << " done ";
table.unlock();
```

```
return iterations;
```

```
}
```





```
// EPOS Semaphore Component Test Program
```

```
#include <machine/display.h>
#include <time.h>
#include <synchronizer.h>
#include <process.h>
```

```
using namespace EPOS;
```

```
const int iterations = 10;
```

```
Mutex table;
```

```
Thread * phil[5];
Semaphore * chopstick[5];
```

```
OStream cout;
```

```
int philosopher(int n, int l, int c)
```

```
{
    implicit exit() if a thread
    for(int i = iterations; i > 0; i--) {
```

```
        table.lock();
        Display::position(l, c);
        cout << "thinking";
        table.unlock();
```

```
        Delay thinking(2000000);
```

```
        // get first chopstick
        chopstick[first]->p();
        // get second chopstick
        chopstick[second]->p();
```

```
        table.lock();
        Display::position(l, c);
        cout << "eating ";
        table.unlock();
```

```
        Delay eating(1000000);
```


```
        // release first chopstick
        chopstick[first]->v();
        // release second chopstick
        chopstick[second]->v();
```

```
    }

    table.lock();
    Display::position(l, c);
    cout << "done ";
    table.unlock();
```

```
    return iterations;
```

```
}
```



```
int main()
{
    table.lock();
    Display::clear();
    Display::position(0, 0);
    cout << "The Philosopher's Dinner:" << endl;

    for(int i = 0; i < 5; i++)
        chopstick[i] = new Semaphore;

    phil[0] = new Thread(&philosopher, 0, 5, 32);
    phil[1] = new Thread(&philosopher, 1, 10, 44);
    phil[2] = new Thread(&philosopher, 2, 16, 39);
    phil[3] = new Thread(&philosopher, 3, 16, 24);
    phil[4] = new Thread(&philosopher, 4, 10, 20);

    cout << "Philosophers are alive and hungry!"
         << endl;

    Display::position(7, 44);
    cout << '/';
    Display::position(13, 44);
    cout << '\\';
    Display::position(16, 35);
    cout << '|';
    Display::position(13, 27);
    cout << '/';
    Display::position(7, 27);
    cout << '\\';
    Display::position(19, 0);
```


```
    cout << "The dinner is served ..." << endl;
    table.unlock();

    for(int i = 0; i < 5; i++) {
        int ret = phil[i]->join();
        table.lock();
        Display::position(20 + i, 0);
        cout << "Philosopher " << i << " ate "
             << ret << " times " << endl;
        table.unlock();
    }

    for(int i = 0; i < 5; i++)
        delete chopstick[i];
    for(int i = 0; i < 5; i++)
        delete phil[i];

    cout << "The end!" << endl;

    return 0;
}
```



```
int main()
{
    table.lock();
    Display::clear();
    Display::position(0, 0);
    cout << "The Philosopher's Dinner:" << endl;

    for(int i = 0; i < 5; i++)
        chopstick[i] = new Semaphore;

    phil[0] = new Thread(&philosopher, 0, 5, 32);
    phil[1] = new Thread(&philosopher, 1, 10, 44);
    phil[2] = new Thread(&philosopher, 2, 16, 39);
    phil[3] = new Thread(&philosopher, 3, 16, 24);
    phil[4] = new Thread(&philosopher, 4, 10, 20);

    cout << "Philosophers are alive and hungry!"
         << endl;

    Display::position(7, 44);
    cout << '/';
    Display::position(13, 44);
    cout << '\\';
    Display::position(16, 35);
    cout << '|';
    Display::position(13, 27);
    cout << '/';
    Display::position(7, 27);
    cout << '\\';
    Display::position(19, 0);
```

semaphore as a critical section  
guard (initialized with 1)

```
cout << "The dinner is served ..." << endl;
table.unlock();


for(int i = 0; i < 5; i++) {
    int ret = phil[i]->join();
    table.lock();
    Display::position(20 + i, 0);
    cout << "Philosopher " << i << " ate "
         << ret << " times " << endl;
    table.unlock();
}

for(int i = 0; i < 5; i++)
    delete chopstick[i];
for(int i = 0; i < 5; i++)
    delete phil[i];

cout << "The end!" << endl;

return 0;
}
```





```
int main()
{
    table.lock();
    Display::clear();
    Display::position(0, 0);
    cout << "The Philosopher's Dinner:" << endl;

    for(int i = 0; i < 5; i++)
        chopstick[i] = new Semaphore;

    phil[0] = new Thread(&philosopher, 0, 5, 32);
    phil[1] = new Thread(&philosopher, 1, 10, 44);
    phil[2] = new Thread(&philosopher, 2, 16, 39);
    phil[3] = new Thread(&philosopher, 3, 16, 24);
    phil[4] = new Thread(&philosopher, 4, 10, 20);

    cout << "Philosophers are alive and hungry!"
         << endl;

    Display::position(7, 44);
    cout << '/';
    Display::position(13, 44);
    cout << '\\';
    Display::position(16, 35);
    cout << '|';
    Display::position(13, 27);
    cout << '/';
    Display::position(7, 27);
    cout << '\\';
    Display::position(19, 0);
```

```
    cout << "The dinner is served ..." << endl;
    table.unlock();


    for(int i = 0; i < 5; i++) {
        int ret
        table.l
        Display
        cout << "Philosopher " << i << " ate
            << ret << " times " << endl;
        table.unlock();
    }

    for(int i = 0; i < 5; i++)
        delete chopstick[i];
    for(int i = 0; i < 5; i++)
        delete phil[i];

    cout << "The end!" << endl;

    return 0;
}
```

5 threads from a single function with different parameters



```
int main()
{
    table.lock();
    Display::clear();
    Display::position(0, 0);
    cout << "The dinner is served ..." << endl;
```

join waits for threads to finish

```
    phil[0] = new Thread(&philosopher, 0, 5, 32);
    phil[1] = new Thread(&philosopher, 1, 10, 44);
    phil[2] = new Thread(&philosopher, 2, 16, 39);
    phil[3] = new Thread(&philosopher, 3, 16, 24);
    phil[4] = new Thread(&philosopher, 4, 10, 20);
```

```
    cout << "Philosophers are alive and hungry!"
         << endl;
```

```
    Display::position(7, 44);
    cout << '/';
    Display::position(13, 44);
    cout << '\\';
    Display::position(16, 35);
    cout << '|';
    Display::position(13, 27);
    cout << '/';
    Display::position(7, 27);
    cout << '\\';
    Display::position(19, 0);
```

```
    cout << "The dinner is served ..." << endl;
    table.unlock();
```


```
    for(int i = 0; i < 5; i++) {
        int ret = phil[i]->join();
        table.lock();
        Display::position(20 + i, 0);
        cout << "Philosopher " << i << " ate "
             << ret << " times " << endl;
        table.unlock();
    }
```

```
    for(int i = 0; i < 5; i++)
        delete chopstick[i];
    for(int i = 0; i < 5; i++)
        delete phil[i];
```

```
    cout << "The end!" << endl;
```

```
    return 0;
```

```
}
```



```
int main()
{
    table.lock();
    Display::clear();
    Display::position(0, 0);
    cout << "The Philosopher's Dinner:" << endl;

    for(int i = 0; i < 5; i++)
        chopstick[i] = new Semaphore;

    phil[0] = new Thread(&philosopher, 0, 5, 32);
    phil[1] = new Thread(&philosopher, 1, 10, 44);
    phil[2] = new Thread(&philosopher, 2, 16, 39);
    phil[3] = new Thread(&philosopher, 3, 16, 24);
    phil[4] = new Thread(&philosopher, 4, 10, 20);

    cout << "Philosophers are alive and hungry!"
         << endl;

    Display::position(7, 44);
    cout << '/';
    Display::position(13, 44);
    cout << '\\';
    Display::position(16, 35);
    cout << '|';
    Display::position(13, 27);
    cout << '/';
    Display::position(7, 27);
    cout << '\\';
    Display::position(19, 0);
```

delete what you create, even  
system objects


```
    cout << "The dinner is served ..." << endl;
    table.unlock();

    for(int i = 0; i < 5; i++) {
        int ret = phil[i]->join();
        table.lock();
        Display::position(20 + i, 0);
        cout << "Philosopher " << i << " ate "
             << ret << " times " << endl;
        table.unlock();
    }

    for(int i = 0; i < 5; i++)
        delete chopstick[i];
    for(int i = 0; i < 5; i++)
        delete phil[i];

    cout << "The end!" << endl;

    return 0;
}
```



```
int main()
{
    table.lock();
    Display::clear();
    Display::position(0, 0);
    cout << "The Philosopher's Dinner:" << endl;

    for(int i = 0; i < 5; i++)
        chopstick[i] = new Semaphore;

    phil[0] = new Thread(&philosopher, 0, 5, 32);
    phil[1] = new Thread(&philosopher, 1, 10, 44);
    phil[2] = new Thread(&philosopher, 2, 16, 39);
    phil[3] = new Thread(&philosopher, 3, 16, 24);
    phil[4] = new Thread(&philosopher, 4, 10, 20);

    cout << "Philosophers are alive and hungry!"
         << endl;
```

implicit **exit()**  
+ globals' **destructors**

```
Display::position(10, 44);
cout << '\\';
Display::position(16, 35);
cout << '|';
Display::position(13, 27);
cout << '/';
Display::position(7, 27);
cout << '\\';
Display::position(19, 0);
```

```
cout << "The dinner is served ..." << endl;
table.unlock();
```

```
for(int i = 0; i < 5; i++) {
    int ret = phil[i]->join();
    table.lock();
    Display::position(20 + i, 0);
    cout << "Philosopher " << i << " ate "
         << ret << " times " << endl;
    table.unlock();
}
```

```
for(int i = 0; i < 5; i++)
    delete chopstick[i];
for(int i = 0; i < 5; i++)
    delete phil[i];
```

```
cout << "The end!" << endl;
```

```
return 0;
```

```
}
```

```
// EPOS Synchronizer Component Test Program
```

```
#include <time.h>
#include <synchronizer.h>
#include <process.h>
```

```
using namespace EPOS;
```

```
const int iterations = 100;
```

```
OStream cout;
```

```
const int BUF_SIZE = 16;
char buffer[BUF_SIZE];
Semaphore empty(BUF_SIZE);
Semaphore full(0);
```

```
int consumer()
{
    int out = 0;
    for(int i = 0; i < iterations; i++) {
        full.p();
        cout << "C<- " << buffer[out] << "\t";
        out = (out + 1) % BUF_SIZE;
        Alarm::delay(100000);
        empty.v();
    }

    return 0;
}
```

```
int main()
{
    cout << "Producer x Consumer" << endl;

    Thread * cons = new Thread(&consumer);


    // producer
    int in = 0;
    for(int i = 0; i < iterations; i++) {
        empty.p();
        Alarm::delay(100000);
        buffer[in] = 'a' + in;
        cout << "P->" << buffer[in] << "\t";
        in = (in + 1) % BUF_SIZE;
        full.v();
    }

    cons->join();

    cout << "The end!" << endl;

    delete cons;

    return 0;
}
```



## Producer x Consumer

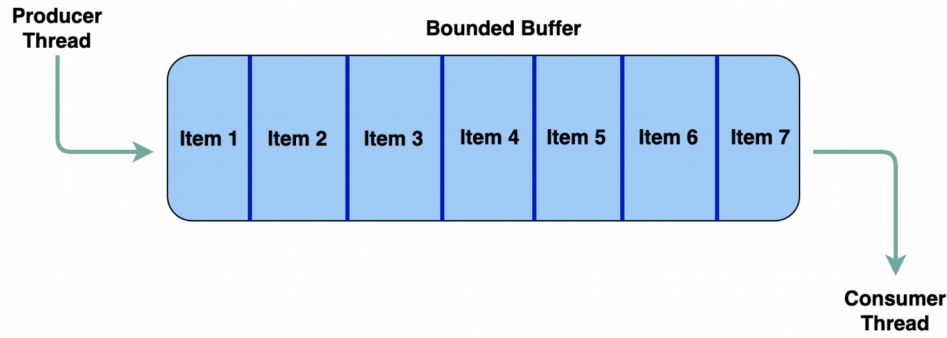



Figure by educative.io

```
Semaphore empty(BUF_SIZE);  
Semaphore full(0);
```

```
int consumer()  
{  
    int out = 0;  
    for(int i = 0; i < iterations; i++) {  
        full.p();  
        cout << "C<->" << buffer[out] << "\t";  
        out = (out + 1) % BUF_SIZE;  
        Alarm::delay(100000);  
        empty.v();  
    }  
  
    return 0;  
}
```

```
int main()  
{  
    cout << "Producer x Consumer" << endl;  
  
    Thread * cons = new Thread(&consumer);  
  
    // producer  
    int in = 0;  
    for(int i = 0; i < iterations; i++) {  
        empty.p();  
        Alarm::delay(100000);  
        buffer[in] = 'a' + in;  
        cout << "P->" << buffer[in] << "\t";  
        in = (in + 1) % BUF_SIZE;  
        full.v();  
    }  
  
    cons->join();  
  
    cout << "The end!" << endl;  
  
    delete cons;  
  
    return 0;  
}
```



## // EPOS Synchronizer Component Test Program

```
#include <time.h>
#include <synchronizer.h>
#include <process.h>

using namespace EPOS;

const int iterations = 100;

OStream cout;

const int BUF_SIZE = 16;
char buffer[BUF_SIZE];
Semaphore empty(BUF_SIZE);
Semaphore full(0);

int consumer()
{
    int out = 0;
    for(int i = 0; i < iterations; i++) {
        full.p();
        cout << "C<- " << buffer[out] << "\t";
        out = (out + 1) % BUF_SIZE;
        Alarm::delay(100000);
        empty.v();
    }

    return 0;
}
```

semaphores as atomic resource  
counters

```
int main()
{
    cout << "Producer x Consumer" << endl;

    Thread * cons = new Thread(&consumer);


    // producer
    int in = 0;
    for(int i = 0; i < iterations; i++) {
        empty.p();
        Alarm::delay(100000);
        buffer[in] = 'a' + in;
        cout << buffer[in] << "\t";
        in = (in + 1) % BUF_SIZE;
        full.v();
    }

    cons->join();

    cout << "The end!" << endl;

    delete cons;

    return 0;
}
```



## // EPOS Synchronizer Component Test Program

```
#include <time.h>
#include <synchronizer.h>
#include <process.h>
```

```
using namespace EPOS;
```

```
const int iterations = 100;
```

```
OStream cout;
```

```
const int BUF_SIZE = 16;
char buffer[BUF_SIZE];
Semaphore empty(BUF_SIZE);
Semaphore full(0);
```

```
int consumer()
{
    int out = 0;
    for(int i = 0; i < iterations; i++) {
        full.p();
        cout << "C<- " << buffer[out] << "\t";
        out = (out + 1) % BUF_SIZE;
        Alarm::delay(100000);
        empty.v();
    }

    return 0;
}
```

consumer  
consumes from buffer


```
int main()
{
    cout << "Producer x Consumer" << endl;

    Thread * cons = new Thread(&consumer);

    // producer
    int in = 0;
    for(int i = 0; i < iterations; i++) {
        empty.p();
        Alarm::delay(100000);
        buffer[in] = 'a' + in;
        cout << "P->" << buffer[in] << "\t";
        in = (in + 1) % BUF_SIZE;
        full.v();
    }

    delete cons;

    return 0;
}
```





producer produces into buffer

```
#include <synchronizer.h>
#include <process.h>
```

```
using namespace EPOS;
```

```
const int iterations = 100;
```

```
OStream cout;
```

```
const int BUF_SIZE = 16;
char buffer[BUF_SIZE];
Semaphore empty(BUF_SIZE);
Semaphore full(0);
```

```
int consumer()
{
    int out = 0;
    for(int i = 0; i < iterations; i++) {
        full.p();
        cout << "C<- " << buffer[out] << "\t";
        out = (out + 1) % BUF_SIZE;
        Alarm::delay(100000);
        empty.v();
    }

    return 0;
}
```

st Program

```
int main()
{
    cout << "Producer x Consumer" << endl;

    Thread * cons = new Thread(&consumer);


    // producer
    int in = 0;
    for(int i = 0; i < iterations; i++) {
        empty.p();
        Alarm::delay(100000);
        buffer[in] = 'a' + in;
        cout << "P->" << buffer[in] << "\t";
        in = (in + 1) % BUF_SIZE;
        full.v();
    }

    cons->join();

    cout << "The end!" << endl;

    delete cons;

    return 0;
}
```



## // EPOS Synchronizer Component Test Program

```
#include <time.h>
#include <synchronizer.h>
#include <process.h>
```

```
using namespace EPOS;
```

```
const int iterations = 100;
```

```
OStream cout;
```

```
const int BUF_SIZE = 16;
char buffer[BUF_SIZE];
```

join waits for threads to finish

```
int consumer()
{
    int out = 0;
    for(int i = 0; i < iterations; i++) {
        full.p();
        cout << "C<- " << buffer[out] << "\t";
        out = (out + 1) % BUF_SIZE;
        Alarm::delay(100000);
        empty.v();
    }

    return 0;
}
```

```
int main()
{
    cout << "Producer x Consumer" << endl;

    Thread * cons = new Thread(&consumer);


    // producer
    int in = 0;
    for(int i = 0; i < iterations; i++) {
        empty.p();
        Alarm::delay(100000);
        buffer[in] = 'a' + in;
        cout << "P->" << buffer[in] << "\t";
        in = (in + 1) % BUF_SIZE;
        full.v();
    }

    cons->join();

    cout << "The end!" << endl;

    delete cons;

    return 0;
}
```



```
// EPOS Synchronizer Component Test Program
```

```
#include <time.h>
#include <synchronizer.h>
#include <process.h>
```

```
using namespace EPOS;
```

```
const int iterations = 100;
```

```
OStream cout;
```

```
const int BUF_SIZE = 16;
char buffer[BUF_SIZE];
Semaphore empty(BUF_SIZE);
Semaphore full(0);
```

**delete** what you create, even  
system objects

```
for(int i = 0; i < iterations; i++) {
    full.p();
    cout << "C<- " << buffer[out] << "\t";
    out = (out + 1) % BUF_SIZE;
    Alarm::delay(100000);
    empty.v();
}

return 0;
}
```

```
int main()
{
    cout << "Producer x Consumer" << endl;

    Thread * cons = new Thread(&consumer);


    // producer
    int in = 0;
    for(int i = 0; i < iterations; i++) {
        empty.p();
        Alarm::delay(100000);
        buffer[in] = 'a' + in;
        cout << "P->" << buffer[in] << "\t";
        in = (in + 1) % BUF_SIZE;
        full.v();
    }

    cons->join();

    cout << "The end!" << endl;

    delete cons;

    return 0;
}
```



## // EPOS Synchronizer Component Test Program

```
#include <time.h>
#include <synchronizer.h>
#include <process.h>
```

```
using namespace EPOS;
```

```
const int iterations = 100;
```

```
OStream cout;
```

```
const int BUF_SIZE = 16;
char buffer[BUF_SIZE];
Semaphore empty(BUF_SIZE);
Semaphore full(0);
```

```
int consumer()
```

```
{
    int out = 0;
    for(int i = 0; i < iterations; i++) {
        full.p();
```

implicit **exit()**  
+ globals' destructors

```
        cout << "P->" << buffer[in] << "\t";
        empty.v();
    }

    return 0;
}
```

```
int main()
```

```
{
    cout << "Producer x Consumer" << endl;

    Thread * cons = new Thread(&consumer);


    // producer
    int in = 0;
    for(int i = 0; i < iterations; i++) {
        empty.p();
        Alarm::delay(100000);
        buffer[in] = 'a' + in;
        cout << "P->" << buffer[in] << "\t";
        in = (in + 1) % BUF_SIZE;
        full.v();
    }

    cons->join();

    cout << "The end!" << endl;

    delete cons;

    return 0;
}
```



```
// EPOS Thread Component Declarations
```

```
#ifndef __process_h
#define __process_h
```

```
#include <architecture.h>
#include <utility/queue.h>
#include <utility/handler.h>
#include <machine/common.h>
#include <machine/timer.h>
```

```
extern "C" { void __exit(); }
```

```
__BEGIN_SYS
```

```
class Thread
{
```

```
    friend class Init_First;
    friend class Init_System;
    friend class Synchronizer_Common;
    friend class Alarm;
    friend class System;
```

```
protected:
```

```
    static const bool reboot =
        Traits<System>::reboot;
    static const unsigned int QUANTUM =
        Traits<Thread>::QUANTUM;
    static const unsigned int STACK_SIZE =
        Traits<Application>::STACK_SIZE;
```

```
typedef CPU::Log_Addr Log_Addr;
typedef CPU::Context Context;
```

```
public:
```

```
    // Thread State
```

```
    enum State { RUNNING, READY, SUSPENDED,
                 WAITING, FINISHING };
```

```
    // Thread Priority
```

```
    typedef unsigned int Priority;
```

```
    enum {
        HIGH = 0,
        NORMAL = 15,
        LOW = 31
    };
```


```
    // Thread Configuration
```

```
    struct Configuration {
        Configuration(const State & s = READY,
                     const Priority & p = NORMAL,
                     unsigned int ss = STACK_SIZE)
            : state(s), priority(p), stack_size(ss) {}
    };
```

```
        State state;
        Priority priority;
        unsigned int stack_size;
```

```
    // Thread Queue
```

```
    typedef Ordered_Queue<Thread, Priority> Queue;
```



```
// EPOS Thread Component Declarations
```

```
#ifndef __process_h  
#define __process_h
```

```
#include <architecture.h>  
#include <utility/queue.h>  
#include <utility/handler.h>  
#include <machine/common.h>  
#include <machine/timer.h>
```

```
extern "C" { void __exit(); }
```

```
__BEGIN_SYS
```

```
class Thread  
{
```

```
    friend class Init_First;  
    friend class Init_System;  
    friend class Synchronizer_Common;  
    friend class Alarm;  
    friend class System;
```

```
protected:
```

```
    static const bool reboot =  
        Traits<System>::reboot;  
    static const unsigned int QUANTUM =  
        Traits<Thread>::QUANTUM;  
    static const unsigned int STACK_SIZE =  
        Traits<Application>::STACK_SIZE;
```

header file's #include guard

```
public:
```

```
// Thread State
```

```
enum State { RUNNING, READY, SUSPENDED,  
             WAITING, FINISHING };
```

```
// Thread Priority
```

```
typedef unsigned int Priority;
```

```
enum {  
    HIGH = 0,  
    NORMAL = 15,  
    LOW = 31
```

```
};
```

```
// Thread Configuration
```


```
struct Configuration {  
    Configuration(const State & s = READY,  
                  const Priority & p = NORMAL,  
                  unsigned int ss = STACK_SIZE)  
        : state(s), priority(p), stack_size(ss) {}
```

```
    State state;  
    Priority priority;  
    unsigned int stack_size;
```

```
};
```

```
// Thread Queue
```

```
typedef Ordered_Queue<Thread, Priority> Queue;
```



```
// EPOS Thread Component Declarations
```

```
#ifndef __process_h  
#define __process_h
```

```
#include <architecture.h>  
#include <utility/queue.h>  
#include <utility/handler.h>  
#include <machine/common.h>  
#include <machine/timer.h>
```

```
extern "C" { void __exit(); }
```

```
__BEGIN_SYS
```

```
class Thread  
{
```

```
    friend class Init_First;  
    friend class Init_System;  
    friend class Synchronizer_Common;  
    friend class Alarm;  
    friend class System;
```

```
protected:
```

```
    static const bool reboot =  
        Traits<System>::reboot;  
    static const unsigned int QUANTUM =  
        Traits<Thread>::QUANTUM;  
    static const unsigned int STACK_SIZE =  
        Traits<Application>::STACK_SIZE;
```

```
typedef CPU::Log_Addr Log_Addr;  
typedef CPU::Context Context;
```

```
public:
```

```
    // Thread State
```

```
    enum State { RUNNING, READY, SUSPENDED,  
                WAITING, FINISHING };
```

```
    // Thread Priority
```

helper function with C linkage

```
        NORMAL = 15,  
        LOW = 31
```

```
};
```

```
    // Thread Configuration
```


```
    struct Configuration {  
        Configuration(const State & s = READY,  
                      const Priority & p = NORMAL,  
                      unsigned int ss = STACK_SIZE)  
            : state(s), priority(p), stack_size(ss) {}
```

```
        State state;  
        Priority priority;  
        unsigned int stack_size;
```

```
};
```

```
    // Thread Queue
```

```
    typedef Ordered_Queue<Thread, Priority> Queue;
```



```
// EPOS Thread Component Declarations
```

```
#ifndef __process_h
#define __process_h
```

```
#include <architecture.h>
#include <utility/queue.h>
#include <utility/handler.h>
#include <machine/common.h>
#include <machine/timer.h>
```

```
extern "C" { void __exit(); }
```

```
__BEGIN_SYS
```

```
class Thread
{
```

```
    friend class Init_First;
    friend class Init_System;
    friend class Synchronizer_Common;
    friend class Alarm;
    friend class System;
```

```
protected:
```

```
    static const bool reboot =
        Traits<System>::reboot;
    static const unsigned int QUANTUM =
        Traits<Thread>::QUANTUM;
    static const unsigned int STACK_SIZE =
        Traits<Application>::STACK_SIZE;
```

```
typedef CPU::Log_Addr Log_Addr;
typedef CPU::Context Context;
```

```
public:
```

```
// Thread State
```

```
enum State { RUNNING, READY, SUSPENDED,
             WAITING, FINISHING };
```

```
// Thread Priority
```

```
typedef unsigned int Priority;
```

namespaces: API, UTIL, SYS

```
};
```

```
// Thread Configuration
```


```
struct Configuration {
    Configuration(const State & s = READY,
                  const Priority & p = NORMAL,
                  unsigned int ss = STACK_SIZE)
        : state(s), priority(p), stack_size(ss) {}
```

```
    State state;
    Priority priority;
    unsigned int stack_size;
```

```
};
```

```
// Thread Queue
```

```
typedef Ordered_Queue<Thread, Priority> Queue;
```





```
// EPOS Thread Component Declarations
```

```
#ifndef __process_h  
#define __process_h
```

```
#include <architecture.h>  
#include <utility/queue.h>  
#include <utility/handler.h>  
#include <machine/common.h>  
#include <machine/timer.h>
```

```
extern "C" { void __exit(); }
```

```
__BEGIN_SYS
```

```
class Thread  
{
```

```
    friend class Init_First;  
    friend class Init_System;  
    friend class Synchronizer_Common;  
    friend class Alarm;  
    friend class System;
```

```
protected:
```

```
    static const bool reboot =  
        Traits<System>::reboot;  
    static const unsigned int QUANTUM =  
        Traits<Thread>::QUANTUM;  
    static const unsigned int STACK_SIZE  
        Traits<Application>::STACK_SIZE;
```

```
class Class  
{  
    friend class .... ;
```

```
private:  
    flag ... ;  
    CONSTANT ... ;  
    Type .... ;
```

```
public:  
    CONSTANT ... ;  
    Types .... ;
```

```
public:  
    constructor();  
    destructor();  
    method();
```

```
    static method();
```


```
private:  
    method();  
    static method();
```

```
    static void init();
```

```
private:  
    attribute;
```

```
    static attribute;
```

```
}
```



```
SPENDED,  
};
```

```
s = READY,  
y & p = NORMAL,  
ss = STACK_SIZE)  
ack_size(ss) {}
```

```
riority> Queue;
```

```
// EPOS Thread Component Declarations
```

```
#ifndef __process_h  
#define __process_h
```

```
#include <architecture.h>  
#include <utility/queue.h>  
#include <utility/handler.h>  
#include <machine/common.h>  
#include <machine/timer.h>
```

```
extern "C" { void __exit(); }
```

```
__BEGIN_SYS
```

```
class Thread  
{
```

```
    friend class Init_First;  
    friend class Init_System;  
    friend class Synchronizer_Common;  
    friend class Alarm;  
    friend class System;
```

```
protected:
```

```
    static const bool reboot =  
        Traits<System>::reboot;  
    static const unsigned int QUANTUM =  
        Traits<Thread>::QUANTUM;  
    static const unsigned int STACK_SIZE =  
        Traits<Application>::STACK_SIZE;
```

```
class Class  
{  
    friend class .... ;
```

```
private:  
    flag ... ;  
    CONSTANT ... ;  
    Type .... ;
```


```
public:  
    CONSTANT ... ;  
    Types .... ;
```

```
public:  
    constructor();  
    destructor();  
    method();  
  
    static method();
```

```
private:  
    method();  
    static method();  
  
    static void init();
```

```
private:  
    attribute;  
  
    static attribute;  
}
```

I  
N  
T  
E  
R  
F  
A  
C  
E



```
SPENDED,  
};
```

```
s = READY,  
y & p = NORMAL,  
ss = STACK_SIZE)  
ack_size(ss) {}
```

```
riority> Queue;
```

```
// EPOS Thread Component Declarations
```

```
#ifndef __process_h
#define __process_h
```

```
#include <architecture.h>
#include <utility/queue.h>
#include <utility/handler.h>
#include <machine/common.h>
#include <machine/timer.h>
```

```
extern "C" { void __exit(); }
```

```
__BEGIN_SYS
```

```
class Thread
{
```

```
    friend class Init_First;
    friend class Init_System;
    friend class Synchronizer_Common;
    friend class Alarm;
    friend class System;
```

```
protected:
```

```
    static const bool reboot =
        Traits<System>::reboot;
    static const unsigned int QUANTUM =
        Traits<Thread>::QUANTUM;
    static const unsigned int STACK_SIZE =
        Traits<Application>::STACK_SIZE;
```

```
typedef CPU::Log_Addr Log_Addr;
typedef CPU::Context Context;
```

```
public:
```

```
// Thread State
```

```
enum State { RUNNING, READY, SUSPENDED,
             WAITING, FINISHING };
```

```
// Thread Priority
```

```
typedef unsigned int Priority;
```

```
enum {
    HIGH = 0,
    NORMAL = 15,
    LOW = 31
};
```


friend classes to prevent  
encapsulation breaking

```
    State state;
    Priority priority;
    unsigned int stack_size;
    : state(s), priority(p), stack_size(ss) {}
```

```
    State state;
    Priority priority;
    unsigned int stack_size;
};
```

```
// Thread Queue
```

```
typedef Ordered_Queue<Thread, Priority> Queue;
```



```
// EPOS Thread Component Declarations
```

```
#ifndef __process_h
#define __process_h
```

```
#include <architecture.h>
#include <utility/queue.h>
#include <utility/handler.h>
#include <machine/common.h>
#include <machine/timer.h>
```

```
extern "C" { void __exit(); }
```

```
__BEGIN_SYS
```

```
class Thread
{
```

```
    friend class Init_First;
    friend class Init_System;
    friend class Synchronizer_Common;
    friend class Alarm;
    friend class System;
```

```
protected:
```

```
    static const bool reboot =
        Traits<System>::reboot;
    static const unsigned int QUANTUM =
        Traits<Thread>::QUANTUM;
    static const unsigned int STACK_SIZE =
        Traits<Application>::STACK_SIZE;
```

```
typedef CPU::Log_Addr Log_Addr;
typedef CPU::Context Context;
```

```
public:
```

```
// Thread State
```

```
enum State { RUNNING, READY, SUSPENDED,
             WAITING, FINISHING };
```

```
// Thread Priority
```

```
typedef unsigned int Priority;
```

```
enum {
    HIGH = 0,
    NORMAL = 15,
    LOW = 31
};
```

```
// Thread Configuration
```


```
struct Configuration {
    Configuration(const State & s = READY,
                  const Priority & p = NORMAL,
                  unsigned int ss = STACK_SIZE)
        : state(s), priority(p), stack_size(ss) {}
    State state;
    Priority priority;
    unsigned int stack_size;
```

static const int instead of #define

```
};
```

```
// Thread Queue
```

```
typedef Ordered_Queue<Thread, Priority> Queue;
```



```
// EPOS Thread Component Declarations
```

```
#ifndef __process_h
#define __process_h
```

```
#include <architecture.h>
#include <utility/queue.h>
#include <utility/handler.h>
#include <machine/common.h>
#include <machine/timer.h>
```

```
extern "C" { void __exit(); }
```

```
__BEGIN_SYS
```

```
class Thread
{
```

```
    friend class Init_First;
    friend class Init_System;
    friend class Synchronizer_Common;
    friend class Alarm;
    friend class System;
```

```
protected:
```

```
    static const bool reboot =
        Traits<System>::reboot;
    static const unsigned int QUANTUM =
        Traits<Thread>::QUANTUM;
    static const unsigned int STACK_SIZE =
        Traits<Application>::STACK_SIZE;
```

```
typedef CPU::Log_Addr Log_Addr;
typedef CPU::Context Context;
```

```
public:
```

```
    // Thread State
```

```
    enum State { RUNNING, READY, SUSPENDED,
                 WAITING, FINISHING };
```

```
    // Thread Priority
```

```
    typedef unsigned int Priority;
```

```
    enum {
        HIGH = 0,
        NORMAL = 15,
        LOW = 31
    };
```


```
    // Thread Configuration
```

```
    struct Configuration {
        Configuration(const State & s = READY,
                     const Priority & p = NORMAL,
                     unsigned int ss = STACK_SIZE)
            : state(s), priority(p), stack_size(ss) {}
    };
```

```
    State state;
    Priority priority;
    unsigned int stack_size;
```

```
    // Thread Queue
```

```
    typedef Ordered_Queue<Thread, Priority> Queue;
```



Traits<> like in the STL (Stepanov)

imports

```
#define __process_n
```

```
#include <architecture.h>
#include <utility/queue.h>
#include <utility/handler.h>
#include <machine/common.h>
#include <machine/timer.h>
```

```
extern "C" { void __exit(); }
```

```
__BEGIN_SYS
```

```
class Thread
{
```

```
    friend class Init_First;
    friend class Init_System;
    friend class Synchronizer_Common;
    friend class Alarm;
    friend class System;
```

```
protected:
```

```
    static const bool reboot =
        Traits<System>::reboot;
    static const unsigned int QUANTUM =
        Traits<Thread>::QUANTUM;
    static const unsigned int STACK_SIZE =
        Traits<Application>::STACK_SIZE;
```

ions

```
typedef CPU::Log_Addr Log_Addr;
typedef CPU::Context Context;
```

```
public:
```

```
// Thread State
```

```
enum State { RUNNING, READY, SUSPENDED,
             WAITING, FINISHING };
```

```
// Thread Priority
```

```
typedef unsigned int Priority;
```

```
enum {
    HIGH = 0,
    NORMAL = 15,
    LOW = 31
};
```

```
// Thread Configuration
```


```
struct Configuration {
    Configuration(const State & s = READY,
                  const Priority & p = NORMAL,
                  unsigned int ss = STACK_SIZE)
        : state(s), priority(p), stack_size(ss) {}

```

```
    State state;
    Priority priority;
    unsigned int stack_size;
};
```

```
// Thread Queue
```

```
typedef Ordered_Queue<Thread, Priority> Queue;
```



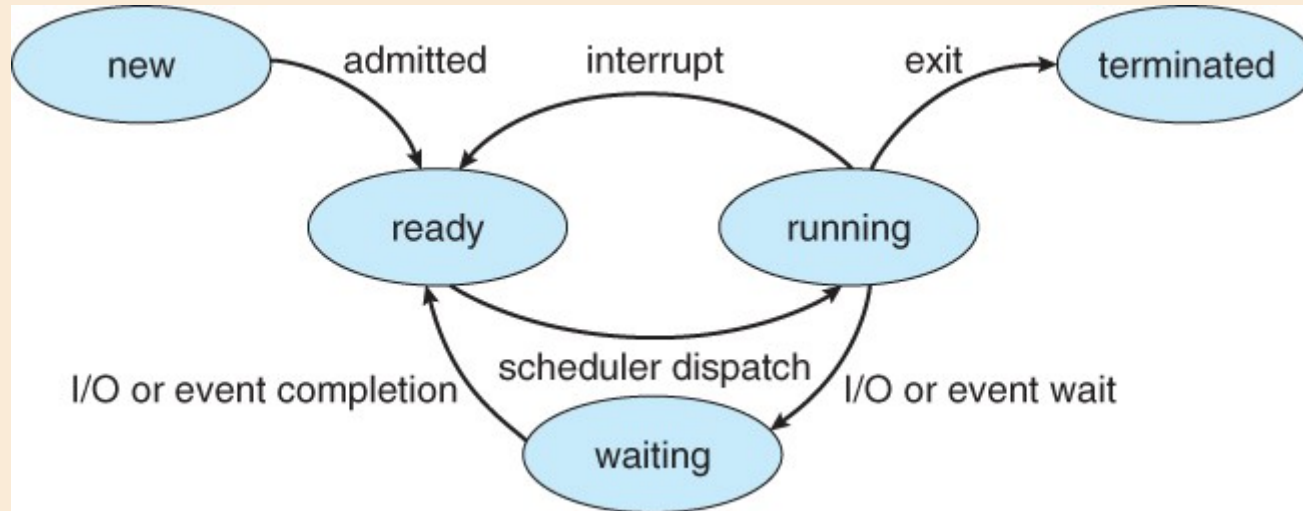
```
// EPOS Thread Component Declarations
```

```
#ifndef __process_h
#define __process_h
```

```
#include <architecture.h>
#include <utility/queue.h>
#include <utility/handler.h>
#include <machine/common.h>
#include <machine/timer.h>
```

```
extern "C" { void exit(): }
```

## States as in Silberschatz



```
typedef CPU::Log_Addr Log_Addr;
typedef CPU::Context Context;
```

```
public:
```

```
// Thread State
```

```
enum State { RUNNING, READY, SUSPENDED,
             WAITING, FINISHING };
```

```
// Thread Priority
```

```
typedef unsigned int Priority;
```

```
enum {
```

```
    HIGH = 0,
```

```
duration
```


```
tion {
```

```
h(const State & s = READY,
   const Priority & p = NORMAL,
   unsigned int ss = STACK_SIZE)
priority(p), stack_size(ss) {}
```

```
riority;
```

```
stack_size;
```

```
Queue<Thread, Priority> Queue;
```



```
// EPOS Thread Component Declarations
```

```
#ifndef __process_h  
#define __process_h
```

```
#include <architecture.h>  
#include <utility/queue.h>
```

unsigned int + anonymous enum for  
standard operations

```
extern "C" { void __exit(); }
```

```
__BEGIN_SYS
```

```
class Thread  
{
```

```
    friend class Init_First;  
    friend class Init_System;  
    friend class Synchronizer_Common;  
    friend class Alarm;  
    friend class System;
```

```
protected:
```

```
    static const bool reboot =  
        Traits<System>::reboot;  
    static const unsigned int QUANTUM =  
        Traits<Thread>::QUANTUM;  
    static const unsigned int STACK_SIZE =  
        Traits<Application>::STACK_SIZE;
```

```
typedef CPU::Log_Addr Log_Addr;  
typedef CPU::Context Context;
```

```
public:
```

```
// Thread State
```

```
enum State { RUNNING, READY, SUSPENDED,  
            WAITING, FINISHING };
```

```
// Thread Priority
```

```
typedef unsigned int Priority;
```

```
enum {  
    HIGH = 0,  
    NORMAL = 15,  
    LOW = 31  
};
```

```
// Thread Configuration
```


```
struct Configuration {  
    Configuration(const State & s = READY,  
                  const Priority & p = NORMAL,  
                  unsigned int ss = STACK_SIZE)  
        : state(s), priority(p), stack_size(ss) {}
```

```
    State state;  
    Priority priority;  
    unsigned int stack_size;
```

```
};
```

```
// Thread Queue
```

```
typedef Ordered_Queue<Thread, Priority> Queue;
```





```
// EPOS Thread Component Declarations
```

```
#ifndef __process_h
#define __process_h
```

```
#include <architecture.h>
#include <utility/queue.h>
#include <utility/handler.h>
#include <machine/common.h>
#include <machine/timer.h>
```

```
extern "C" { void __exit(); }
```

```
__BEGIN_SYS
```

```
class Thread
{
```

```
    friend class Init_First;
    friend class Init_System;
    friend class Synchronizer_Common;
    friend class Alarm;
    friend class System;
```

```
protected:
```

```
    static const bool reboot =
        Traits<System>::reboot;
    static const unsigned int QUANTUM =
```

READY queue

STACK\_SIZE =  
\_SIZE;

```
typedef CPU::Log_Addr Log_Addr;
typedef CPU::Context Context;
```

```
public:
```

```
// Thread State
```

```
enum State { RUNNING, READY, SUSPENDED,
             WAITING, FINISHING };
```

```
// Thread Priority
```

```
typedef unsigned int Priority;
```

```
enum {
    HIGH = 0,
    NORMAL = 15,
    LOW = 31
};
```


```
// Thread Configuration
```

```
struct Configuration {
    Configuration(const State & s = READY,
                  const Priority & p = NORMAL,
                  unsigned int ss = STACK_SIZE)
        : state(s), priority(p), stack_size(ss) {}
```

```
    State state;
    Priority priority;
    unsigned int stack_size;
```

```
// Thread Queue
```

```
typedef Ordered_Queue<Thread, Priority> Queue;
```



```
// EPOS Synchronizer Components
```

```
#ifndef __synchronizer_h  
#define __synchronizer_h
```

```
#include <architecture.h>  
#include <utility/handler.h>  
#include <process.h>
```

```
__BEGIN_SYS
```

```
class Synchronizer_Common  
{
```

```
protected:  
    Synchronizer_Common() {}
```

```
    // Atomic operations
```

```
    bool tsl(volatile bool & lock) {  
        return CPU::tsl(lock); }  
    int finc(volatile int & number) {  
        return CPU::finc(number); }  
    int fdec(volatile int & number) {  
        return CPU::fdec(number); }
```

```
    // Thread operations
```

```
    void begin_atomic() { Thread::lock(); }  
    void end_atomic() { Thread::unlock(); }
```

```
    void sleep() { Thread::yield(); }  
    void wakeup() { end_atomic(); }  
    void wakeup_all() { end_atomic(); }
```

```
};
```

```
class Mutex: protected Synchronizer_Common  
{  
public:  
    Mutex();  
    ~Mutex();
```

```
    void lock();  
    void unlock();
```

```
private:  
    volatile bool _locked;  
};
```

```
class Semaphore: protected Synchronizer_Common  
{
```


```
public:  
    Semaphore(int v = 1);  
    ~Semaphore();
```

```
    void p();  
    void v();
```

```
private:  
    volatile int _value;  
};
```

```
__END_SYS
```

```
#endif
```



```
// EPOS Synchronizer Components
```

```
#ifndef __synchronizer_h
#define __synchronizer_h

#include <architecture.h>
#include <utility/handler.h>
#include <process.h>
```

```
__BEGIN_SYS
```

```
class Synchronizer_Common
{
```

```
protected:
    Synchronizer_Common() {}
```

```
    // Atomic operations
```

```
    bool tsl(volatile bool & lock) {
        return CPU::tsl(lock); }
    int finc(volatile int & number) {
        return CPU::finc(number); }
    int fdec(volatile int & number) {
        return CPU::fdec(number); }
```

```
    // Thread operations
```

```
    void begin_atomic() { Thread::lock(); }
    void end_atomic() { Thread::unlock(); }
```

```
    void sleep() { Thread::yield(); }
    void wakeup() { end_atomic(); }
    void wakeup_all() { end_atomic(); }
```

```
};
```

```
class Mutex: protected Synchronizer_Common
{
public:
```

common packages are reused by a family of  
abstractions  
(protected constructor)

```
    void unlock();
```

```
private:
```

```
    volatile bool _locked;
```

```
};
```

```
class Semaphore: protected Synchronizer_Common
{
public:
```

```
    Semaphore(int v = 1);
    ~Semaphore();
```

```
    void p();
    void v();
```


```
private:
```

```
    volatile int _value;
```

```
};
```

```
__END_SYS
```

```
#endif
```



```
// EPOS Synchronizer Components
```

```
#ifndef __synchronizer_h
#define __synchronizer_h

#include <architecture.h>
#include <utility/handler.h>
#include <process.h>
```

```
__BEGIN_SYS
```

```
class Synchronizer_Common
{
protected:
```

```
    Synchronizer_Common() {}

    // Atomic operations
    bool tsl(volatile bool & lock) {
        return CPU::tsl(lock); }
    int finc(volatile int & number) {
        return CPU::finc(number); }
    int fdec(volatile int & number) {
        return CPU::fdec(number); }
```

```
    // Thread operations
```

```
    void begin_atomic() { Thread::lock(); }
    void end_atomic() { Thread::unlock(); }
```

```
    void sleep() { Thread::yield(); }
    void wakeup() { end_atomic(); }
    void wakeup_all() { end_atomic(); }
```

```
};
```

```
class Mutex: protected Synchronizer_Common
{
public:
    Mutex();
    ~Mutex();
```

```
    void lock();
    void unlock();
```

```
private:
    volatile bool _locked;
};
```

```
class Semaphore
{
public:
```


```
    Semaphore();
    ~Semaphore();
```

```
    void p();
    void v();
```

```
private:
    volatile int _value;
};
```

```
__END_SYS
```

```
#endif
```



common packages are reused by a family of abstractions (private inheritance)

```
// EPOS Synchronizer Components
```

```
#ifndef __synchronizer_h
#define __synchronizer_h

#include <architecture.h>
#include <utility/handler.h>
#include <process.h>
```

```
__BEGIN_SYS
```

```
class Synchronizer_Common
{
```

```
protected:
    Synchronizer_Common() {}
```

```
    // Atomic operations
```

```
    bool tsl(volatile bool & lock) {
        return CPU::tsl(lock); }
    int finc(volatile int & number) {
```

```
        return CPU::finc(number); }
    int fdec(volatile int & number) {
```

```
        return CPU::fdec(number); }
```

```
    // Thread operations
```

```
    void begin_atomic() { Thread::lock(); }
    void end_atomic() { Thread::unlock(); }
```

```
    void sleep() { Thread::yield(); }
```

```
    void wakeup() { end_atomic(); }
```

```
    void wakeup_all() { end_atomic(); }
```

```
};
```

```
class Mutex: protected Synchronizer_Common
{
public:
    Mutex();
    ~Mutex();
```

```
    void lock();
    void unlock();
```

```
private:
```

```
    volatile
```

atomic operations provided by the  
architecture

```
};

class Semaphore: protected Synchronizer_Common
{
```

```
public:
    Semaphore(int v = 1);
    ~Semaphore();
```


```
    void p();
    void v();
```

```
private:
```

```
    volatile int _value;
};
```

```
__END_SYS
```

```
#endif
```



```
// EPOS Synchronizer Components
```

```
#ifndef __synchronizer_h
#define __synchronizer_h

#include <architecture.h>
#include <utility/handler.h>
#include <process.h>
```

```
__BEGIN_SYS
```

```
class Synchronizer_Common
{
```

```
protected:
    Synchronizer_Common() {}
```

```
    // Atomic operations
```

```
    bool tsl(volatile bool & lock) {
        return CPU::tsl(lock); }
    int finc(volatile int & number) {
        return CPU::finc(number); }
    int fdec(volatile int & number) {
        return CPU::fdec(number); }
```

```
    // Thread operations
```

```
    void begin_atomic() { Thread::lock(); }
    void end_atomic() { Thread::unlock(); }
```

```
    void sleep() { Thread::yield(); }
    void wakeup() { end_atomic(); }
    void wakeup_all() { end_atomic(); }
```

```
};
```

```
class Mutex: protected Synchronizer_Common
{
public:
    Mutex();
    ~Mutex();
```

```
    void lock();
    void unlock();
```

```
private:
    volatile bool _locked;
};
```


```
class Semaphore: protected Synchronizer_Common
{
public:
    Semaphore(int v = 1);
    ~Semaphore();
```

```
    void {
    void {
        preemption prevention operations
        provided by thread
```

```
private:
    volatile int _value;
};
```

```
__END_SYS
```

```
#endif
```



```
// EPOS Synchronizer Components
```

```
#ifndef __synchronizer_h
#define __synchronizer_h

#include <architecture.h>
#include <utility/handler.h>
#include <process.h>
```

```
__BEGIN_SYS
```

```
class Synchronizer_Common
{
protected:
```

```
    Synchronizer_Common() {}
```

```
    // Atomic operations
```

```
    bool tsl(volatile bool & lock) {
        return CPU::tsl(lock); }
    int finc(volatile int & number) {
        return CPU::finc(number); }
    int fdec(volatile int & number) {
        return CPU::fdec(number); }
```


```
    // Thread operations
```

```
    void begin_atomic() { Thread::lock(); }
    void end_atomic() { Thread::unlock(); }
```

```
    void sleep() { Thread::yield(); }
    void wakeup() { end_atomic(); }
    void wakeup_all() { end_atomic(); }
```

```
};
```

```
class Mutex: protected Synchronizer_Common
{
public:
    Mutex();
    ~Mutex();
```



## Friend Classes

**Thread::lock()**, which disables preemptive scheduling operations is clearly not part of the interface of class **Thread**. It is therefore declared as a private member function and the **Synchronizer\_Common** class is declared a friend of class **Thread**.

```
    Semaphore(int v = 1);
    ~Semaphore();
```

```
    void p();
    void v();
```

```
private:
    volatile int _value;
};
```

```
__END_SYS
```

```
#endif
```

classic Mutex with a volatile lock

```
#define __SYNCHRONIZER_H__

#include <architecture.h>
#include <utility/handler.h>
#include <process.h>

__BEGIN_SYS

class Synchronizer_Common
{
protected:
    Synchronizer_Common() {}

    // Atomic operations
    bool tsl(volatile bool & lock) {
        return CPU::tsl(lock); }
    int finc(volatile int & number) {
        return CPU::finc(number); }
    int fdec(volatile int & number) {
        return CPU::fdec(number); }

    // Thread operations
    void begin_atomic() { Thread::lock(); }
    void end_atomic() { Thread::unlock(); }

    void sleep() { Thread::yield(); }
    void wakeup() { end_atomic(); }
    void wakeup_all() { end_atomic(); }
};
```

```
class Mutex: protected Synchronizer_Common
{
public:
    Mutex();
    ~Mutex();

    void lock();
    void unlock();

private:
    volatile bool _locked;
};


class Semaphore: protected Synchronizer_Common
{
public:
    Semaphore(int v = 1);
    ~Semaphore();

    void p();
    void v();

private:
    volatile int _value;
};

__END_SYS

#endif
```





```
// EPOS Synchronizer Components
```

```
#ifndef __synchronizer_h
#define __synchronizer_h

#include <architecture.h>
#include <utility/handler.h>
#include <process.h>
```

Dijkstra's Semaphore with a volatile int

```
{
protected:
    Synchronizer_Common() {}

    // Atomic operations
    bool tsl(volatile bool & lock) {
        return CPU::tsl(lock); }
    int finc(volatile int & number) {
        return CPU::finc(number); }
    int fdec(volatile int & number) {
        return CPU::fdec(number); }

    // Thread operations
    void begin_atomic() { Thread::lock(); }
    void end_atomic() { Thread::unlock(); }

    void sleep() { Thread::yield(); }
    void wakeup() { end_atomic(); }
    void wakeup_all() { end_atomic(); }
};
```

```
class Mutex: protected Synchronizer_Common
{
public:
    Mutex();
    ~Mutex();
```

```
    void lock();
    void unlock();
```

```
private:
    volatile bool _locked;
};
```


```
class Semaphore: protected Synchronizer_Common
{
public:
    Semaphore(int v = 1);
    ~Semaphore();
```

```
    void p();
    void v();
```

```
private:
    volatile int _value;
};
```

```
__END_SYS
```

```
#endif
```



```
// EPOS Semaphore Component Implementation
```

```
#include <synchronizer.h>
```

```
__BEGIN_SYS
```

```
Semaphore::Semaphore(int v): _value(v)
{
    db<Synchronizer>(TRC) << "Semaphore(value="
                        << _value << ") => "
                        << this << endl;
}
```

```
Semaphore::~Semaphore()
{
    db<Synchronizer>(TRC) << "~Semaphore(this="
                        << this << ")" << endl;
}
```


```
void Semaphore::p()
{
    db<Synchronizer>(TRC) << "Semaphore::p(this="
                        << this << ",value="
                        << _value << ")" << endl;

    begin_atomic();
    while(_value < 1)
        sleep(); // implicit end_atomic()
    fdec(_value);
    end_atomic();
}
```

```
void Semaphore::v()
{
    db<Synchronizer>(TRC)
        << "Semaphore::v(this=" << this
        << ",value="
        << _value << ")" << endl;

    begin_atomic();
    finc(_value);
    if(_value < 1)
        wakeup(); // implicit end_atomic()
    else
        end_atomic();
}

__END_SYS
```



```
// EPOS Semaphore Component Implementation
```

```
#include <synchronizer.h>
```

```
__BEGIN_SYS
```

```
Semaphore::Semaphore(int v): _value(v)
```

```
{  
    db<Synchronizer>(TRC) << "Semaphore(value=" << _value << ") => "  
    << this << endl;  
}
```

```
Semaphore::~Semaphore()
```

```
{  
    db<Synchronizer>(TRC) << "~Semaphore(this=" << this << ")" << endl;  
}
```

```
void Semaphore::p()
```


```
{  
    db<Synchronizer>(TRC) << "Semaphore::p(this=" << this << ",value=" << _value << ")" << endl;  
  
    begin_atomic();  
    while(_value < 1)  
        sleep(); // implicit end_atomic()  
    fdec(_value);  
    end_atomic();  
}
```

initialization

```
<< "Semaphore::v(this=" << this  
<< ",value=" << _value << ")" << endl;
```

```
begin_atomic();  
finc(_value);  
if(_value < 1)  
    wakeup(); // implicit end_atomic()  
else  
    end_atomic();  
}
```

```
__END_SYS
```



```
// EPOS Semaphore Component Implementation
```

```
#include <synchronizer.h>
```

```
__BEGIN_SYS
```

```
Semaphore::Semaphore(int v): _value(v)
{
    db<Synchronizer>(TRC) << "Semaphore(value="
                        << _value << ") => "
                        << this << endl;
}
```

```
Semaphore::~Semaphore()
{
    db<Synchronizer>(TRC) << "~Semaphore(this="
                        << this << ")" << endl;
}
```


```
void Semaphore::p()
{
    db<Synchronizer>(TRC) << "Semaphore::p(this="
                        << this << ",value="
                        << _value << ")" << endl;

    begin_atomic();
    while(_value < 1)
        sleep(); // implicit end_atomic()
    fdec(_value);
    end_atomic();
}
```

```
void Semaphore::v()
{
    db<Synchronizer>(TRC)
        << "Semaphore::v(this=" << this
        << ",value="
        << _value << ")" << endl;

    begin_atomic();
    finc(_value);
    if(_value < 1)
        wakeup(); // implicit end_atomic()
    else
        end_atomic();
}

__END_SYS
```



disable preemption via Thread::lock()

```
// EPOS Semaphore Component Implementation
```

```
#include <synchronizer.h>
```

```
__BEGIN_SYS
```

```
Semaphore::Semaphore(int v): _value(v)
{
    db<Synchronizer>(TRC) << "Semaphore(value="
                        << _value << ") => "
                        << this << endl;
}
```

```
Semaphore::~Semaphore()
{
    db<Synchronizer>(TRC) << "~Semaphore(this="
                        << this << ")" << endl;
}
```


```
void Semaphore::p()
{
    db<Synchronizer>(TRC) << "Semaphore::p(this="
                        << this << ",value="
                        << _value << ")" << endl;

    begin_atomic();
    while(_value < 1)
        sleep(); // implicit end_atomic()
    fdec(_value);
    end_atomic();
}
```

```
void Semaphore::v()
{
    db<Synchronizer>(TRC)
        << "Semaphore::v(this=" << this
        << ",value="
        << _value << ")" << endl;

    begin_atomic();
    finc(_value);
    if(_value < 1)
        wakeup(); // implicit end_atomic()
    else
        end_atomic();
}

__END_SYS
```



atomic decrement

```
// EPOS Semaphore Component Implementation
```

```
#include <synchronizer.h>
```

```
__BEGIN_SYS
```

```
Semaphore::Semaphore(int v): _value(v)
{
    db<Synchronizer>(TRC) << "Semaphore(value="
                        << _value << ") => "
                        << this << endl;
}
```

```
Semaphore::~Semaphore()
{
    db<Synchronizer>(TRC) << "~Semaphore(this="
                        << this << ")" << endl;
}
```

```
void Semaphore::p()
{
    db<Synchronizer>(TRC) << "Semaphore::p(this="
                        << this << ",value="
                        << _value << ")" << endl;
```


```
    begin_atomic();
    while(_value < 1)
        sleep(); // implicit end_atomic()
    fdec(_value);
    end_atomic();
}
```

```
void Semaphore::v()
{
    db<Synchronizer>(TRC)
        << "Semaphore::v(this=" << this
        << ",value="
        << _value << ")" << endl;

    begin_atomic();
    finc(_value);
    if(_value < 1)
        wakeup(); // implicit end_atomic()
    else
        end_atomic();
}
```

```
__END_SYS
```

```
void sleep() { Thread::yield(); }
```



```
// EPOS Semaphore Component Implementation
```

```
#include <synchronizer.h>
```

```
__BEGIN_SYS
```

```
Semaphore::Semaphore(int v): _value(v)
```

```
{
```

atomic increment

```
    Semaphore(value="v", _value=v) => "v" << endl;
```

```
}
```

```
Semaphore::~Semaphore()
```

```
{
```

```
    db<Synchronizer>(TRC) << "~Semaphore(this=" << this << ")" << endl;
```

```
}
```

```
void Semaphore::p()
```

```
{
```

```
    db<Synchronizer>(TRC) << "Semaphore::p(this=" << this << ", value=" << _value << ")" << endl;
```

```
    begin_atomic();
```

```
    while(_value < 1)
```

```
        sleep(); // implicit end_atomic()
```

```
    fdec(_value);
```

```
    end_atomic();
```

```
}
```

```
void Semaphore::v()
```

```
{
```

```
    db<Synchronizer>(TRC)
```

```
        << "Semaphore::v(this=" << this
```

```
        << ", value="
```

```
        << _value << ")" << endl;
```

```
    begin_atomic();
```

```
    finc(_value);
```

```
    if(_value < 1)
```


```
        wakeup(); // implicit end_atomic()
```

```
    else
```

```
        end_atomic();
```

```
}
```

```
__END_SYS
```



```
// EPOS Semaphore Component Implementation
```

```
#include <synchronizer.h>
```

```
__BEGIN_SYS
```

```
Semaphore::Semaphore(int v): value(v)
```

```
void wakeup() { end_atomic(); }
```

```
db<Synchronizer>(TRC) << "Semaphore(value="
                        << _value << ") => "
                        << this << endl;
```

```
}
```

```
Semaphore::~Semaphore()
```

```
{
```

```
db<Synchronizer>(TRC) << "~Semaphore(this="
                        << this << ")" << endl;
```

```
}
```

```
void Semaphore::p()
```

```
{
```

```
db<Synchronizer>(TRC) << "Semaphore::p(this="
                        << this << ",value="
                        << _value << ")" << endl;
```

```
begin_atomic();
```

```
while(_value < 1)
```

```
sleep(); // implicit end_atomic()
```

```
fdec(_value);
```

```
end_atomic();
```

```
}
```

```
void Semaphore::v()
```

```
{
```

```
db<Synchronizer>(TRC)
```

```
<< "Semaphore::v(this=" << this
```

```
<< ",value="
```

```
<< _value << ")" << endl;
```

```
begin_atomic();
```

```
finc(_value);
```

```
if(_value < 1)
```


```
wakeup(); // implicit end_atomic()
```

```
else
```

```
end_atomic();
```

```
}
```

```
__END_SYS
```





```
// EPOS Mutex Component Implementation
```

```
#include <synchronizer.h>
```

```
__BEGIN_SYS
```

```
Mutex::Mutex(): _locked(false)
```

```
{  
    db<Synchronizer>(TRC) << "Mutex() => "  
                        << this << endl;  
}
```

```
Mutex::~Mutex()
```

```
{  
    db<Synchronizer>(TRC) << "~Mutex(this=" << this << ")" << endl;  
}
```


```
void Mutex::lock()
```

```
{  
    db<Synchronizer>(TRC) << "Mutex::lock(this=" << this << ")" << endl;  
  
    begin_atomic();  
    if(tsl(_locked))  
        while(tsl(_locked))  
            sleep(); // implicit end_atomic()  
    else  
        end_atomic();  
}
```

```
void Mutex::unlock()
```

```
{  
    db<Synchronizer>(TRC)  
        << "Mutex::unlock(this=" << this  
        << ")" << endl;  
  
    begin_atomic();  
    _locked = false;  
    wakeup(); // implicit end_atomic()  
}
```

```
__END_SYS
```



```
// EPOS Mutex Component Implementation
```

```
#include <synchronizer.h>
```

```
__BEGIN_SYS
```

```
Mutex::Mutex(): _locked(false)
```

```
{  
    db<Synchronizer>(TRC) << "Mutex() => "  
    << this << endl;  
}
```

```
Mutex::~Mutex()
```

```
{  
    db<Synchronizer>(TRC) << "~Mutex(this=" << this << ")" << endl;  
}
```

```
void Mutex::lock()
```


```
{  
    db<Synchronizer>(TRC) << "Mutex::lock(this=" << this << ")" << endl;  
  
    begin_atomic();  
    if(tsl(_locked))  
        while(tsl(_locked))  
            sleep(); // implicit end_atomic()  
    else  
        end_atomic();  
}
```

initialization

```
<< "Mutex::unlock(this=" << this  
<< ")" << endl;
```

```
begin_atomic();  
_locked = false;  
wakeup(); // implicit end_atomic()  
}
```

```
__END_SYS
```



```
// EPOS Mutex Component Implementation
```

```
#include <synchronizer.h>
```

```
__BEGIN_SYS
```

```
Mutex::Mutex(): _locked(false)
```

```
{  
    db<Synchronizer>(TRC) << "Mutex() => "  
    << this << endl;  
}
```

```
Mutex::~Mutex()
```

```
{  
    db<Synchronizer>(TRC) << "~Mutex(this=" << this << ")" << endl;  
}
```

```
void Mutex::lock()
```


```
{  
    db<Synchronizer>(TRC) << "Mutex::lock(this=" << this << ")" << endl;  
  
    begin_atomic();  
    if(tsl(_locked))  
        while(tsl(_locked))  
            sleep(); // implicit end_atomic()  
    else  
        end_atomic();  
}
```

```
void Mutex::unlock()
```

```
{  
    db<Synchronizer>(TRC)  
    << "Mutex::unlock(this=" << this  
    << ")" << endl;  
  
    begin_atomic();  
    _locked = false;  
    wakeup(); // implicit end_atomic()  
}
```

```
__END_SYS
```

tsl



```
// EPOS Mutex Component Implementation
```

```
#include <synchronizer.h>
```

```
__BEGIN_SYS
```

```
Mutex::Mutex(): _locked(false)
```

```
{  
    db<Synchronizer>(TRC) << "Mutex() => "  
    << this << endl;  
}
```

```
Mutex::~Mutex()
```

```
{  
    db<Synchronizer>(TRC) << "~Mutex(this=" << this << ")" << endl;  
}
```

```
void Mutex::lock()
```


```
{  
    db<Synchronizer>(TRC) << "Mutex::lock(this=" << this << ")" << endl;  
  
    begin_atomic();  
    if(tsl(_locked))  
        while(tsl(_locked))  
            sleep(); // implicit end_atomic()  
    else  
        end_atomic();  
}
```

```
void Mutex::unlock()
```

```
{  
    db<Synchronizer>(TRC)  
    << "Mutex::unlock(this=" << this  
    << ")" << endl;  
  
    begin_atomic();  
    _locked = false;  
    wakeup(); // implicit end_atomic()  
}
```

```
__END_SYS
```

```
void sleep() { Thread::yield(); }
```



```
// EPOS Mutex Component Implementation
```

```
#include <synchronizer.h>
```

```
__BEGIN_SYS
```

```
Mutex::Mutex(): _locked(false)
```

```
{ void wakeup() { end_atomic(); }
```

```
db<Synchronizer>(TRC) << "Mutex::Mutex()" << this << endl;
```

```
}
```

```
Mutex::~Mutex()
```

```
{
```

```
db<Synchronizer>(TRC) << "~Mutex(this=" << this << ")" << endl;
```

```
}
```

```
void Mutex::lock()
```

```
{
```

```
db<Synchronizer>(TRC) << "Mutex::lock(this=" << this << ")" << endl;
```

```
begin_atomic();
```

```
if(tsl(_locked))
```

```
while(tsl(_locked))
```

```
sleep(); // implicit end_atomic()
```

```
else
```

```
end_atomic();
```

```
}
```

```
void Mutex::unlock()
```

```
{
```

```
db<Synchronizer>(TRC)
```

```
<< "Mutex::unlock(this=" << this
```

```
<< ")" << endl;
```


```
begin_atomic();
```

```
_locked = false;
```

```
wakeup(); // implicit end_atomic()
```

```
}
```

```
__END_SYS
```



```

void Thread::yield()
{
    lock();

    db<Thread>(TRC) << "Thread::yield(running="
                << _running << ")" << endl;

    if(!_ready.empty()) {
        Thread * prev = _running;
        prev->_state = READY;
        _ready.insert(&prev->_link);

        _running = _ready.remove()->object();
        _running->_state = RUNNING;

        dispatch(prev, _running);
    } else
        idle(); // implicit unlock()

    unlock();
}

void Thread::dispatch(Thread * prev, Thread * next)
{
    if(prev != next) {
        assert(prev->_state != RUNNING);
        assert(next->_state == RUNNING);


        db<Thread>(TRC) << "Thread::dispatch(prev="
                        << prev << ",next="
                        << next << ")" << endl;
    }
}

```

```

db<Thread>(INF) << "prev={" << prev
               << ",ctx="
               << *prev->_context
               << "}" << endl;

```



```

db<Thread>(INF) << "next={" << next
               << ",ctx=" << *next->_context
               << "}" << endl;

```

```

CPU::switch_context(&prev->_context,
                   next->_context);

```

```

}

```

```

unlock();

```

```

}

```

```

int Thread::idle()
{

```

```

    db<Thread>(TRC) << "Thread::idle()" << endl;

```

```

    db<Thread>(INF)
    << "There are no runnable threads at the moment!"
    << endl;
    db<Thread>(INF) << "Halting the CPU ..." << endl;

```

```

    unlock();
    CPU::halt();

```

```

    return 0;

```

```

}

```

```

void Thread::yield()
{
    lock();

    db<Thread>(TRC) << "Thread::yield(running="
                  << _running << ")" << endl;


```

**yields** the CPU by moving running into  
READY and removing its head

```

    << prev
context
endl;

```



```

if(!_ready.empty()) {
    Thread * prev = _running;
    prev->_state = READY;
    _ready.insert(&prev->_link);

    _running = _ready.remove()->object();
    _running->_state = RUNNING;

    dispatch(prev, _running);
} else
    idle(); // implicit unlock()

unlock();
}

void Thread::dispatch(Thread * prev, Thread * next)
{
    if(prev != next) {
        assert(prev->_state != RUNNING);
        assert(next->_state == RUNNING);

        db<Thread>(TRC) << "Thread::dispatch(prev="
                        << prev << ",next="
                        << next << ")" << endl;

```

```

    db<Thread>(INF) << "next={" << next
                  << ",ctx=" << *next->_context
                  << "}" << endl;

```

```

    CPU::switch_context(&prev->_context,
                       next->_context);

```

```

}

```

```

unlock();

```

```

}

```

```

int Thread::idle()
{

```

```

    db<Thread>(TRC) << "Thread::idle()" << endl;

```

```

    db<Thread>(INF)
    << "There are no runnable threads at the moment!"
    << endl;
    db<Thread>(INF) << "Halting the CPU ..." << endl;

```

```

    unlock();
    CPU::halt();

```

```

    return 0;

```

```

}

```

```

void Thread::yield()
{
    lock();

    db<Thread>(TRC) << "Thread::yield(running="
        << _running << ")" << endl;

    if(!_ready.empty()) {
        Thread * prev = _running;
        prev->_state = READY;
        _ready.insert(&prev->_link);

        _running = _ready.remove()->object();
        _running->_state = RUNNING;

        dispatch(prev, _running);
    } else
        idle(); // implicit unlock()

    unlock();
}

```

```

void Thread::dispatch(Thread * prev, Thread * next)
{
    if(prev != next) {
        assert(prev->_state != RUNNING);
        assert(next->_state == RUNNING);


        db<Thread>(TRC) << "Thread::dispatch(prev="
            << prev << ",next="
            << next << ")" << endl;
    }
}

```

```

db<Thread>(INF) << "prev={" << prev
    << ",ctx="
    << *prev->_context
    << "}" << endl;

```



dispatches the new running

```

    << "next={" << next
    << ",ctx=" << *next->_context
    << "}" << endl;

    CPU::switch_context(&prev->_context,
        next->_context);
}

```

```

unlock();
}

int Thread::idle()
{
    db<Thread>(TRC) << "Thread::idle()" << endl;

```

```

    db<Thread>(INF)
    << "There are no runnable threads at the moment!"
    << endl;
    db<Thread>(INF) << "Halting the CPU ..." << endl;

    unlock();
    CPU::halt();

    return 0;
}

```



```
void Thread::yield()
{
    lock();

    db<Thread>(TRC) << "Thread::yield(running="
                << _running << ")" << endl;
```

```
    if(!_ready.empty()) {
        Thread * prev = _running;
        prev->_state = READY;
        _ready.insert(&prev->_link);

        _running = _ready.remove()->object();
        _running->_state = RUNNING;


        dispatch(prev, _running);
    } else
        idle(); // implicit unlock()

    unlock();
}
```

```
void Thread::dispatch(Thread * prev, Thread * next)
{
    if(prev != next) {
        assert(prev->_state != RUNNING);
        assert(next->_state == RUNNING);

        db<Thread>(TRC) << "Thread::dispatch(prev="
                        << prev << ",next="
                        << next << ")" << endl;
```

```
        db<Thread>(INF) << "prev={" << prev
                        << ",ctx="
                        << *prev->_context
                        << "}" << endl;
```



```
        db<Thread>(INF) << "next={" << next
                        << ",ctx=" << *next->_context
                        << "}" << endl;
```

or goes into idle if READY is empty

```
        dispatch(&prev->_context,
                 next->_context);
    }

    unlock();
}
```

```
int Thread::idle()
{
    db<Thread>(TRC) << "Thread::idle()" << endl;

    db<Thread>(INF)
    << "There are no runnable threads at the moment!"
    << endl;
    db<Thread>(INF) << "Halting the CPU ..." << endl;

    unlock();
    CPU::halt();

    return 0;
}
```

```
void Thread::yield()
{
    lock();

    db<Thread>(TRC) << "Thread::yield(running="
                << _running << ")" << endl;
```

```
    if(!_ready.empty()) {
        Thread * prev = _running;
        prev->_state = READY;
        _ready.insert(&prev->_link);

        _running = _ready.remove()->object();
        _running->_state = RUNNING;
```

```
        dispatch(prev, _running);
    } else
        idle(); // implicit unlock()
```


```
unlock();
}
```

```
void Thread::dispatch(Thread * prev, Thread * next)
{
```

```
    if(prev != next) {
        assert(prev->_state != RUNNING);
        assert(next->_state == RUNNING);
```

```
        db<Thread>(TRC) << "Thread::dispatch(prev="
                        << prev << ",next="
                        << next << ")" << endl;
```

```
        db<Thread>(INF) << "prev={" << prev
                        << ",ctx="
                        << *prev->_context
                        << "}" << endl;
```



```
        db<Thread>(INF) << "next={" << next
                        << ",ctx=" << *next->_context
                        << "}" << endl;
```

```
        CPU::switch_context(&prev->_context,
                            next->_context);
```

```
    }
```

```
unlock();
```

```
}
```

assert states are correct (see  
#assert definition)

```
db<Thread>(TRC) << "Thread::idle()" << endl;
```

```
db<Thread>(INF)
<< "There are no runnable threads at the moment!"
<< endl;
db<Thread>(INF) << "Halting the CPU ..." << endl;
```

```
unlock();
CPU::halt();
```

```
return 0;
```

```
}
```

```

void Thread::yield()
{
    lock();

    db<Thread>(TRC) << "Thread::yield(running="
        << _running << ")" << endl;

    if(!_ready.empty()) {
        Thread * prev = _running;
        prev->_state = READY;
        _ready.insert(&prev->_link);

        _running = _ready.remove()->object();
        _running->_state = RUNNING;

        dispatch(prev, _running);
    } else
        idle(); // implicit unlock()

    unlock();
}

```

```

void Thread::dispatch(Thread * prev, Thread * next)
{
    if(prev != next) {
        assert(prev->_state != RUNNING);
        assert(next->_state == RUNNING);


        db<Thread>(TRC) << "Thread::dispatch(prev="
            << prev << ",next="
            << next << ")" << endl;
    }
}

```

```

db<Thread>(INF) << "prev={" << prev
    << ",ctx="
    << *prev->_context
    << "}" << endl;

```



```

db<Thread>(INF) << "next={" << next
    << ",ctx=" << *next->_context
    << "}" << endl;

```

```

CPU::switch_context(&prev->_context,
    next->_context);

```

```

}

```

```

unlock();

```

```

}

```

```

int Thread::idle()

```

fine-grain debug control  
<Class>(LEVEL)

```

Thread::idle()" << endl;

```

```

db<Thread>(INF)
    << "There are no runnable threads at the moment!"
    << endl;
db<Thread>(INF) << "Halting the CPU ..." << endl;

unlock();
CPU::halt();

return 0;
}

```

```

void Thread::yield()
{
    lock();

    db<Thread>(TRC) << "Thread::yield(running="
                << _running << ")" << endl;

    if(!_ready.empty()) {
        Thread * prev = _running;
        prev->_state = READY;
        _ready.insert(&prev->_link);

        _running = _ready.remove()->object();
        _running->_state = RUNNING;
    }
}

```

### CPU::halt()

Most CPUs support one or more low-power modes, including a mode in which only registers are preserved and the interrupt signal is monitored. A halted CPU returns to normal operation whenever an interrupt happens. Some CPU support selecting which interrupts are enabled during halt. Returning to normal operation usually takes several CPU cycles, but halting saves a lot of energy.

```


    ( )
    v, Thread * next)
    NING);
    NING);
    ::dispatch(prev="
    ",next="
    ") << endl;
}

```

```

db<Thread>(INF) << "prev={" << prev
               << ",ctx="
               << *prev->_context
               << "}" << endl;

```



```

db<Thread>(INF) << "next={" << next
               << ",ctx=" << *next->_context
               << "}" << endl;

```

```

CPU::switch_context(&prev->_context,
                   next->_context);

```

```

}

```

```

unlock();

```

```

}

```

```

int Thread::idle()
{

```

```

    db<Thread>(TRC) << "Thread::idle()" << endl;

```

```

    db<Thread>(INF)

```

```

    << "There are no runnable threads at the moment!"
    << endl;

```

```

    db<Thread>(INF) << "Halting the CPU ..." << endl;

```

```

    unlock();

```

```

    CPU::halt();

```

```

    return 0;

```

```

}

```